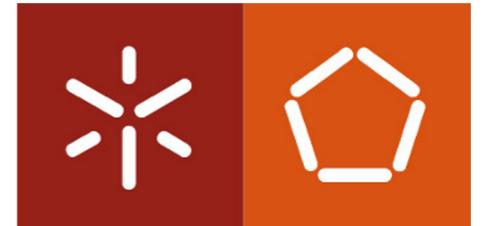


Cloud Computing Applications and Services (Aplicações e Serviços de Computação em Nuvem)

Virtualization Part II

University of Minho
2022/2023



Containers

- Lightweight virtual environment that groups and isolates a set of processes and resources (memory, CPU, disk, ...), from the host and any other containers

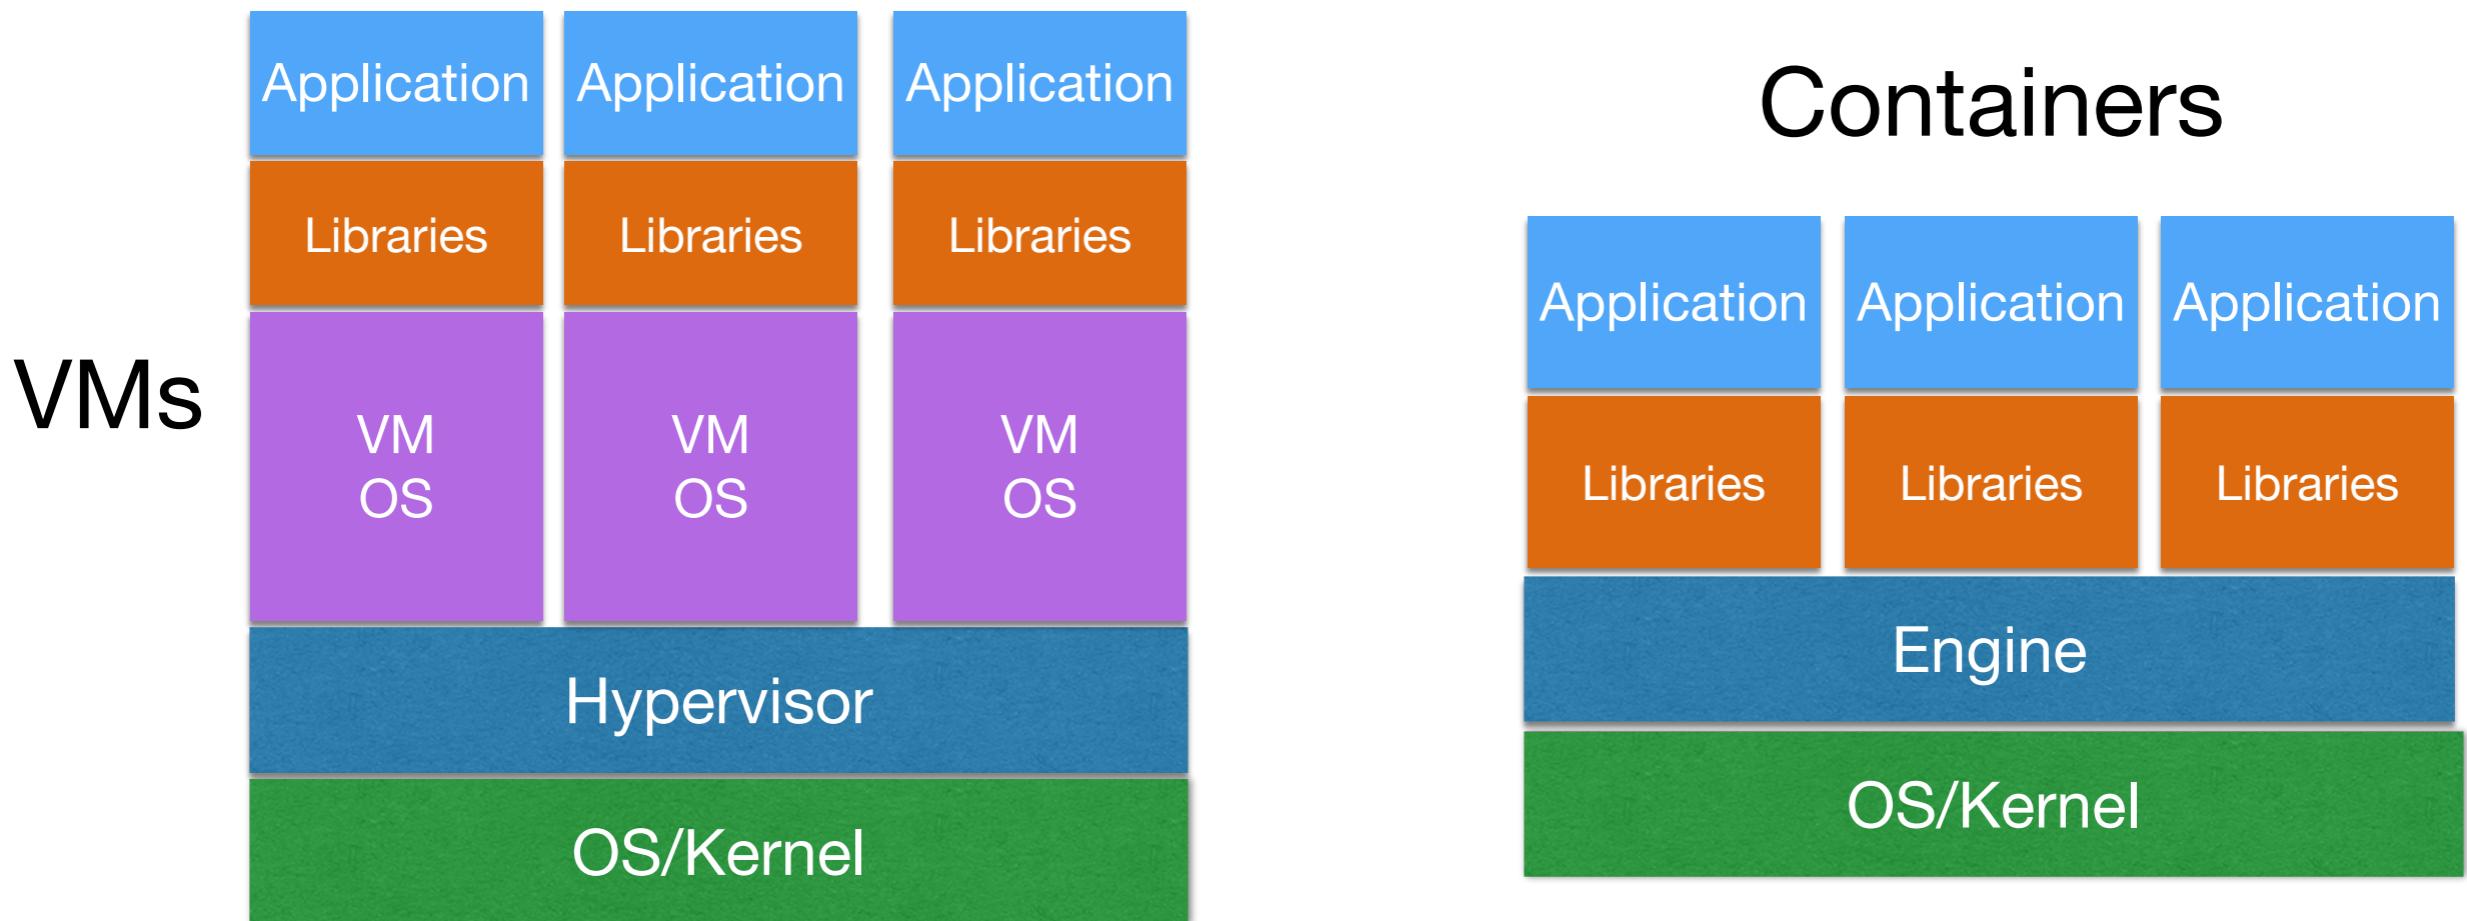
Containers



- Why are containers useful?
 - Running different isolated versions of the same software/application (e.g., database) in a shared OS/Kernel environment
 - Portability/migration across servers
 - Easy packaging of software, applications and their dependencies

Linux Containers

- Containers are lightweight, while not requiring full virtualization of CPU, RAM, network, and storage requests (as in VMs)
- The engine isolates and configures resources (see next slide)



Linux Containers

- The host system is compartmentalised for each container in terms of CPU and I/O (memory, disk, network)
- Each container shares the hardware and kernel/OS with the host system
- Containers are isolated from each other

Linux Containers

Building Blocks

- Namespaces (Isolation)
 - Component of the Linux Kernel that makes a global host resource appear as a dedicated resource for a group of processes running in the same namespace
 - Allows sharing of host resources across different containers without conflicts
 - Network, storage (filesystem), ...

Linux Containers

Building Blocks

- Control Groups (Resource Management)
 - Allows allocating resources (CPU, RAM, Disk I/O, network bandwidth) among groups of processes
 - In other words, it limits the amount of resources used by each container (e.g., storage space)

Linux Containers

Building Blocks

- SELinux (Security)
 - Provides additional security over namespaces so that a container is not able to compromise the host system and other containers
 - Enforces access control and security policies

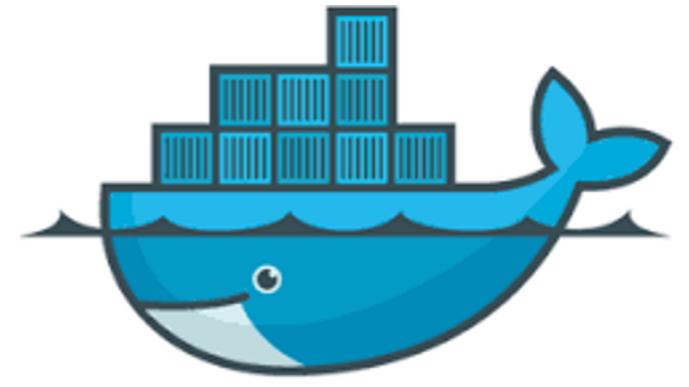
Linux Containers

Types

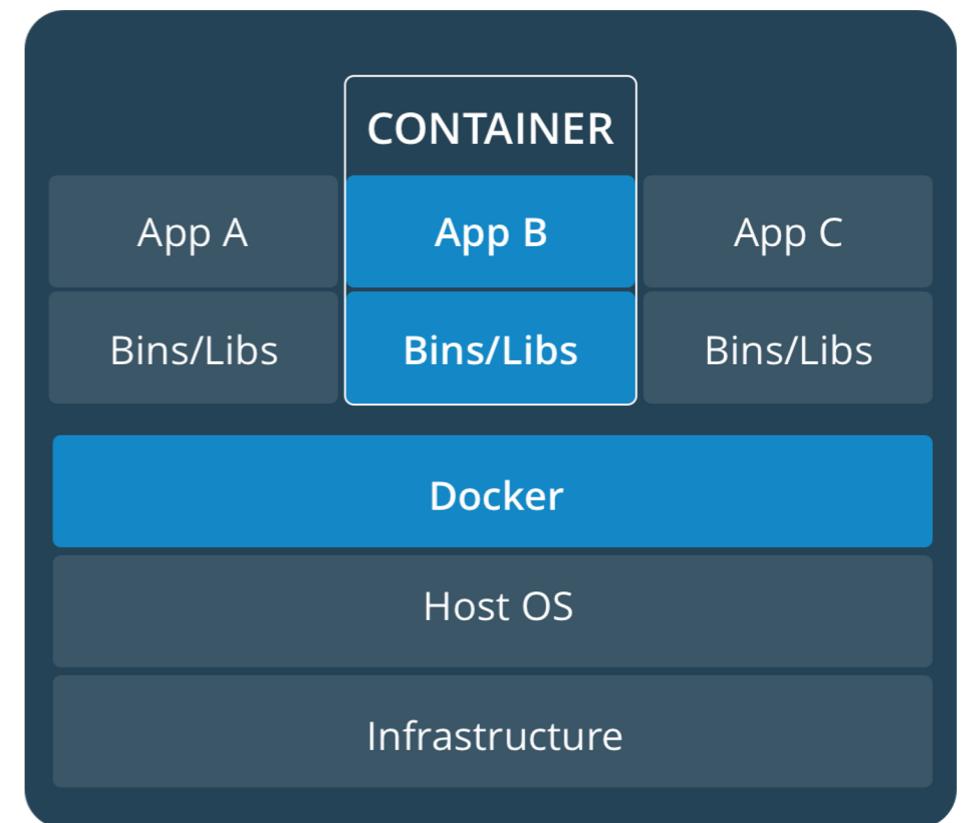
- OS Containers (e.g., LXC)
 - Containers run multiple processes and simulate a “lightweight” operating system
- Application Containers (e.g., Docker)
 - Focused on deploying applications
 - Each application is seen as an independent process

Docker

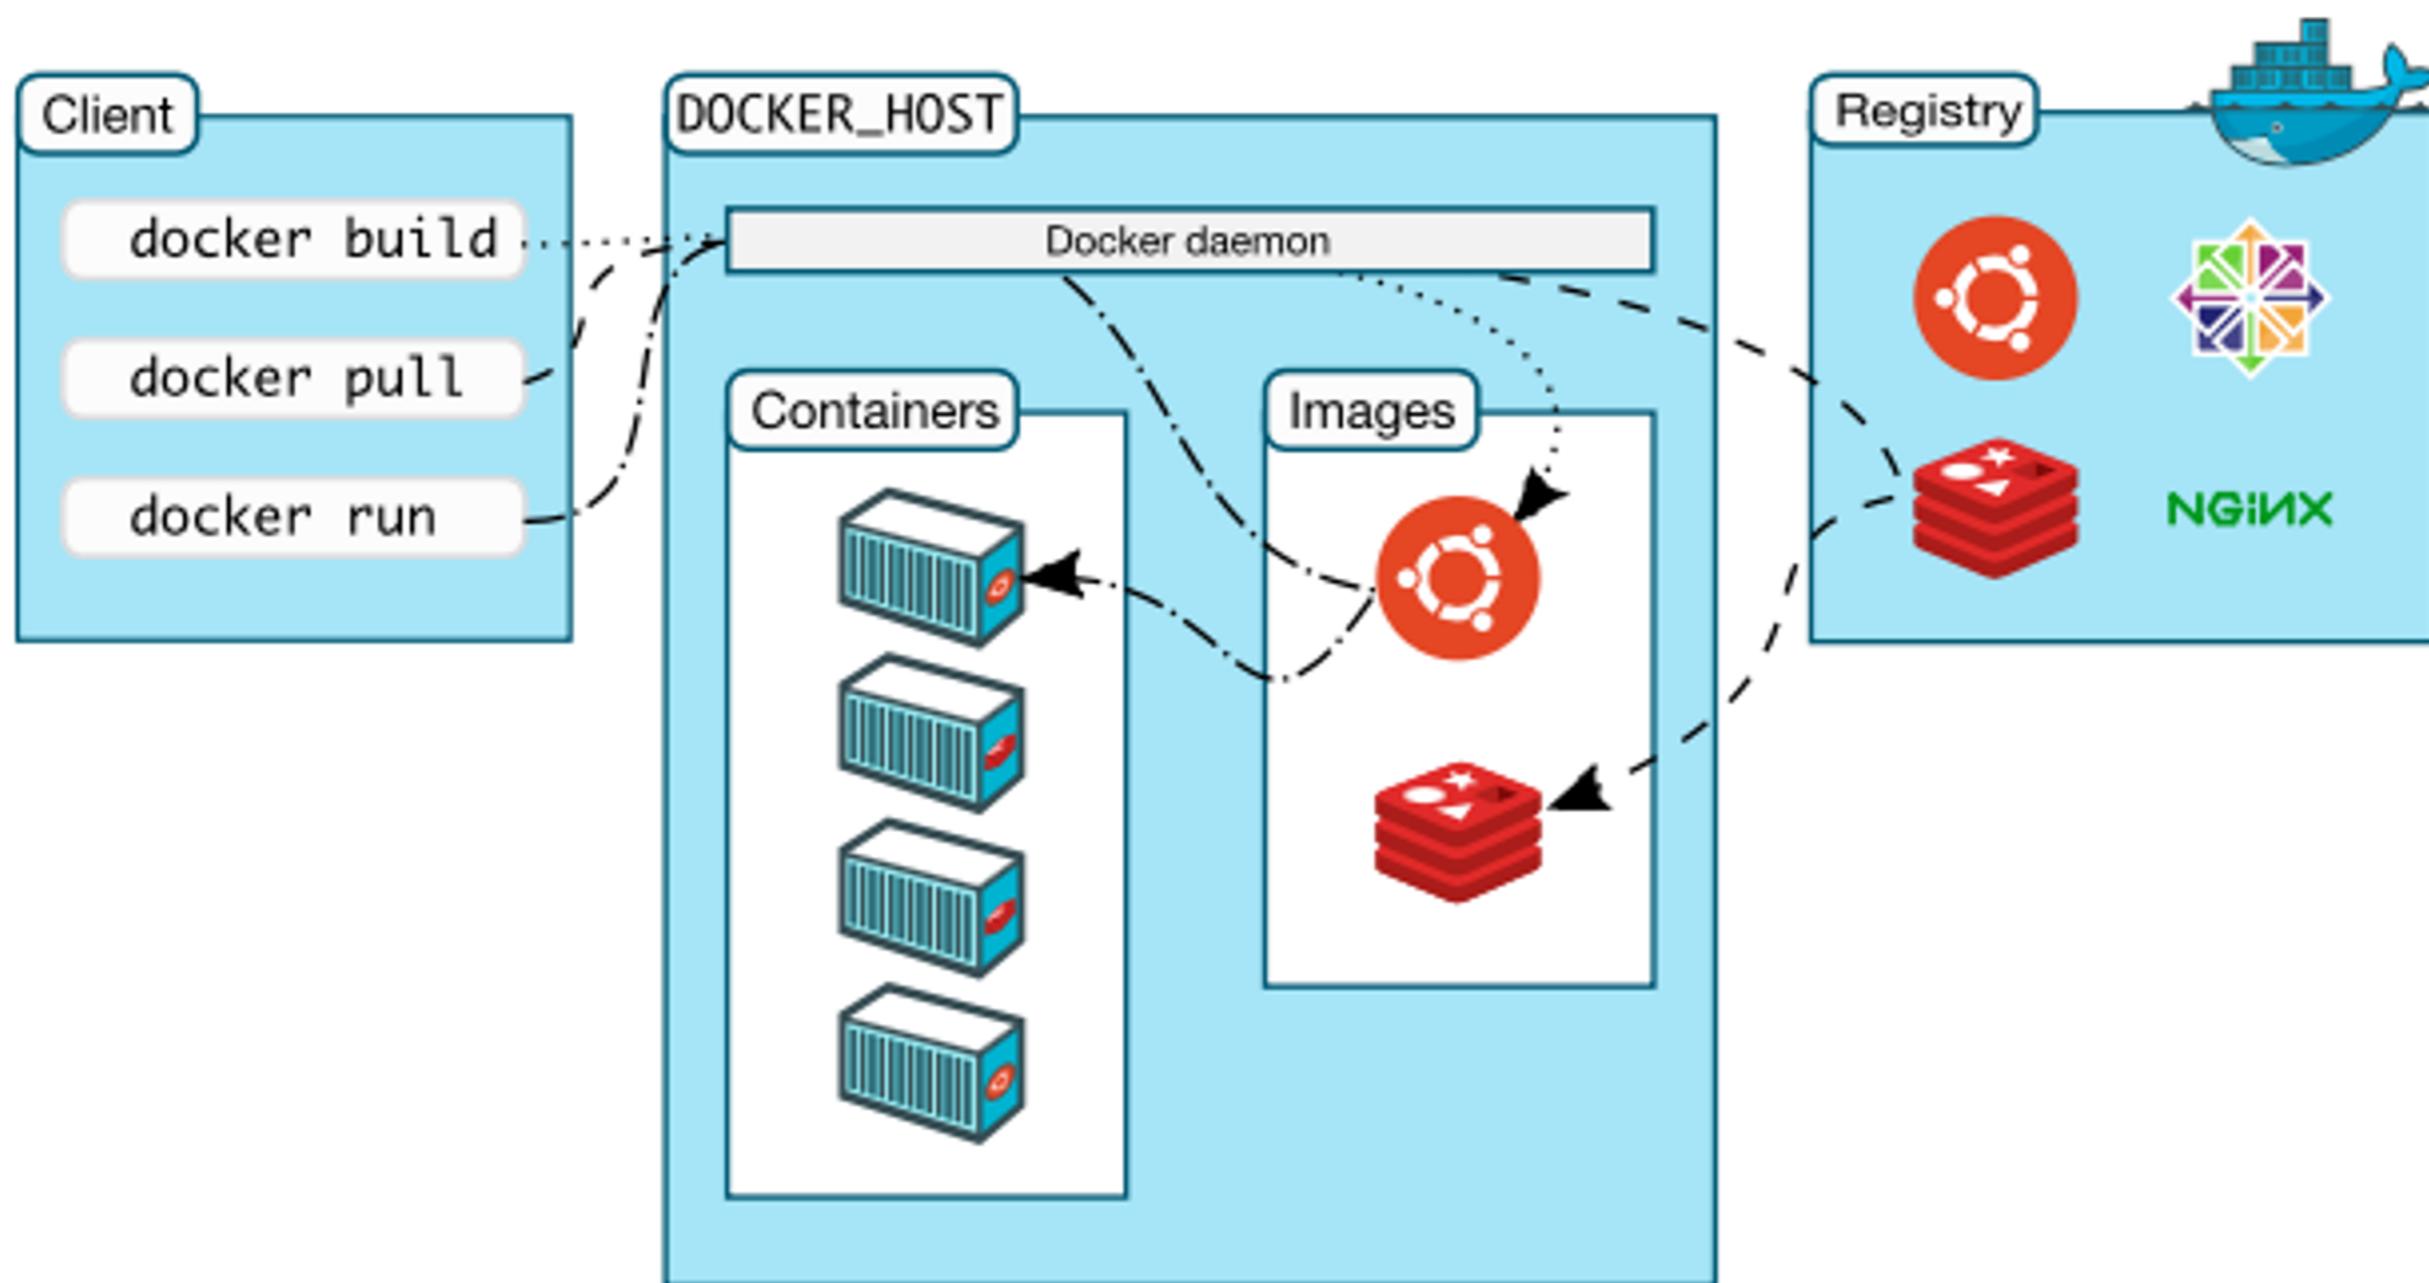
- Most widely-known container platform
- Supports Ubuntu, Fedora, RHEL, CentOS, Windows, etc
- <https://www.docker.com>



docker



Docker



Docker

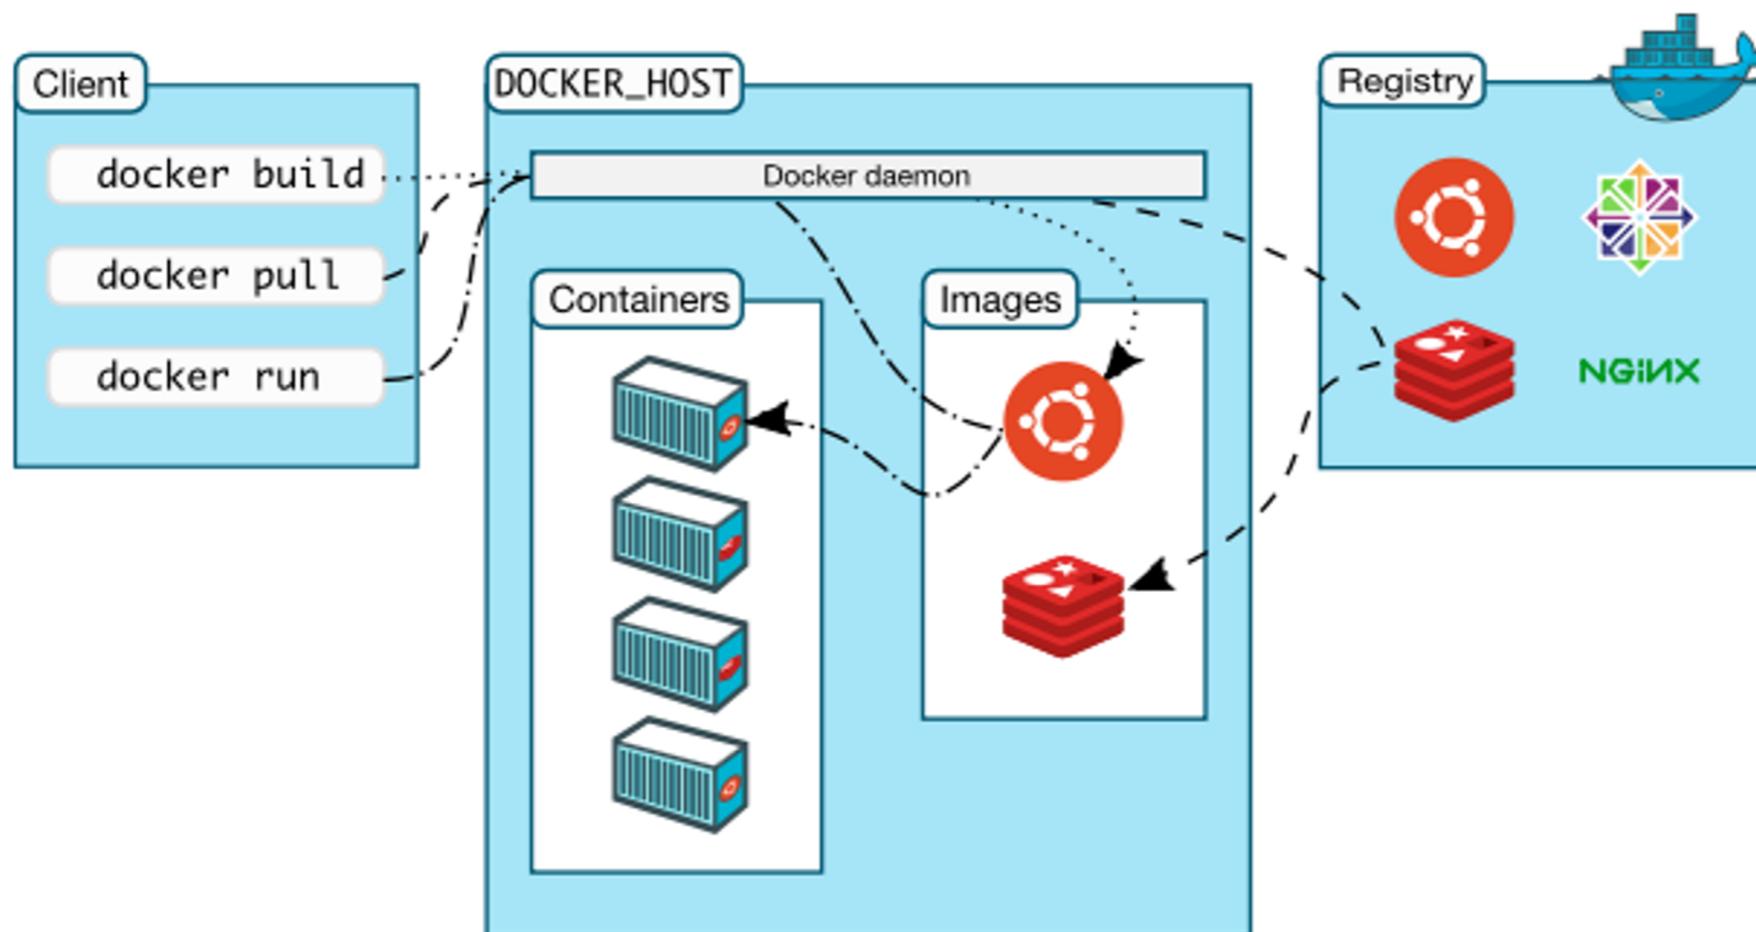
Docker Client

- Component used by users to interact with the Docker platform
- Exposes the Docker API for:
 - Running and managing containers, and networks
 - Reading logs and metrics
 - Pulling and managing images

Docker

Docker Daemon

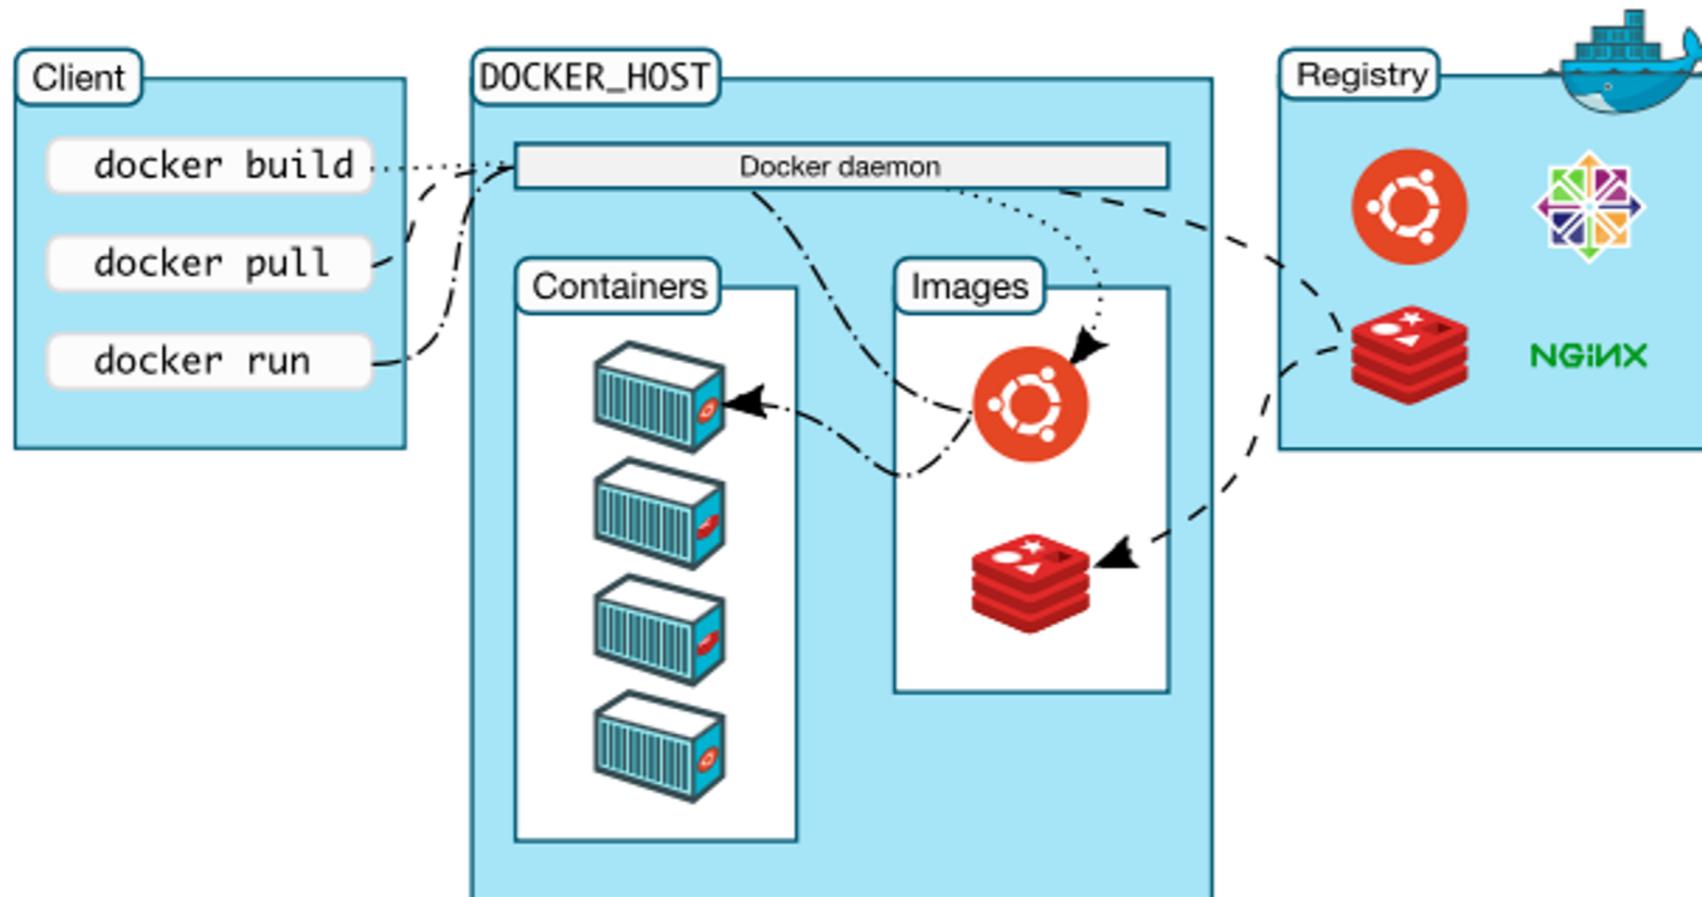
- Uses the Docker API for receiving requests from the Docker Client
- Manages Docker Images, Containers, Networks



Docker

Docker Objects

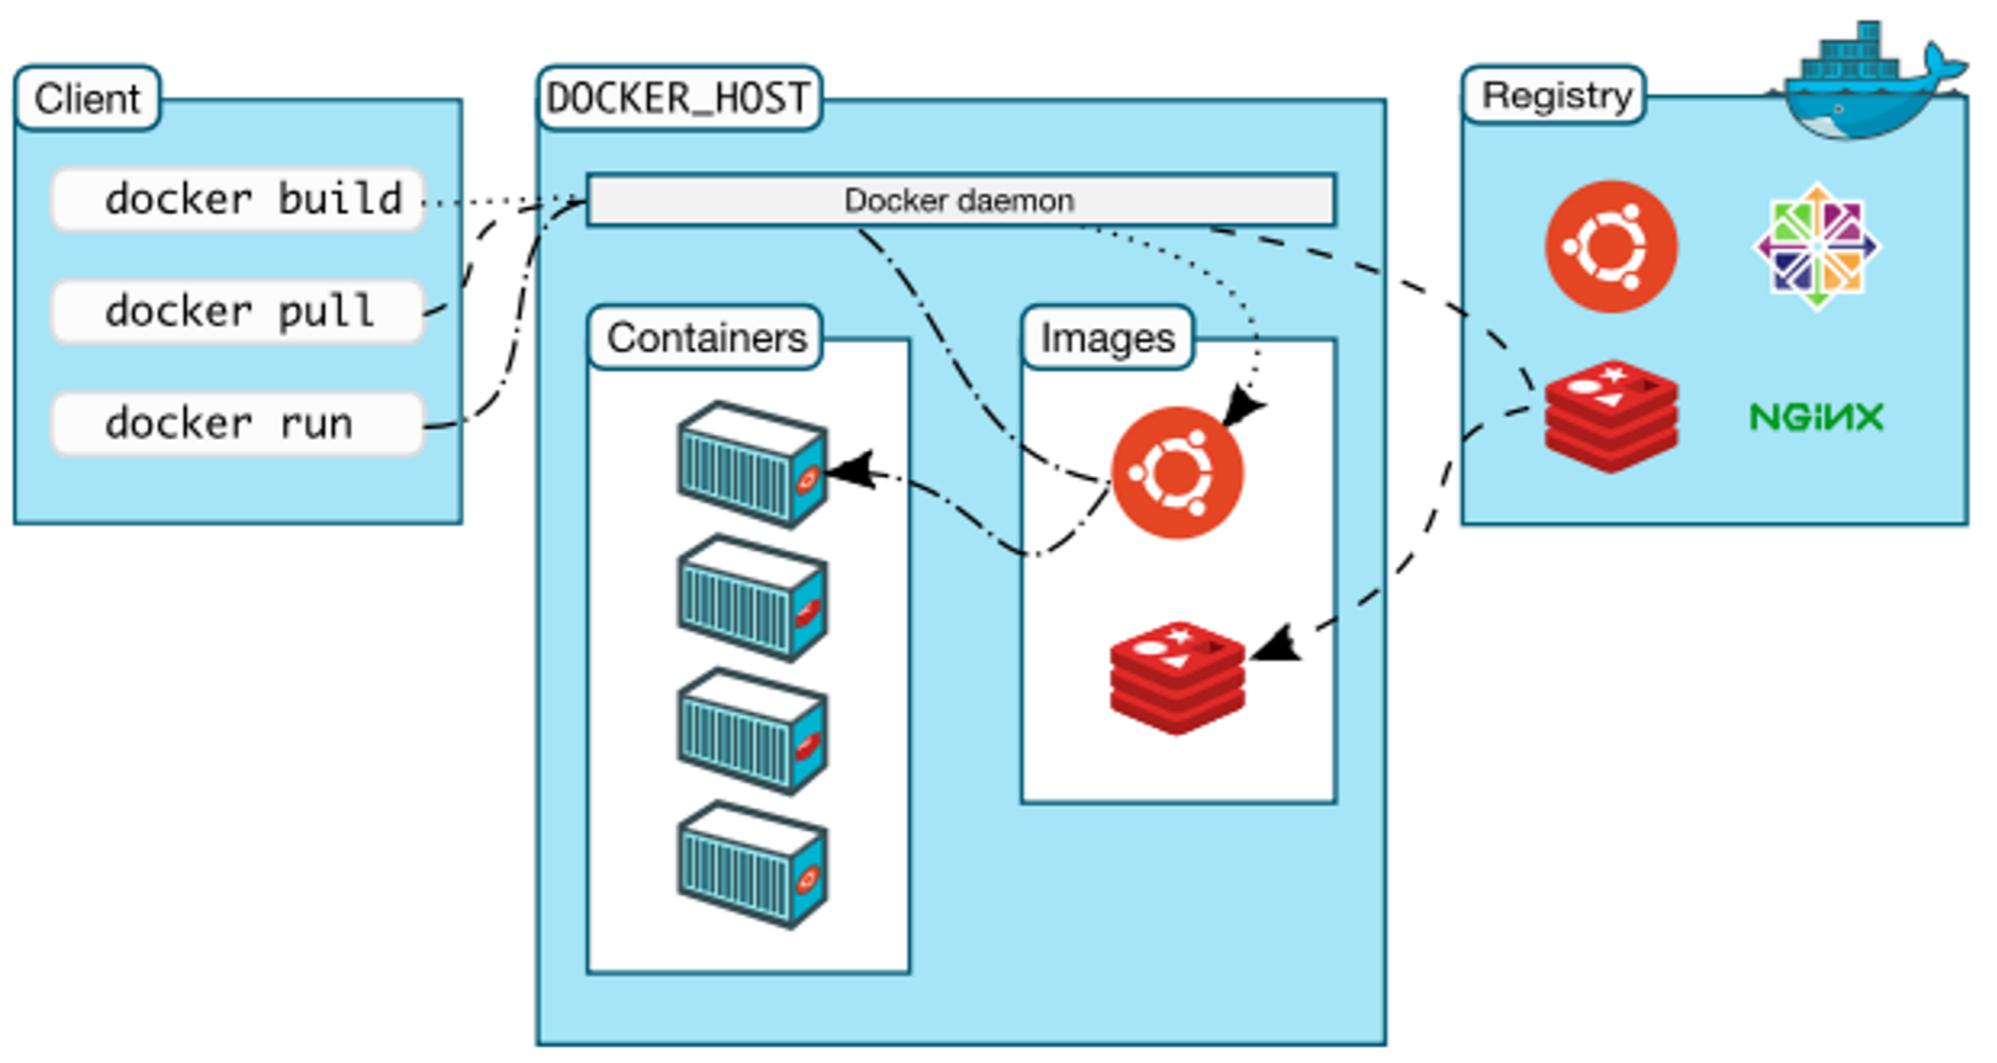
- Image - immutable (unchangeable) file that contains the source code, libraries, and other files needed for an application to run
- Container - runnable instance of an Image



Docker

Docker Registry

- Repository of Docker Images
- E.g., <https://store.docker.com>



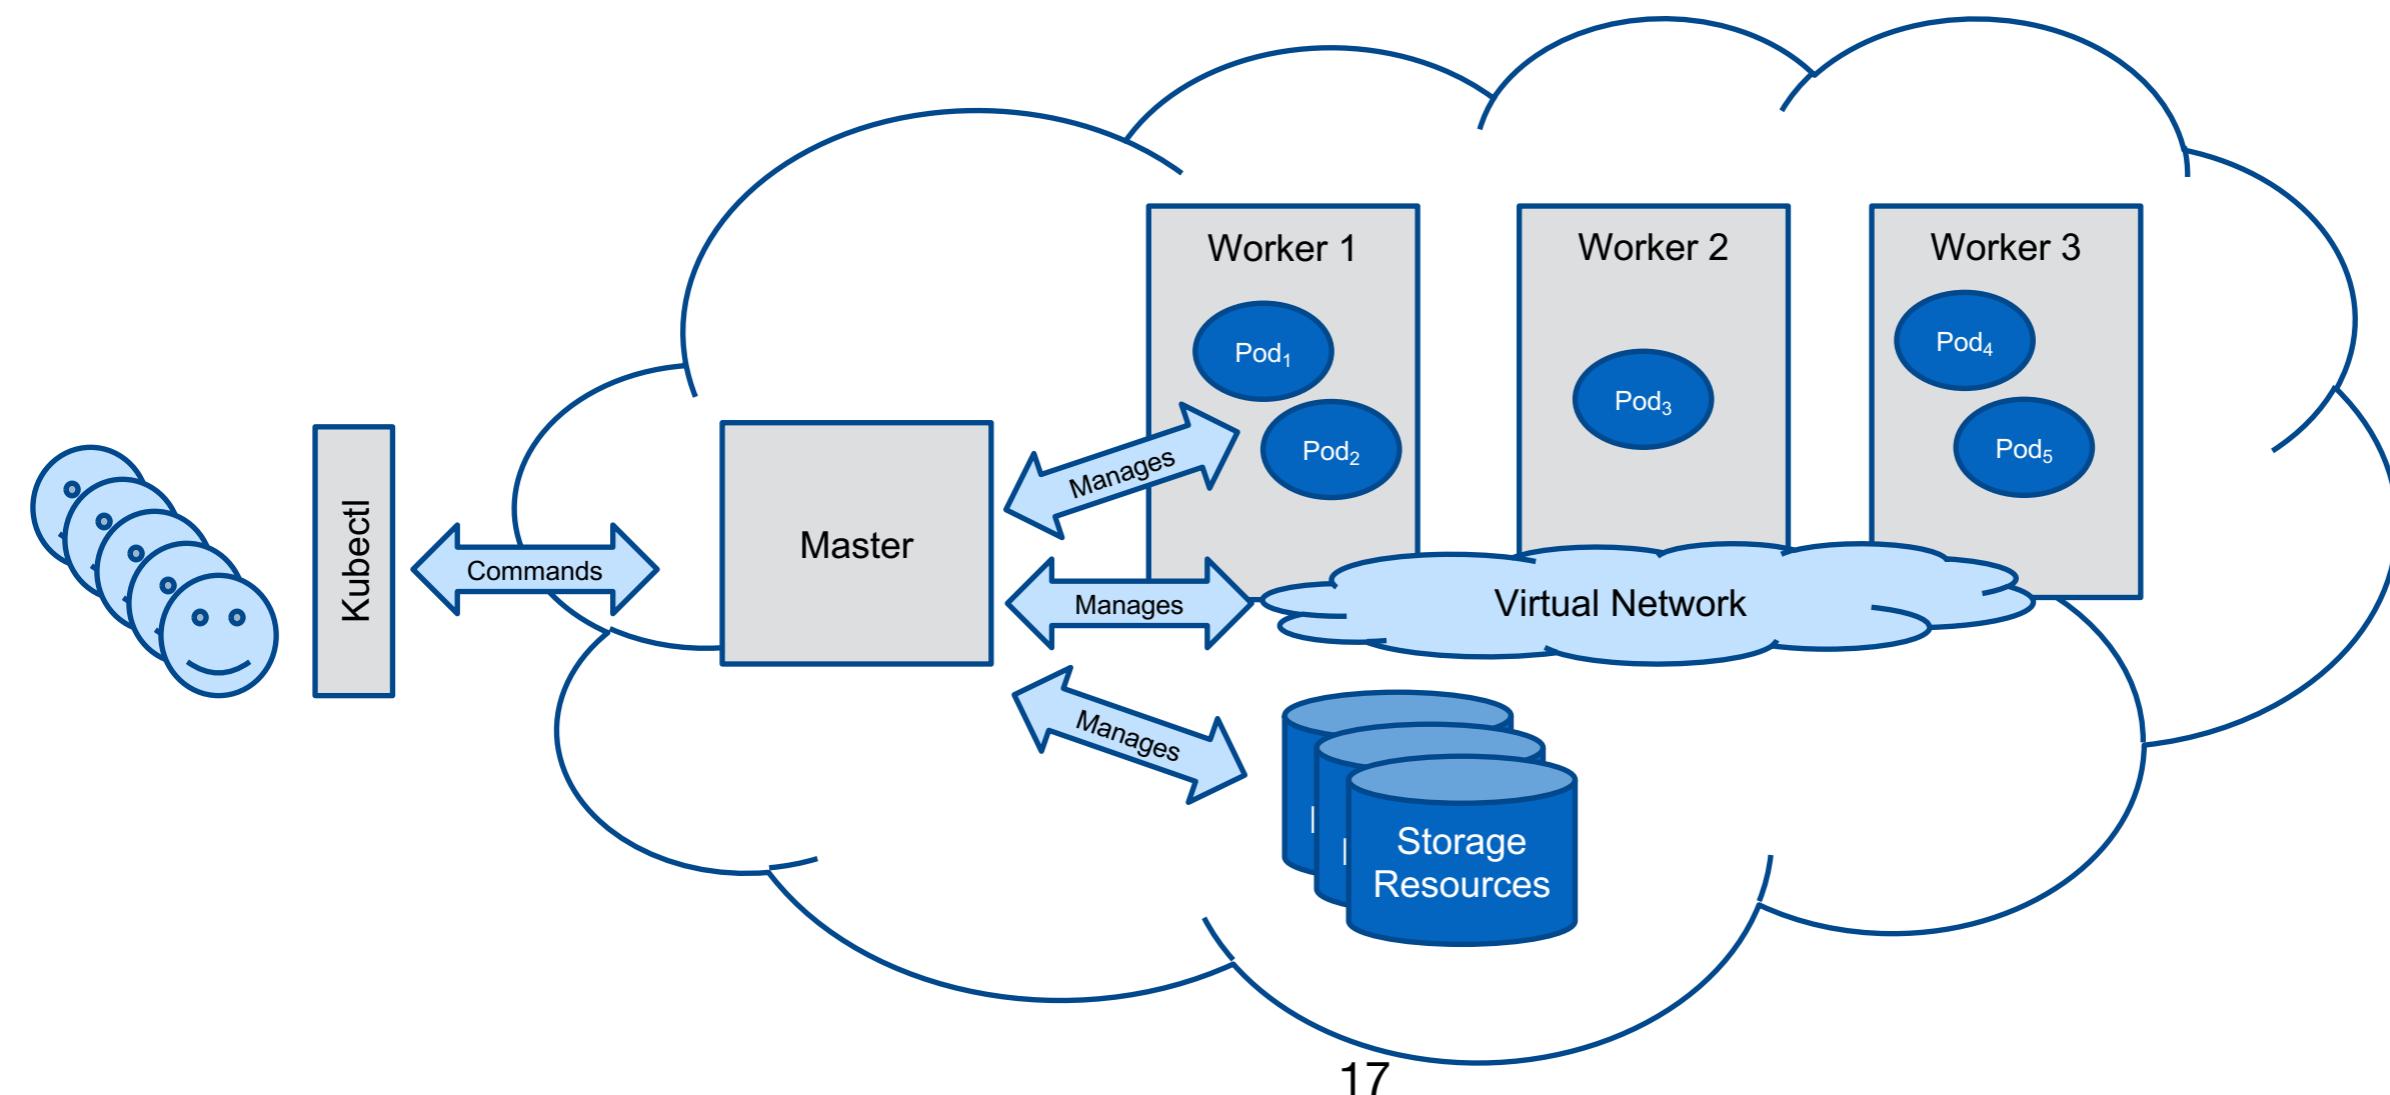
Kubernetes

- Automates the deployment, scaling and management of containers
- Interesting features:
 - Network management
(e.g., service discovery, load balancing)
 - Modular storage orchestration
(e.g., iSCSI, NFS, Ceph, AWS, GCP)
 - Simplified **scheduling, self-healing** and **scale out** for containers
- <https://kubernetes.io>



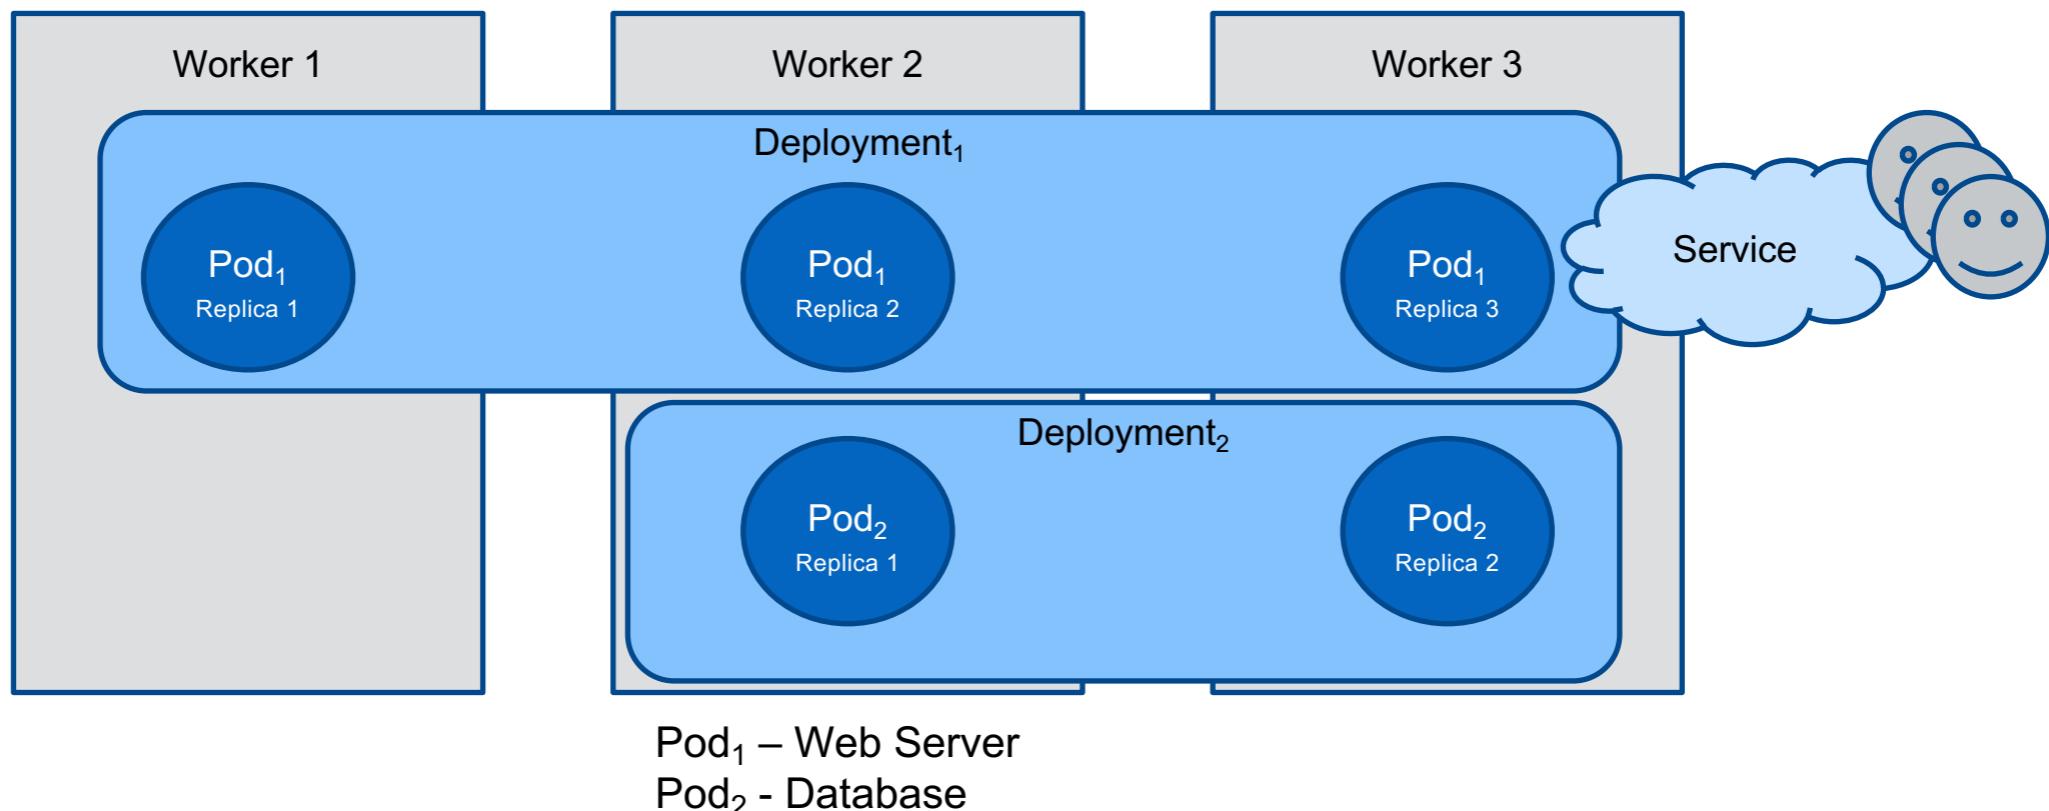
Kubernetes Cluster

- Clients interact with the Master node (e.g., through the *kubectl* command-line tool)
- The Master manages the kubernetes cluster (i.e., worker nodes, pods, network, storage, ...)



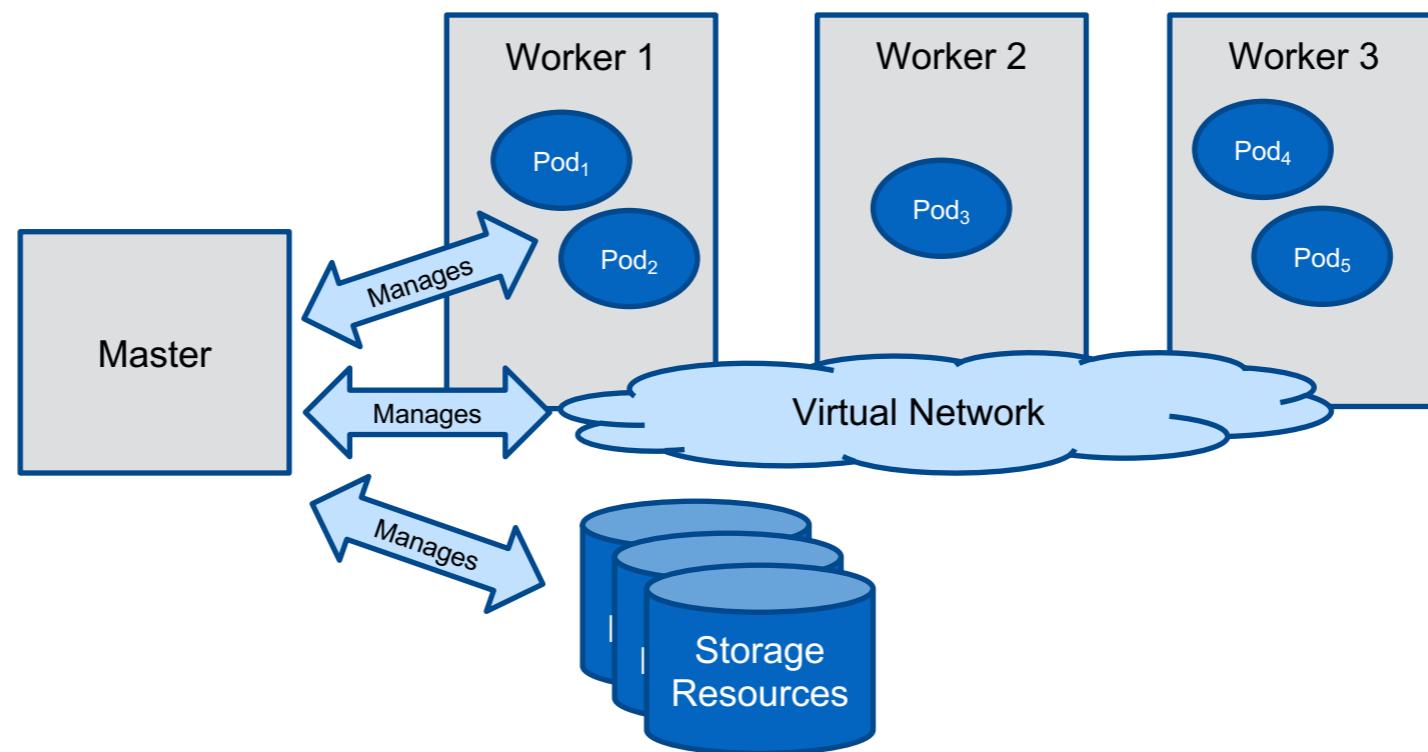
Pod, Deployment and Service

- **Pod** - computing unit composed by one or more application containers (e.g., Docker, containerd, CRI-O)
- **Deployment** – Specifies the runtime environment for pods (i.e., number of replicas)
- **Service** – Exposes a group of pods (application) as a network service



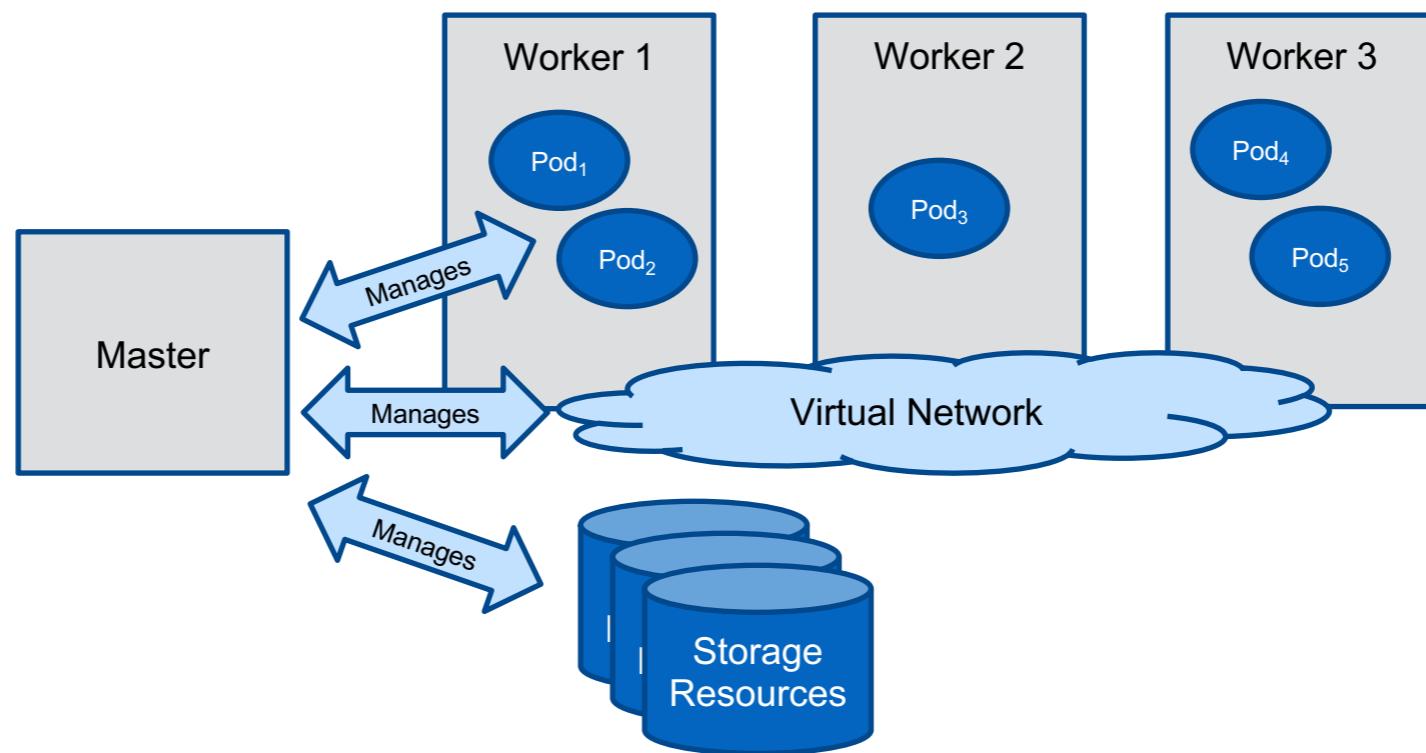
Network

- **Network** – each pod has a unique cluster-wide IP address
 - Networking across pods (even in distinct worker nodes) can be managed through networking overlays (e.g., Flannel, Calico)



Storage

- **Storage** – Pods can access storage volumes provided by different storage backends (e.g., iSCSI, NFS, AWS, GCP)
 - **Ephemeral** storage – available during the pod's lifetime
 - **Persistent** storage – available even if the pod is terminated



Summary

Containers vs VMs – Disclaimer!

Each technology is built with different goals and the best one depends on the targeted use case!

- VMs are useful when full server (OS) virtualization is needed
- Containers are useful for managing different libraries / applications

Advantages Containers over VMs

- Faster testing/provisioning/migration
- Better resource utilization (lightweight) and performance
- Can be deployed on both physical and virtualized servers

Disadvantages Containers over VMs

- Weaker isolation/security (OS/Kernel are shared)
- Less flexibility in running different OSs
(e.g., Linux, Windows, BSD, ...)

Further reading

- S. Alapati. **Modern Linux Administration: How to Become a Cutting-edge Linux Administrator.** O'Reilly, 2016
- K. Morris. ***Infrastructure as Code: Managing Servers in the Cloud.*** O'Reilly, 2016
- B. Burns. **Kubernetes Up & Running (Second Edition).** O'Reilly, 2019.

Questions?