University of Minho
School of Engineering

ISLab

Dados e Aprendizagem Automática

Intro to Data Science & Python/Scikit-learn

DAA @ MEI-1º/MiEI-4º – 1º Semestre

César Analide, Bruno Fernandes, Filipa Ferraz, Filipe Gonçalves, Victor Alves

Part I

Contents

Data Types

Data Types

- Major Types of Data:
  - Numerical
  - Categorical
  - Ordinal

Data Types

Numerical

- Represents some sort of quantitative measurement
  - Heights of people, page load times, stock prices, etc.

- Discrete Data
  - Integer based; often counts of some event.
    - How many purchases did a customer make in a year?
    - How many times did I flip “heads”?

- Continuous Data
  - Has an infinite number of possible values
    - How much time did it take for a user to check out?
    - How much rain fell on a given day?

Data Types

Categorical

- Qualitative data that has no inherent mathematical meaning
  - Gender, Yes/No (Binary Data), Race, State of Residence, Product Category, Political Party, etc.

- You can assign numbers to categories in order to represent them more compactly, but the numbers don't have mathematical meaning

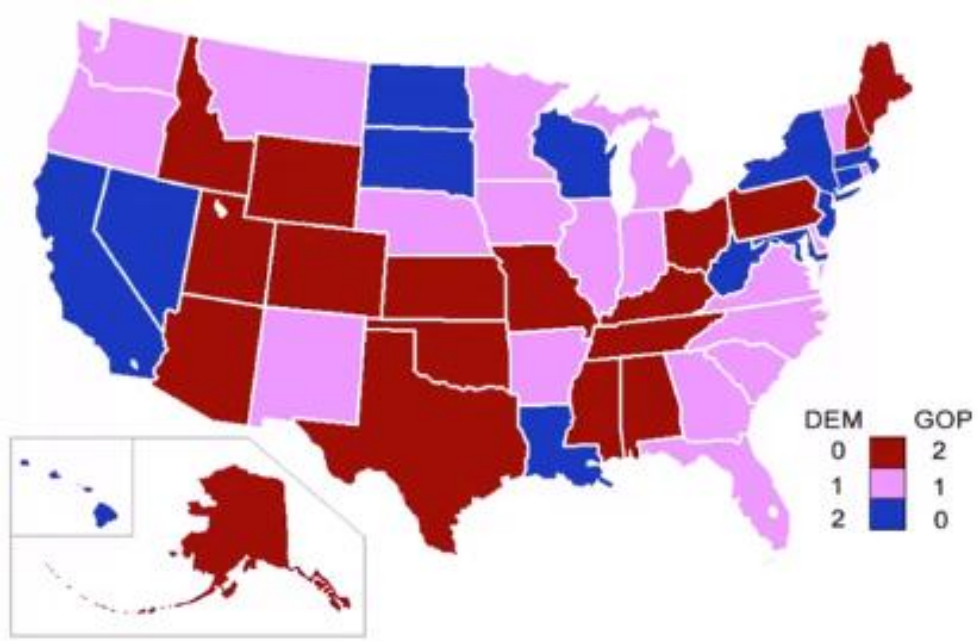

Data Types

Ordinal

- A mixture of numerical and categorical
- Categorical data that has mathematical meaning
- Example: movie ratings on a 1-5 scale.
  - Ratings must be 1,2,3,4 or 5
  - These values have mathematical meaning; 1 means it's a worse movie than a 2.

Data Types

Quick Quiz:

- Are the following types of data numerical, categorical, or ordinal?
  - How much gas is in your gas tank?
  - A rating of your overall health where the choices are 1,2,3 or 4, corresponding to “poor”, “moderate”, “good” and “excellent”
  - The nationalities of your classmates
  - Ages in years
  - Money spent in a store

Mean, Median & Mode

Mean, Median & Mode

Mean

- AKA Average
- Sum / number of samples
- Example:
  - Number of children in each house on my street:

0, 2, 3, 2, 1, 0, 0, 2, 0

The MEAN is (0+2+3+2+1+0+0+2+0) / 9 = **1.11**

Mean, Median & Mode

Median

- Sort the values, and take the value at the midpoint.
- Example:

0, 2, 3, 2, 1, 0, 0, 2, 0

Sort it:

0, 0, 0, 0, 1, 2, 2, 2, 3

- If you have an even number of samples, take the average of the two in the middle.

Mean, Median & Mode

Median

- Median is less susceptible to outliers than the mean
  - Example: mean household income in the USA is $72,641, but the median is only $51,939 – because the mean is skewed by a handful of billionaires.
  - Median represents better the "typical" American in this example.

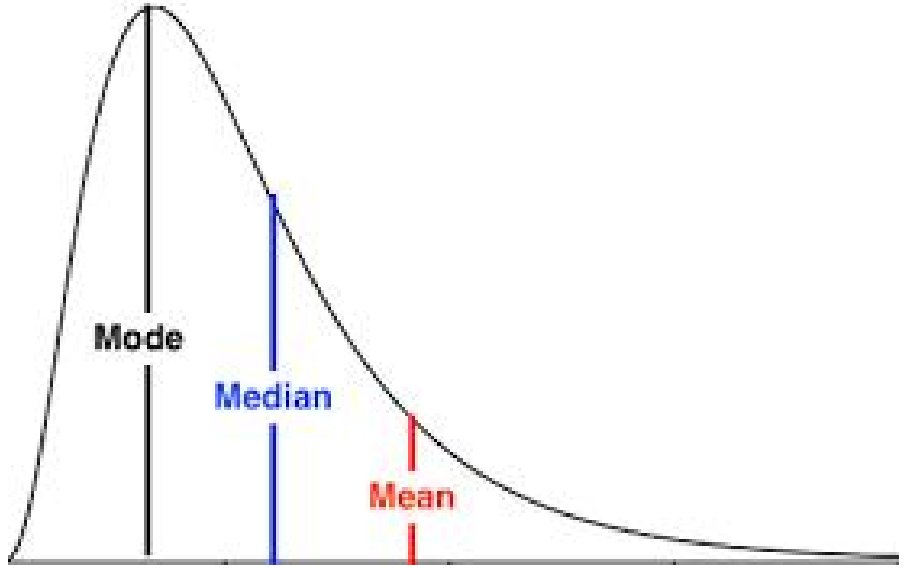

Mean, Median & Mode

Mode

- The most common value in a dataset
  - Not relevant to continuous numerical data
- Number of kids in each house example:

0, 2, 3, 2, 1, 0, 0, 2, 0

How many of each value are there?

0: 4, 1: 1, 2: 3, 3: 1

The MODE is 0

Standard Deviation & Variance

Standard Deviation & Variance

An example of a histogram…

Standard Deviation & Variance

Variance measures how "spread-out" the data is.

- Variance (δ^2) is simply the average of the squared differences from the mean

- Example: What is the variance of the data set (1, 4, 5, 4, 8)?
  - First find the mean: (1+4+5+4+8) / 5 = 4.4
  - Now find the difference from the mean: (-3.4, -0.4, 0.6, -0.4, 3.6)
  - Find the squared differences: (11.56, 0.16, 0.36, 0.16, 12.96)
  - Find the average of the squared differences:
  - δ^2 = (11.56+0.16+0.36+0.16+12.96) / 5 = 5.04

Standard Deviation & Variance

Standard Deviation δ is the square root of the variance.

- Case Study = (1,4,5,4,8)
- Mean = 4.4
- $\delta^2 = 5.04$
- $\delta = 2.24$

- Stand. Dev. Is usually used as a way to identify outliers.
- Data points that lie more than one standard deviation from the mean can be considered unusual.

- You can talk about how extreme a data point is by talking about "how many sigmas" away from the mean it is.

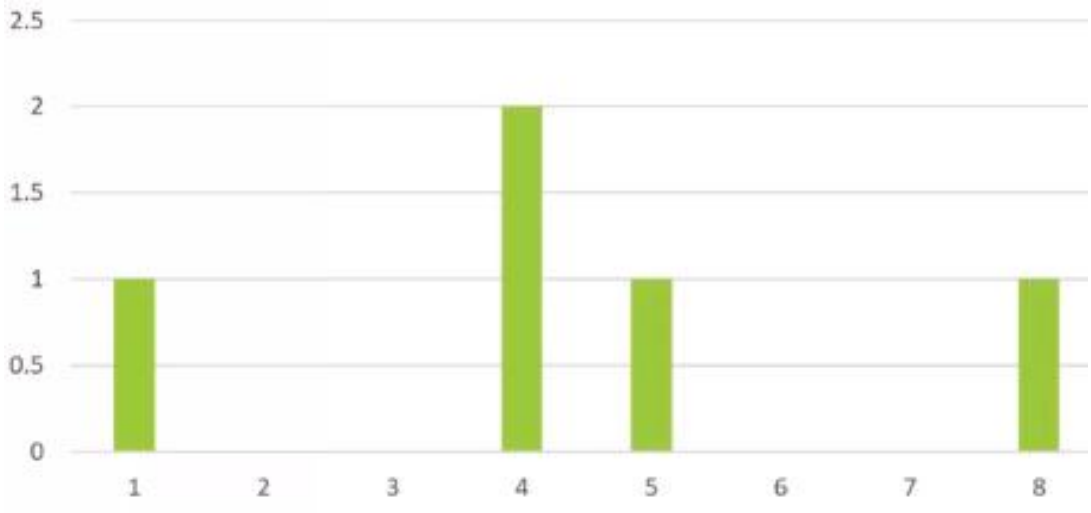

Probability Density Functions

Probability Density Functions

Example: a "normal distribution"

- Gives you the probability of a data point falling within some given range of a given value
- Based on histogram values, a normal probability density function can be calculated

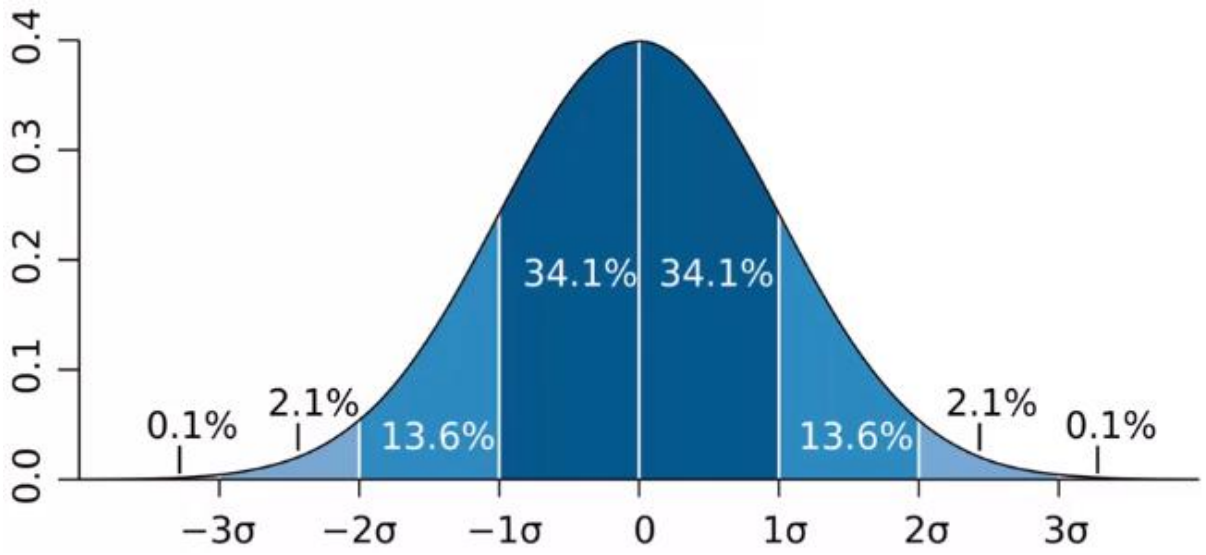

Probability Density Functions

Example: Probability Mass Function

- Used for discrete data
- Based on histogram values, a normal probability density function can be calculated

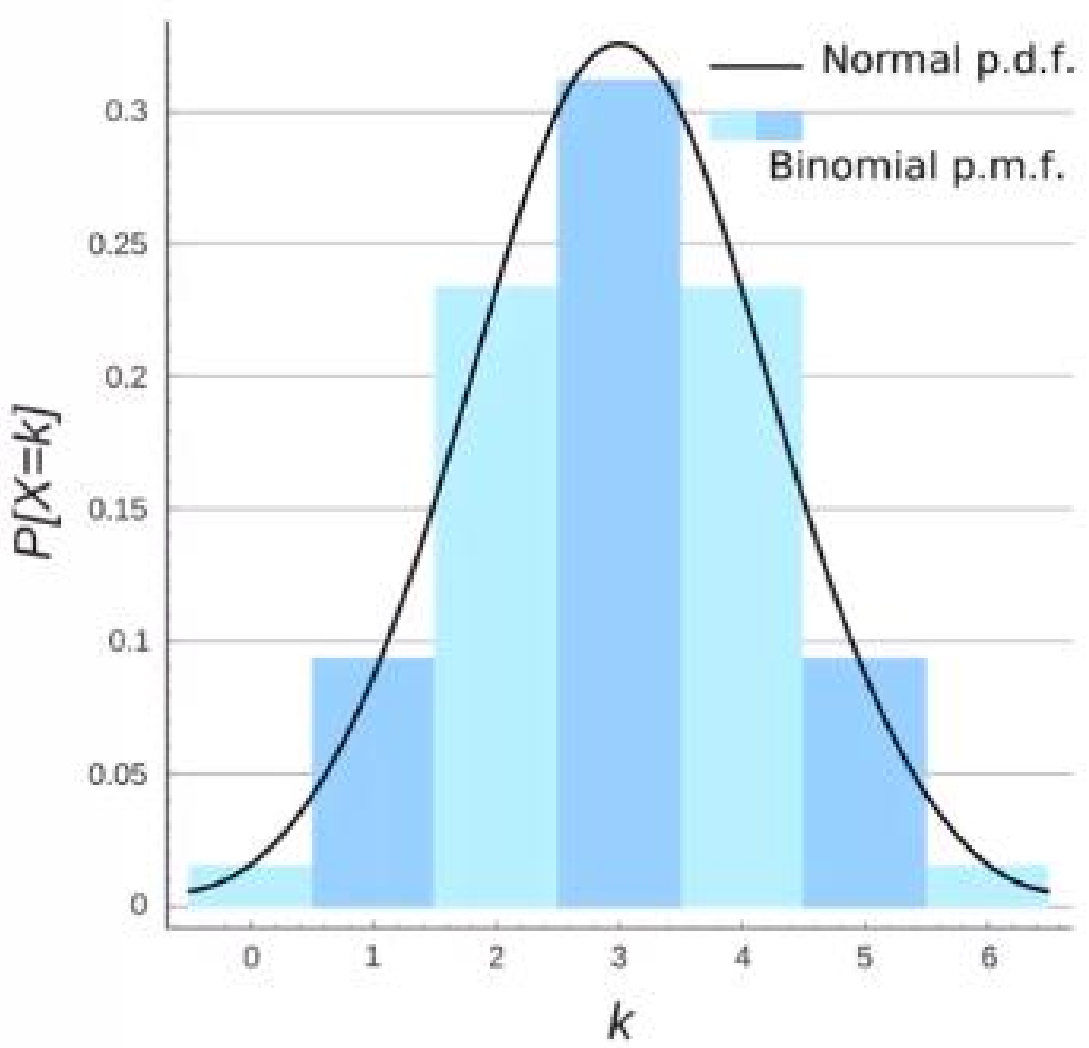

Percentiles

Percentiles

Percentiles

- In a dataset, what's the point at which X% of the values are less than that value?
- Example: income distribution
  - Take all incomes from a country's population and sort them
  - 99th percentile represents the income amount in which 99% of the population gains less then that value (i.e., $506,553)

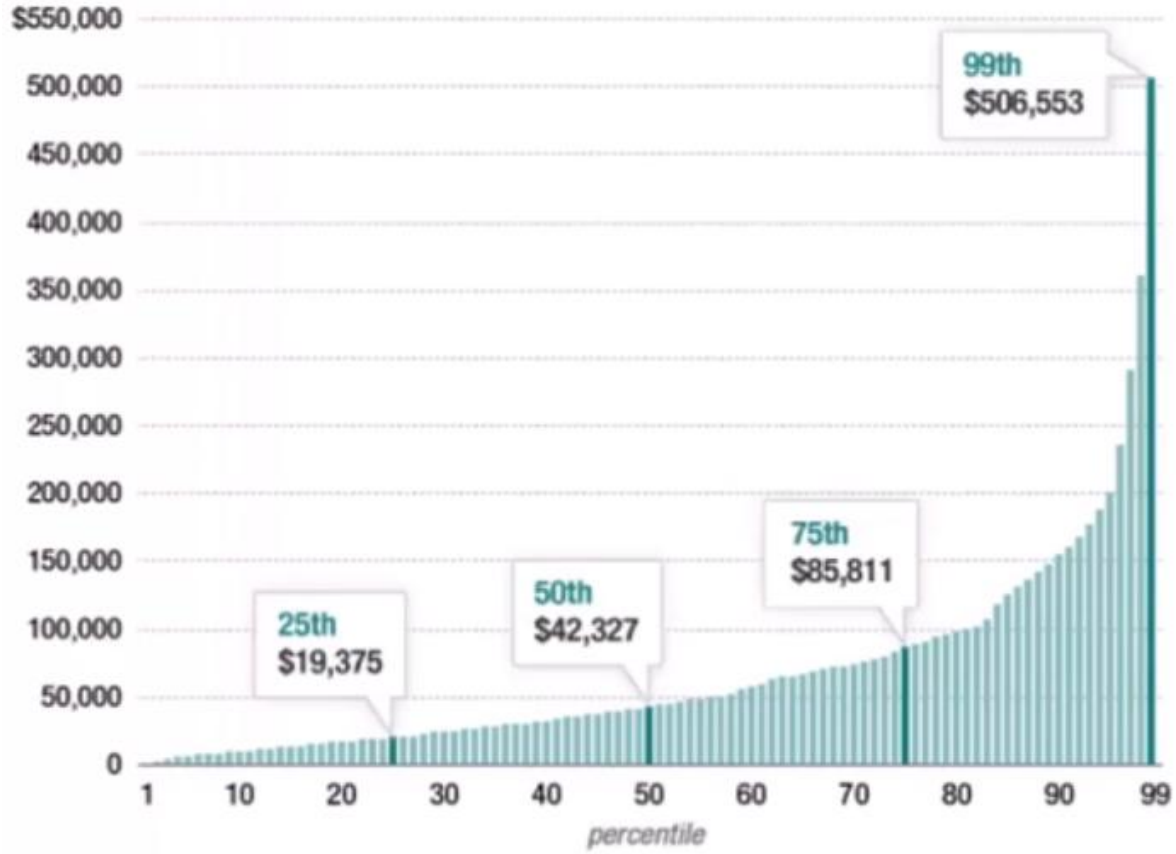

Percentiles

Percentiles in a normal distribution

- Between Quartil 1 & Quartil 3 represents 50% of the data distribution
- IQR (Inter-Quartil Range) represents the area in the middle of the distribution (where data is more focused)

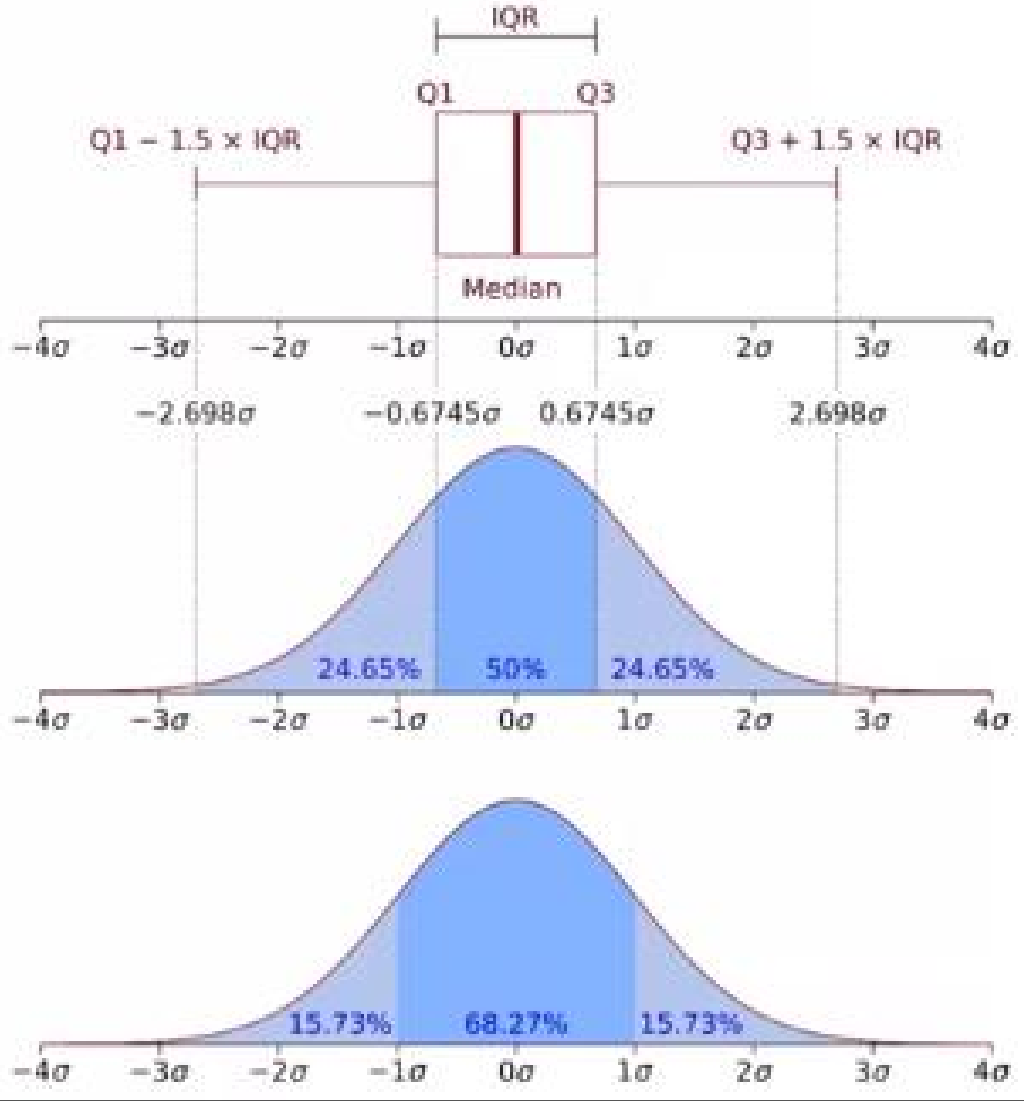

Covariance & Correlation

Covariance & Correlation

Covariance

- Measures how two variables vary in tandem from their means.
- i.e. how 2 attributes depend on each other (left plot – low covariance / right plot – high covariance)

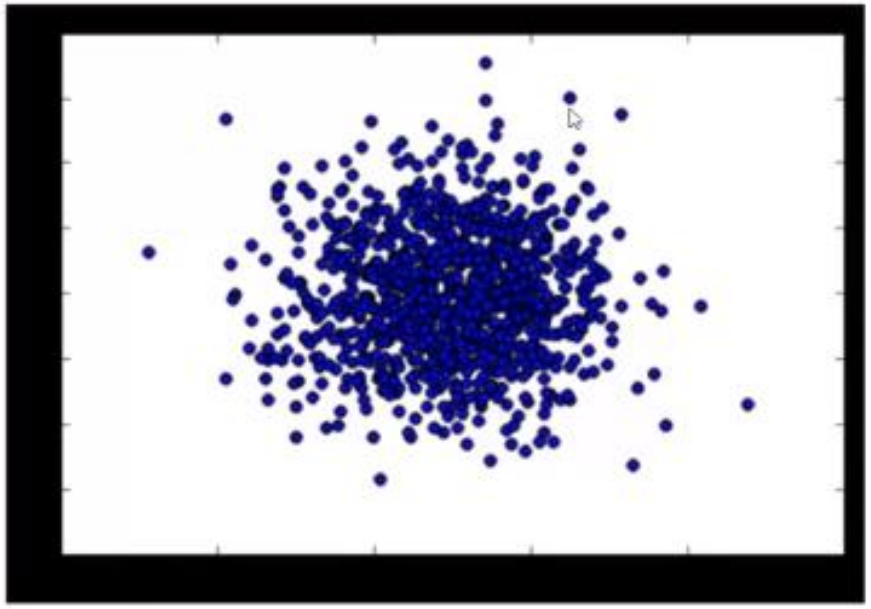

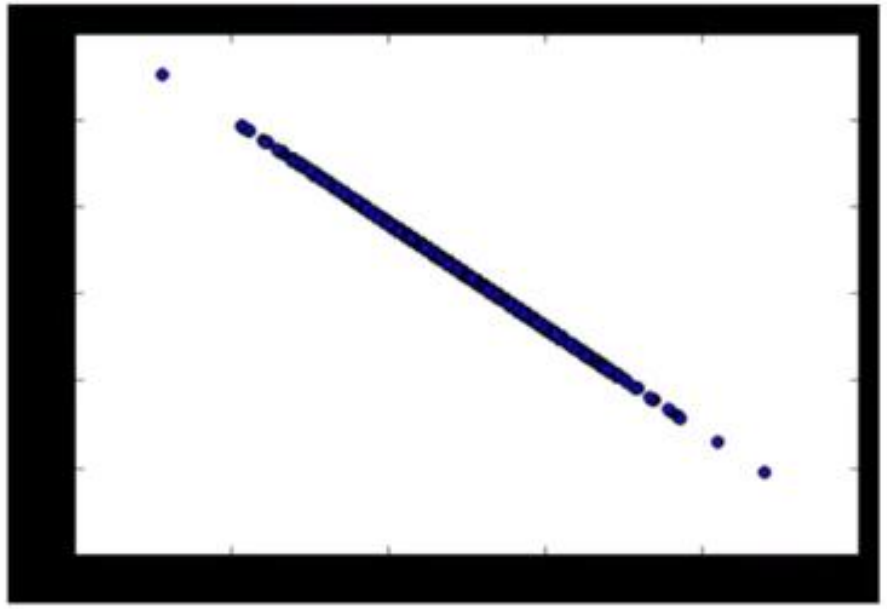

Covariance & Correlation

Measuring covariance

- Think of the datasets for the two variables as high-dimensional vectors
- Convert these to vectors of variances from the mean
- Take the dot product (cosine of the angle between them) of the two vectors
- Divide by the population size

Population Covariance Formula

$$Cov(x,y) = \frac{\Sigma(x_i-\bar{x})(y_i-\bar{y})}{N}$$

Sample Covariance

$$Cov(x,y) = \frac{\Sigma(x_i-\bar{x})(y_i-y)}{N-1}$$

Covariance & Correlation

Interpreting covariance is hard

- Small covariance (close to 0) means there isn't much correlation between the two variables
- Large covariance (far from 0 – can be negative for inverse relationships) means that there is a correlation

Interpreting correlation is easier

- Normalization value of covariance divided by the standard deviations of both variables
  - Correlation of -1: perfect inverse correlation
  - Correlation of 0: no correlation
  - Correlation of 1: perfect correlation

Covariance & Correlation

Correlation does not imply causation!

- Only a controlled, randomized experiment can give you insights on causation.
- Use correlation to decide what experiments to conduct!

Covariance & Correlation

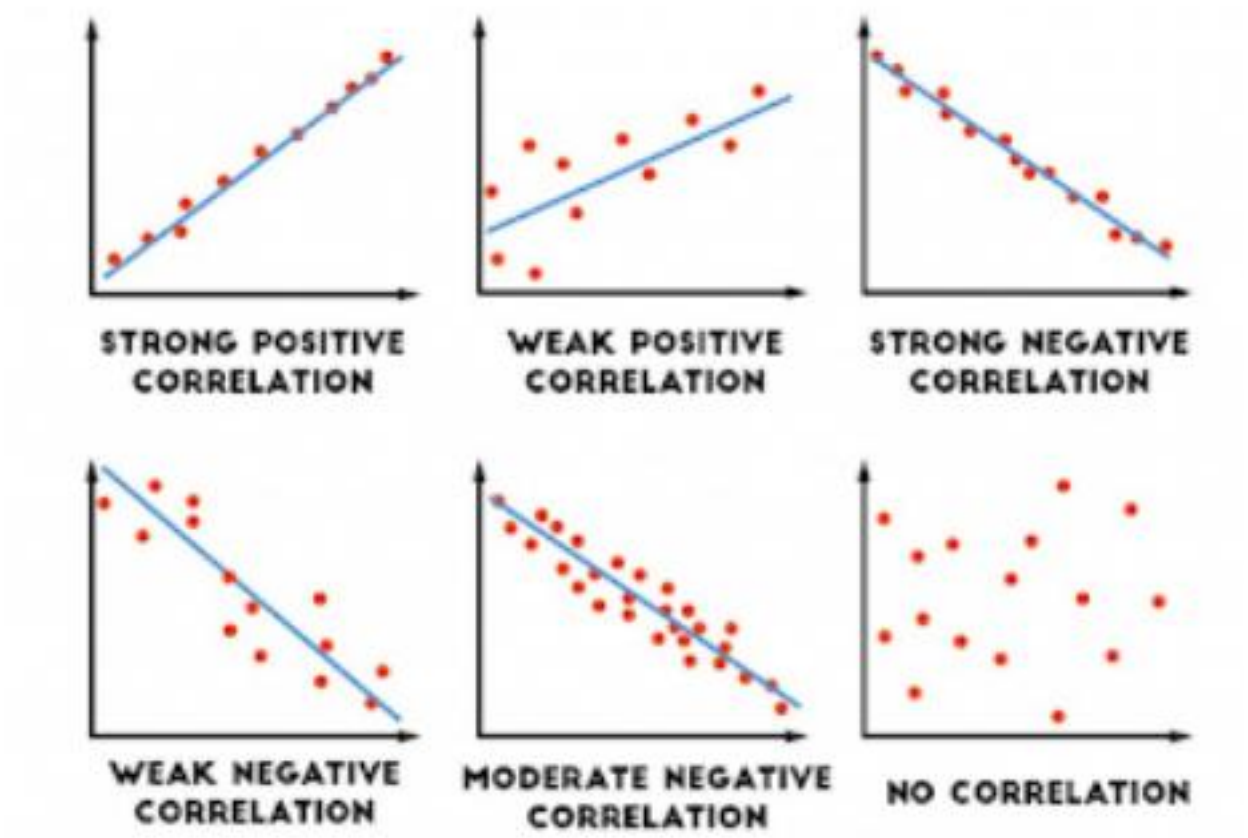

Virtual Environments

Virtual Environments

- Virtual Environments allow you to set up virtual installations of Python and libraries on your computer
- You can have multiple versions of Python or libraries and easily activate or deactivate these environments
- Let's see some examples of why you may want to do this

Virtual Environments

- Sometimes you'll want to program in different versions of a library
- For example:
  - You develop a program with SciKit-Learn 0.17
  - SciKit-Learn 0.18 is released
  - You want to explore 0.18 but don't want your old code to break

- Sometimes you'll want to make sure your library installations are in the correct location

- For example:
  - You want multiple versions of Python on your computer
  - You want one environment with Python 2.7 and another with Python 3.6

Virtual Environments

- Anaconda has a built-in virtual environment manager that makes the whole process really easy
- Check out the resource link for the official documentation:
  - http://conda.pydata.org/docs/using/envs.html

- Command Prompt Example (create env. and activate it):

```
conda create --name mypython3version python=3.6 numpy
conda info --evns
activate mypython3version
python
import numpy as np
import pandas as pd      -> Error
quit()
conda install pandas
deactivate
```

Anaconda Distribution

The World's Most Popular Python/R Data Science Platform

- FOSS
- Share, collaborate on, and reproduce projects
- Highly supported by the community
- Conda, a package, dependency and environment manager
  - Easily create, save, load and switch between environments
  - Easily install, update and run any package (and its dependencies… automatically!)

- Anaconda provides two user clients
  - Anaconda Navigator
  - Anaconda Prompt (or the terminal on Linux and macOS)

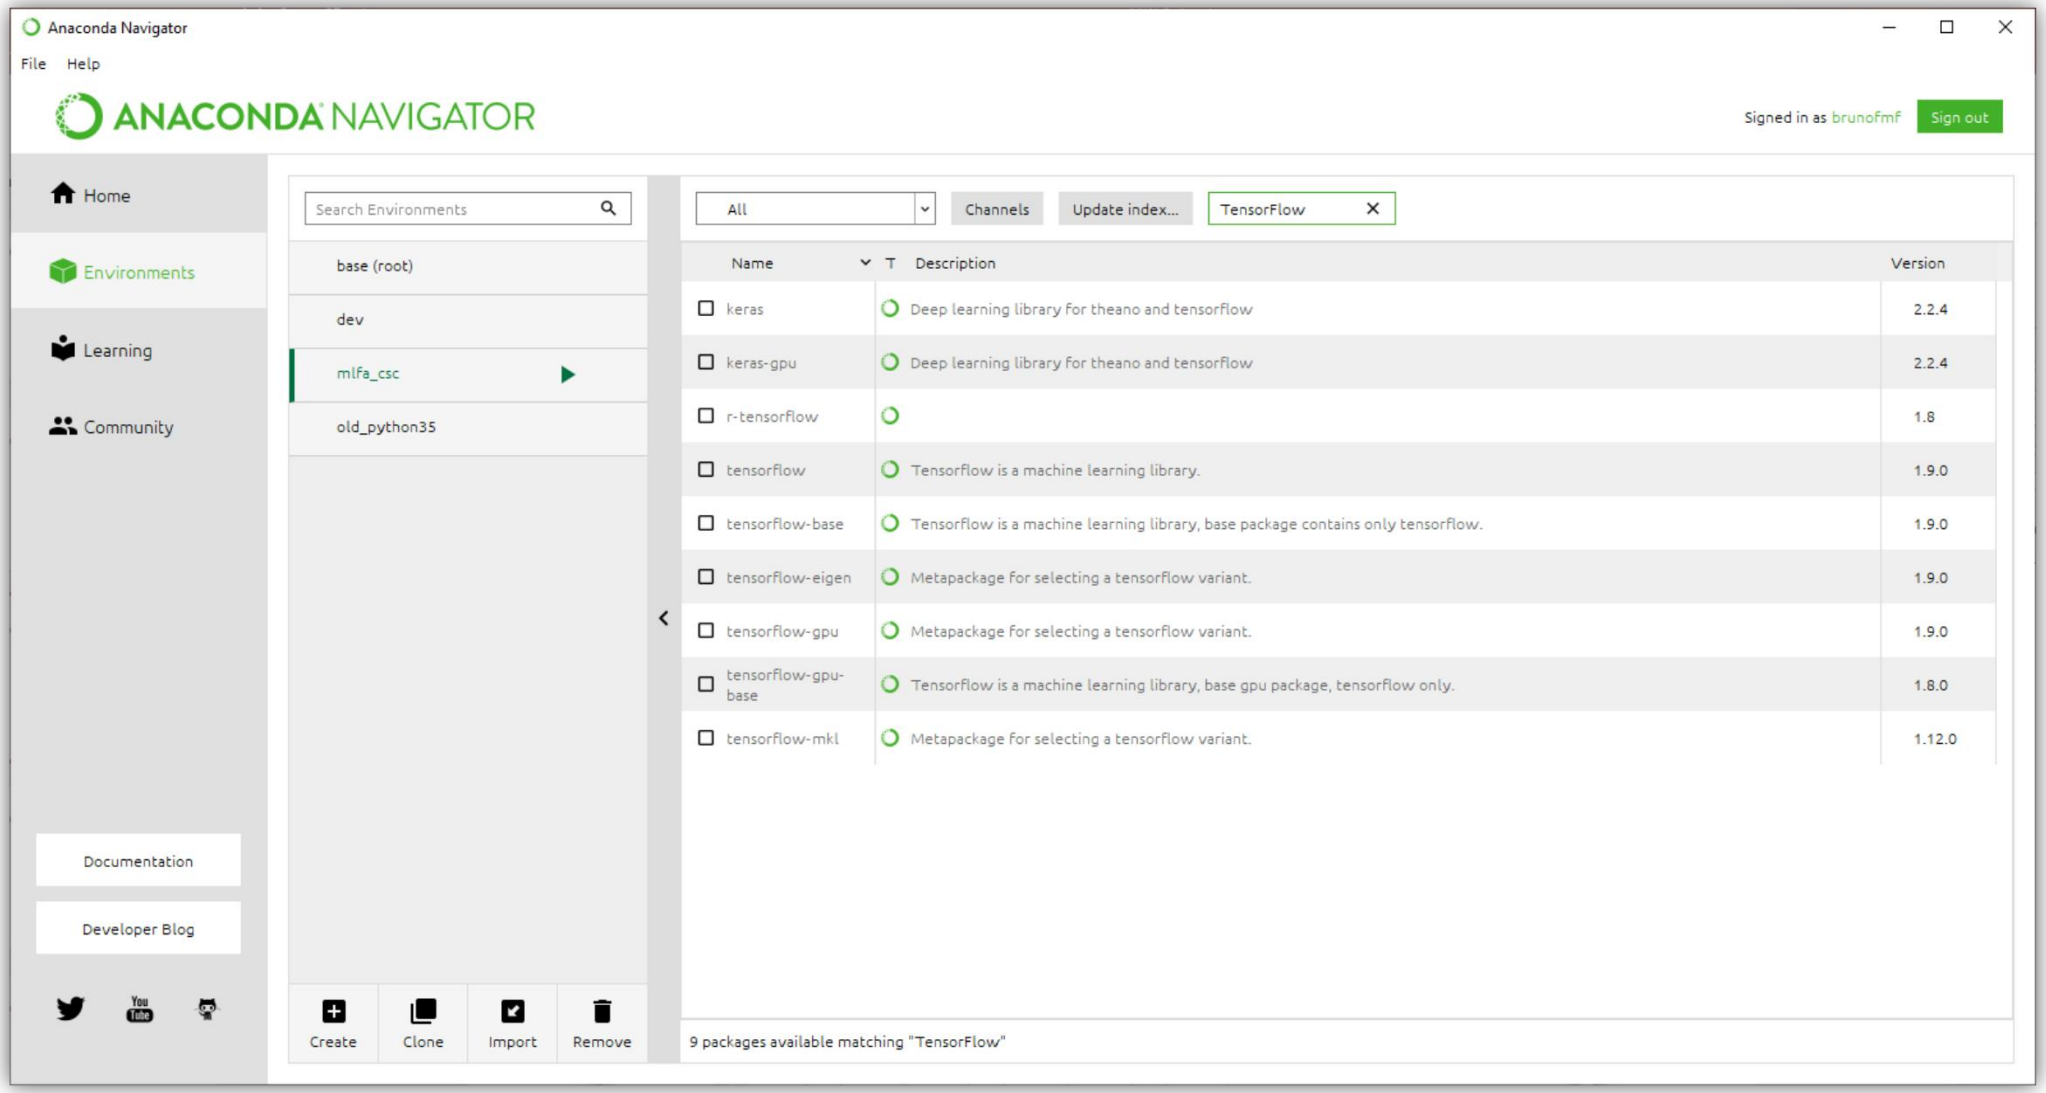
Anaconda Navigator
File Help
ANACONDA NAVIGATOR
Signed in as brunofmf
Sign out
Home
Environments
Learning
Community
Documentation
Developer Blog
Search Environments
base (root)
dev
mlfa_csc
old_python35
Create
Clone
Import
Remove
All
Channels
Update index...
TensorFlow
Name
T
Description
Version
keras
Deep learning library for theano and tensorflow
2.2.4
keras-gpu
Deep learning library for theano and tensorflow
2.2.4
r-tensorflow
1.8
tensorflow
Tensorflow is a machine learning library.
1.9.0
tensorflow-base
Tensorflow is a machine learning library, base package contains only tensorflow.
1.9.0
tensorflow-eigen
Metapackage for selecting a tensorflow variant.
1.9.0
tensorflow-gpu
Metapackage for selecting a tensorflow variant.
1.9.0
tensorflow-gpu-base
Tensorflow is a machine learning library, base gpu package, tensorflow only.
1.8.0
tensorflow-mkl
Metapackage for selecting a tensorflow variant.
1.12.0
9 packages available matching "TensorFlow"

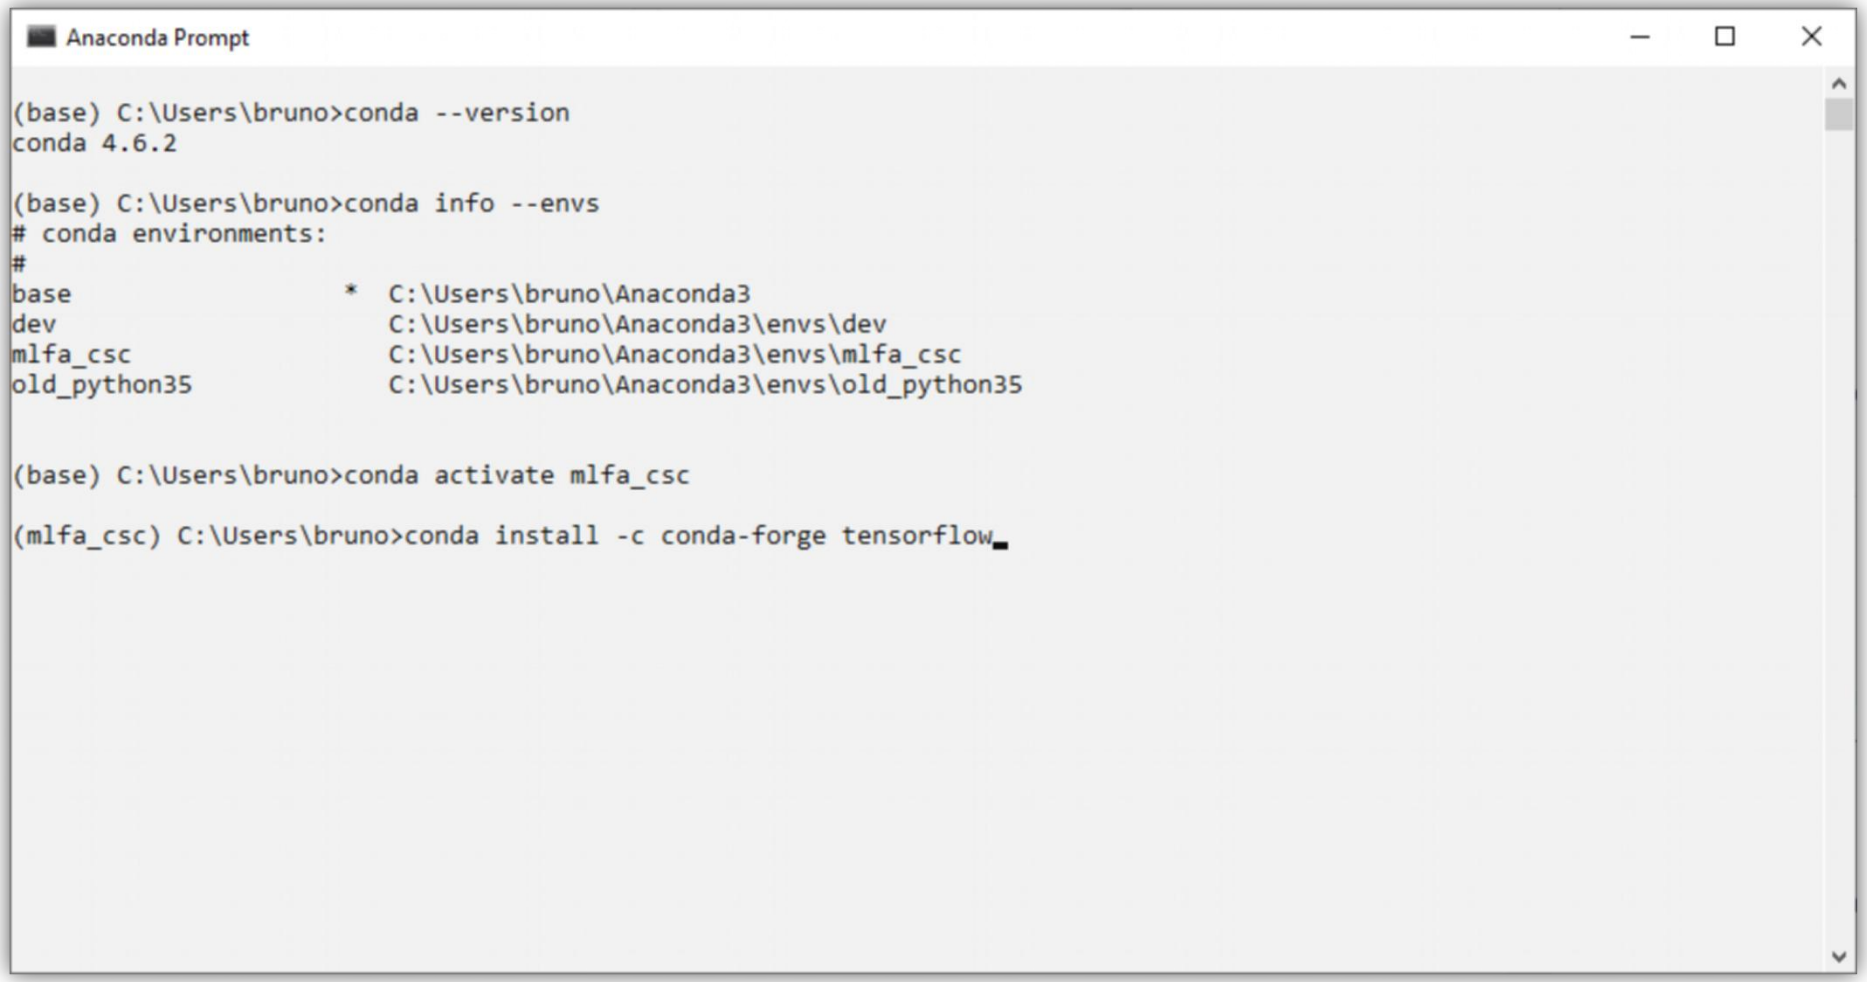
Anaconda Prompt
(base) C:\Users\bruno>conda --version
conda 4.6.2
(base) C:\Users\bruno>conda info --envs
conda environments:
#
base * C:\Users\bruno\Anaconda3
dev C:\Users\bruno\Anaconda3\envs\dev
mlfa_csc C:\Users\bruno\Anaconda3\envs\mlfa_csc
old_python35 C:\Users\bruno\Anaconda3\envs\old_python35
(base) C:\Users\bruno>conda activate mlfa_csc
(mlfa_csc) C:\Users\bruno>conda install -c conda-forge tensorflow

Conda basics

Verify conda is installed, check version number	`conda info`
Update conda to the current version	`conda update conda`
Install a package included in Anaconda	`conda install PACKAGENAME`
Run a package after install, example Spyder*	`spyder`
Update any installed program	`conda update PACKAGENAME`
Command line help	`COMMANDNAME --help` `conda install --help`

Using environments

Create a new environment named py35, install Python 3.5	`conda create --name py35 python=3.5`
Activate the new environment to use it	`WINDOWS: activate py35` `LINUX, macOS: source activate py35`
Get a list of all my environments, active environment is shown with *	`conda env list`
Make exact copy of an environment	`conda create --clone py35 --name py35-2`
List all packages and versions installed in active environment	`conda list`
List the history of each change to the current environment	`conda list --revisions`
Restore environment to a previous revision	`conda install --revision 2`
Save environment to a text file	`conda list --explicit > bio-env.txt`
Delete an environment and everything in it	`conda env remove --name bio-env`
Deactivate the current environment	`WINDOWS: deactivate` `macOS, LINUX: source deactivate`
Create environment from a text file	`conda env create --file bio-env.txt`
Stack commands: create a new environment, name it bio-env and install the biopython package	`conda create --name bio-env biopython`

Installing and updating packages

Install a new package (Jupyter Notebook) in the active environment	`conda install jupyter`
Run an installed package (Jupyter Notebook)	`jupyter-notebook`
Install a new package (toolz) in a different environment (bio-env)	`conda install --name bio-env toolz`
Update a package in the current environment	`conda update scikit-learn`
Install a package (boltons) from a specific channel (conda-forge)	`conda install --channel conda-forge boltons`
Install a package directly from PyPI into the current active environment using pip	`pip install boltons`
Remove one or more packages (toolz, boltons) from a specific environment (bio-env)	`conda remove --name bio-env toolz boltons`

Managing multiple versions of Python

Install different version of Python in a new environment named py34	`conda create --name py34 python=3.4`
Switch to the new environment that has a different version of Python	`Windows: activate py34` `Linux, macOS: source activate py34`
Show the locations of all versions of Python that are currently in the path **NOTE**: The first version of Python in the list will be executed.	`Windows: where python` `Linux, macOS: which -a python`
Show version information for the current active Python	`python --version`

Environment Setup

Environment Setup

- This course will use Jupyter Notebooks/spyder for teaching and to provide notes
  - **Note**: you are free to use **whatever development environment you prefer (e.g., Spyder, PyCharm, ..)**

- We will be using the Python 3.6 for this course through the Anaconda Distribution
- Now let's go over your installation options for Jupyter Notebook!

Environment Setup

- For experienced users who already have Python
  - As an existing Python user, you may wish to install Jupyter and required APIs using Python package manager pip, instead of Anaconda
  - Just go to your command prompt or terminal and use:

pip install jupyter

- For new users, we highly recommend installing Anaconda
  - Anaconda conveniently installs Python, the Jupyter Notebook, and other commonly used packages for scientific computing and data science
  - Let's go to www.jupyter.org to walkthrough the installation steps!

Environment Setup

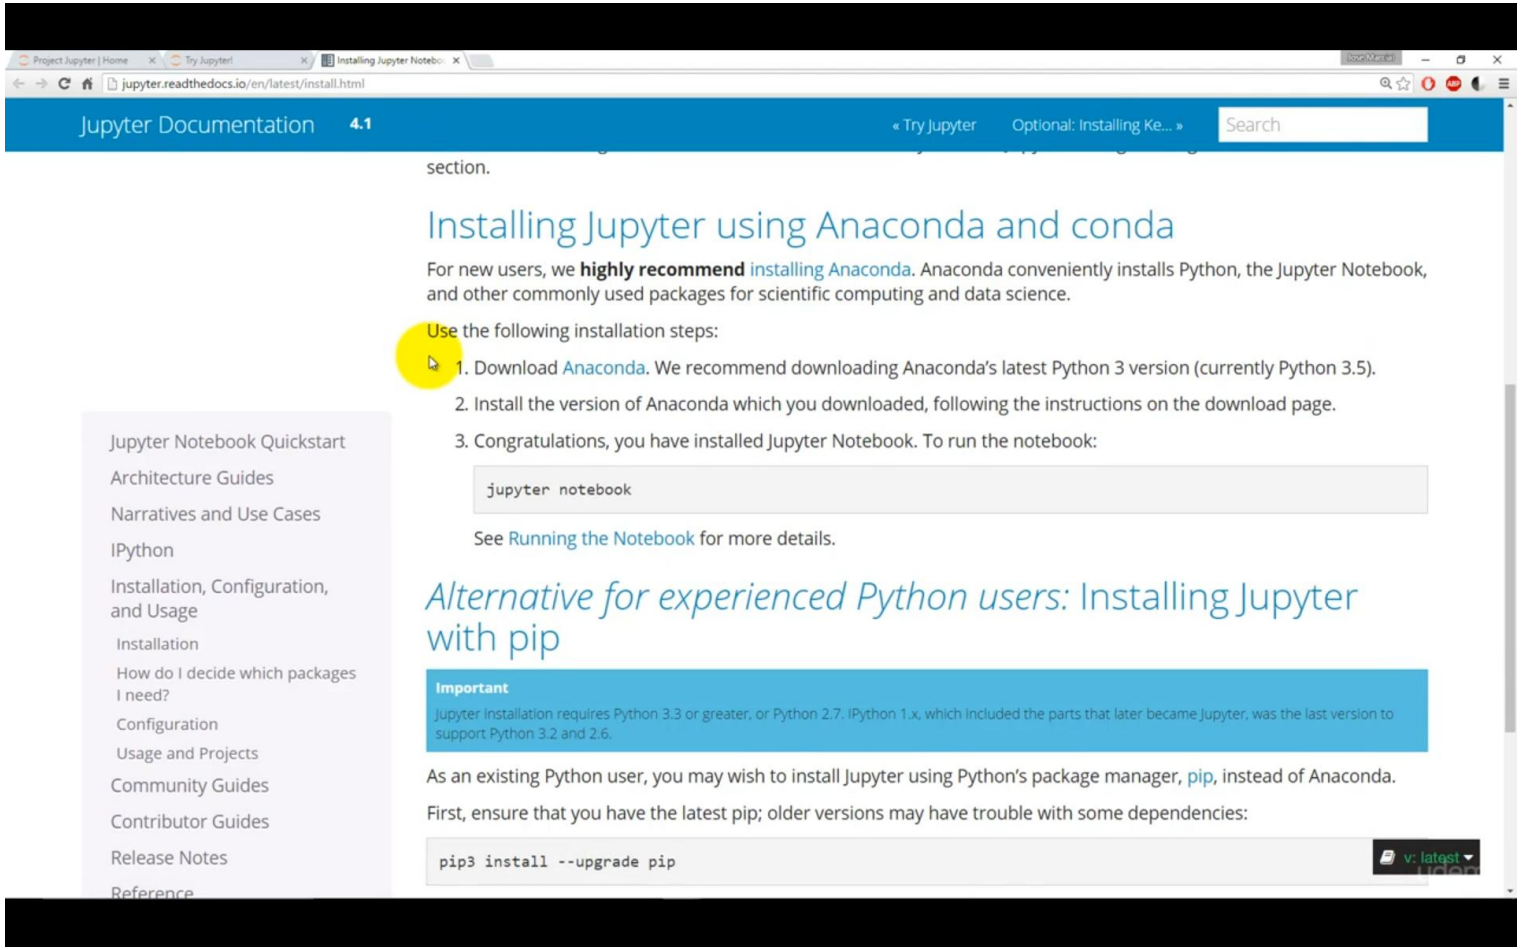

Environment Setup

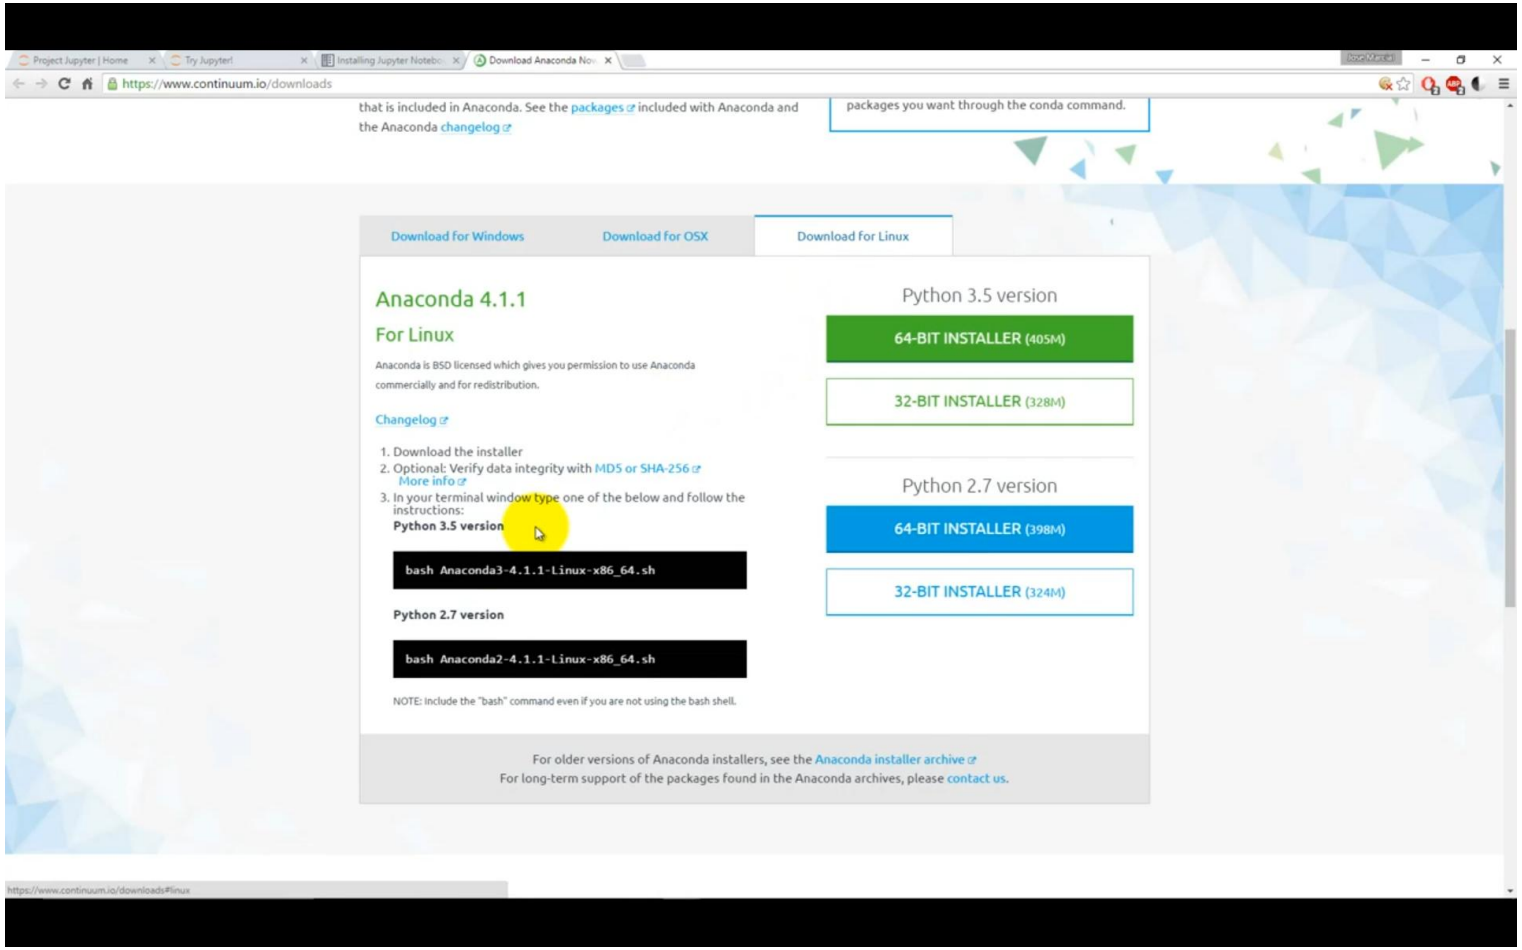

Environment Setup

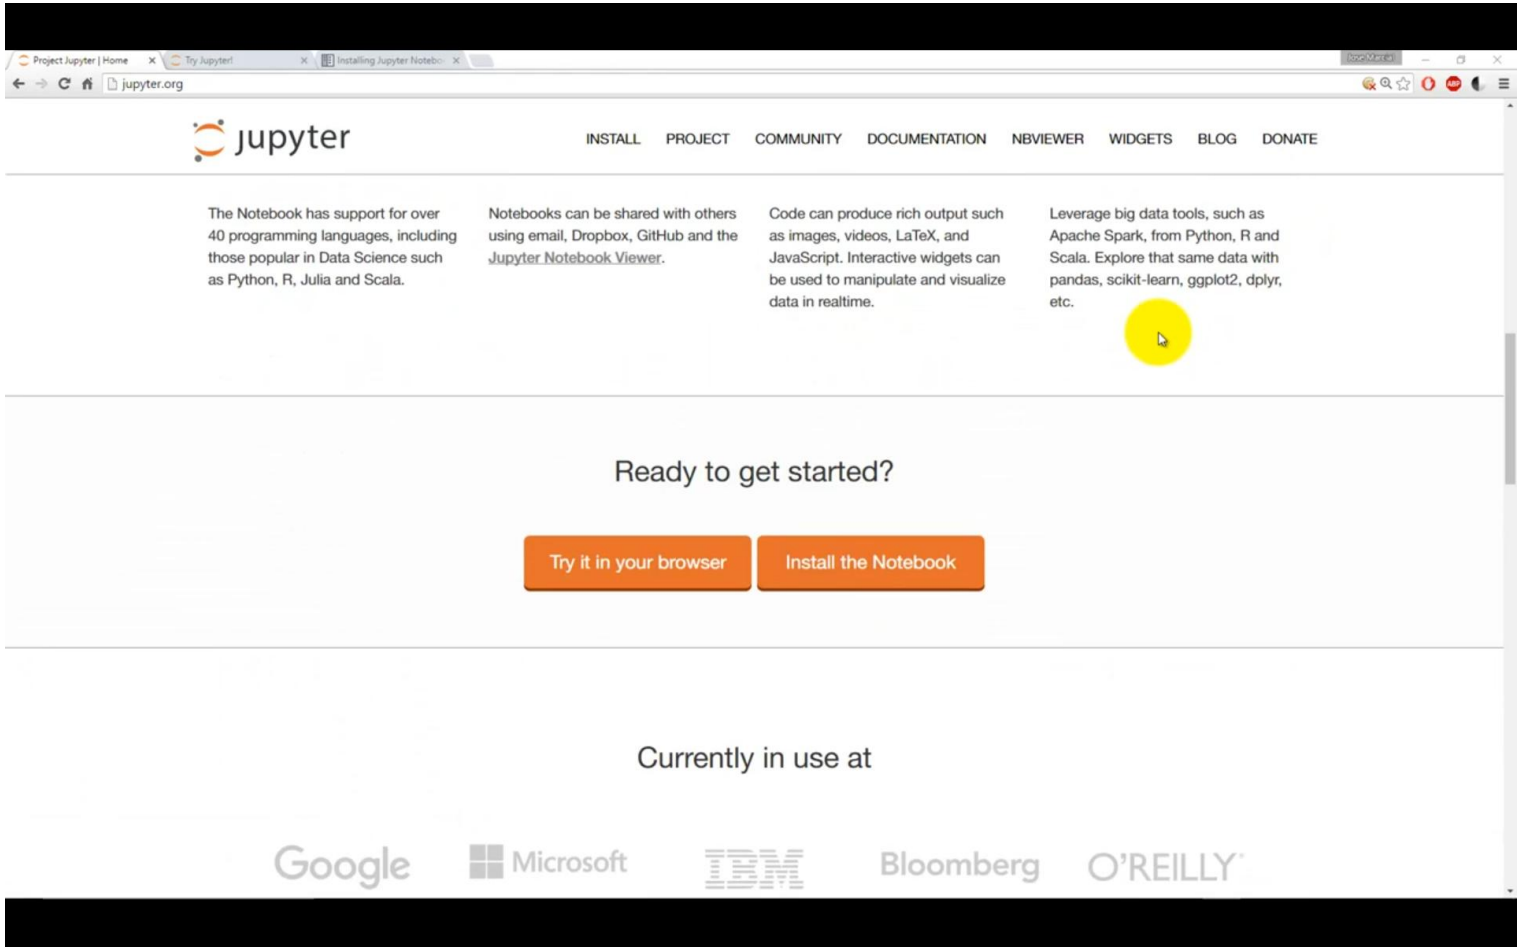

Environment Setup

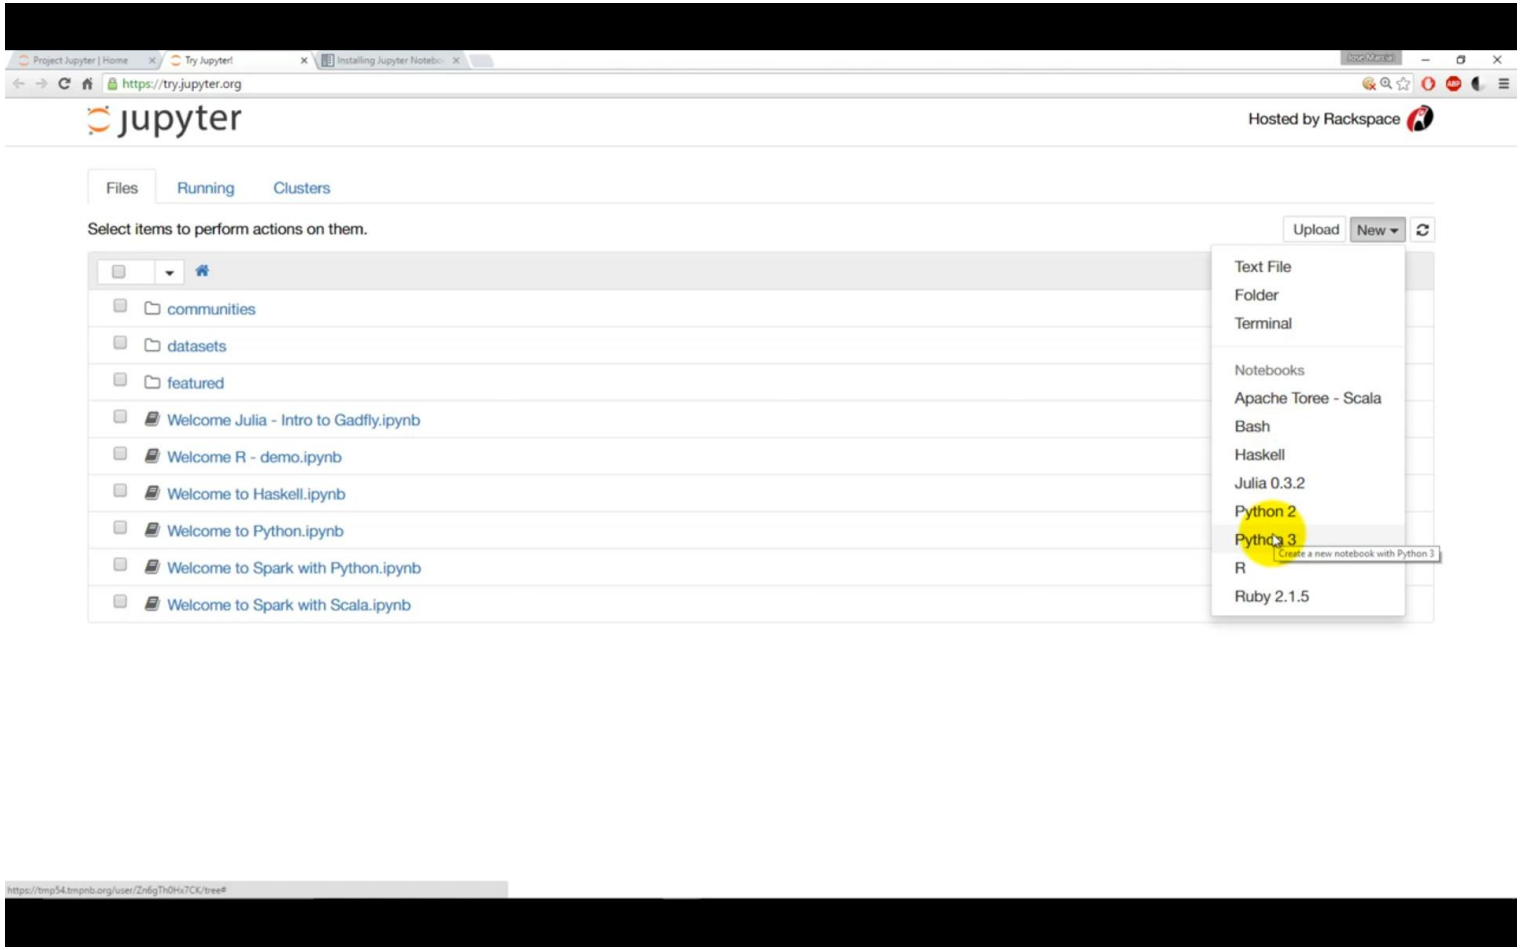

Environment Setup

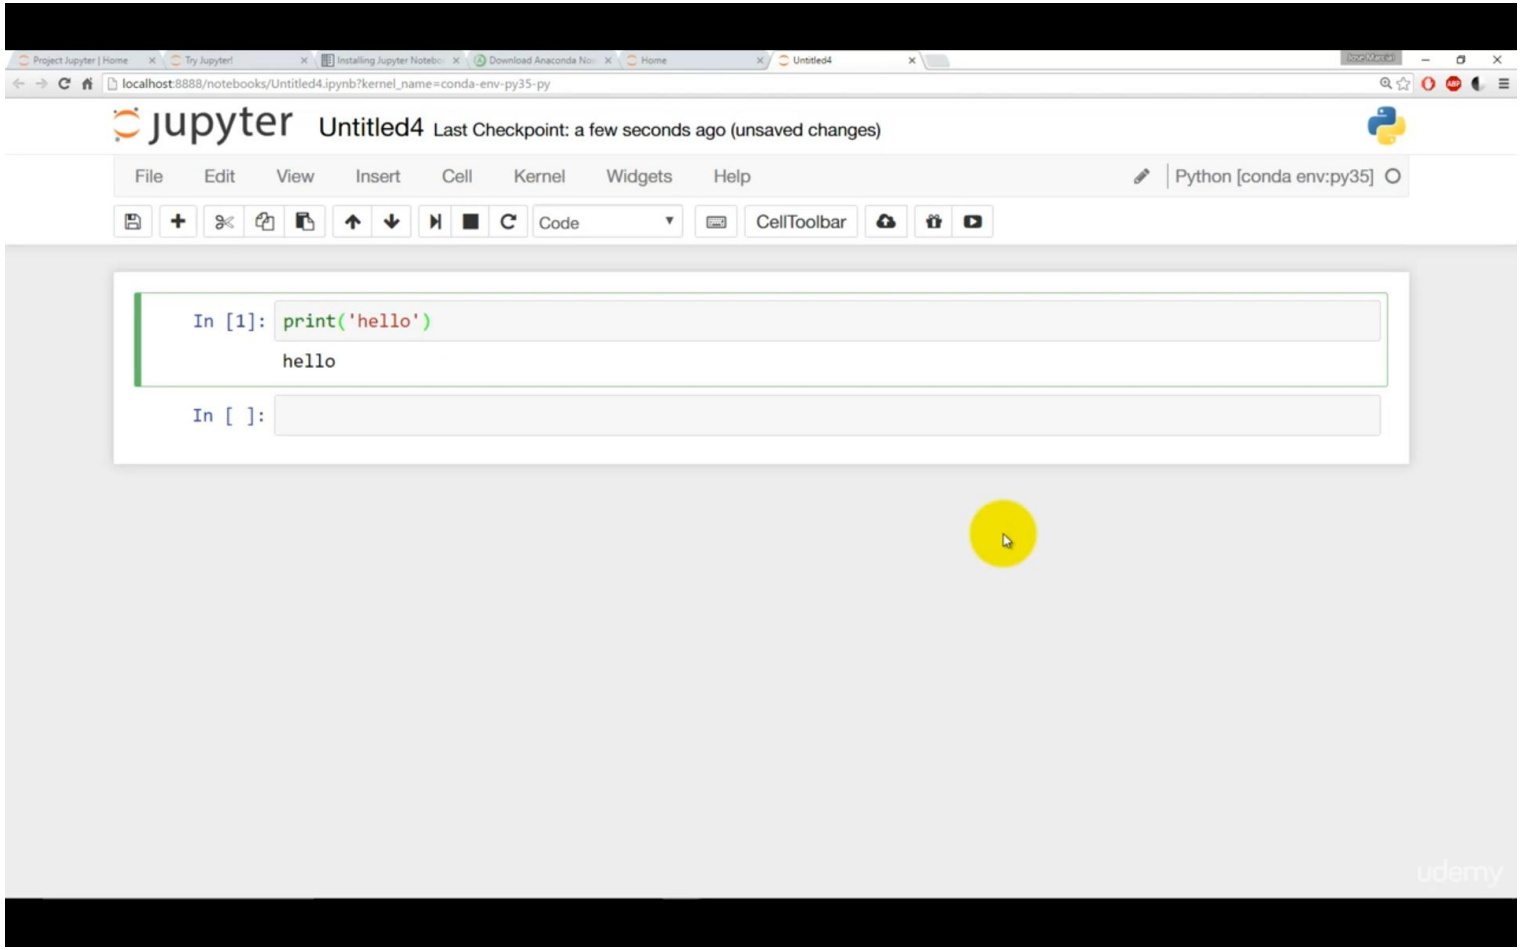

Hands On

Hands On

T1

- We will use scikit-learn/sklearn (Anaconda – Python Management Environment). Download and install the Anaconda Python package for your respective platform (Windows, Mac OS, Linux). The platform is available at https://www.anaconda.com/
  - Anaconda – Python 3.6
  - Deep Learning Libraries **not** required (Theano, Tensorflow, Keras)
  - Required to install Python IDE
  - https://machinelearningmastery.com/setup-python-environment-machine-learning-deep-learning-anaconda/

Hands On

T2

- Start Anaconda and create a virtual Python3.6 environment:
  - Open Terminal & Execute:
    - *conda create --name env python==3.6 numpy pandas xlrd xlwt matplotlib seaborn scikit-learn jupyterlab*
  - To install packages, enter the env. and execute: *conda install PACKAGENAME*
  - To work inside the python environment, execute: *conda activate env*
  - To exit python environment, execute: *conda deactivate*

Hands On

T2

- In this environment, the following libraries must be installed:
    a. Numpy
    b. Pandas
    c. Xlrd
    d. Xlwt
    e. Matplotlib
    f. Seaborn
    g. Scikit-learn
    h. Jupyterlab

Hands On

T3

- Activate the created virtual environment and check the installed libraries; validate the installation of the set of libraries presented in T2

T4

- Briefly check the documentation for each library mentioned in question T2. Identify its relevance in the context of Machine Learning algorithm development.