



Universidade do Minho
Escola de Engenharia
Departamento de Informática

Eng. Serviços de Rede (ESR)

Transport Protocols Recent developments

Stream Control Transmission Protocol (SCTP)

[RFC4960,2007]

Internet transport protocols

SCTP motivation



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- Why is SCTP needed given TCP and UDP are widely available?
- In 1998, an IETF working group (SIGTRAN) was formed to design a mechanism for reliably transporting call control signaling over the Internet.
- During SIGTRAN's work, **two key problems** surfaced in the use of TCP:
 - **head-of-line blocking** and **lack of multihoming** caused unacceptable delays in call signalling

Internet transport protocols

SCTP motivation



Universidade do Minho
Escola de Engenharia
Departamento de Informática

■ Reliability but no ordering

“TCP provides both reliable data transfer and strict order of transmission delivery of data. Some applications need reliable transfer without sequence maintenance, while others would be satisfied with partial ordering of the data. In both of these cases, the head-of-line blocking offered by TCP causes unnecessary delay.”

■ Reliability but no stream-oriented

“The stream-oriented nature of TCP is often an inconvenience. Applications must add their own record marking to delineate their messages, and must make explicit use of the push facility to ensure that a complete message is transferred in a reasonable time.”

Internet transport protocols

SCTP motivation



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- **More scalability and redundancy (multihoming)**

“The limited scope of TCP sockets complicates the task of providing highly-available data transfer capability using multi-homed hosts.”

- **More robustness (to attacks)**

“TCP is relatively vulnerable to denial-of-service attacks, such as SYN flood attacks.”

“When compared to TCP, SCTP aims at enhancing transport performance and robustness”


Internet transport protocols

SCTP introduction



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- **SCTP Stream Control Transmission Protocol**
 - Proposed IETF standard (RFC 4960, 2007)
 - Like TCP, it provides reliable, full-duplex, unicast transfers
 - Unlike TCP and UDP, it offers **new delivery options** that are particularly desirable **for telephony signaling** and **multimedia applications**



Internet transport protocols

SCTP goals



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- SCTP keeps TCP-like reliable retransmissions, congestion control, connection oriented, plus:
 - **framing** - preserve message boundaries
 - **4-way handshake** - to reduce vulnerability to DoS attacks
 - **multistreaming** - up to 64k independent ordered streams
 - **multihoming** - instead of one IP address per endpoint a set of IP addresses per endpoint
 - SCTP uses **multihoming for redundancy, not for load balancing!**

Internet transport protocols

SCTP association



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- The combination of an SCTP port and an IP address defines an “SCTP Transport Address”
- An SCTP **endpoint**
 - an endpoint is the logical end of the SCTP transport protocol - a communicating party
 - may have MORE than one IP address but it always has one and only one port number.

endpoint = [10.1.4.2, 10.1.5.3 : 80]

- In SCTP, the communication relationship between two endpoints is called an **association**

association = { [10.1.61.11: 2223], [10.1.4.2, 10.1.5.3 : 80]}

Internet transport protocols

SCTP framing



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- SCTP **maintains** message boundaries; an app message is maintained as **one or more data chunks**

*The objective is to increase performance by removing **blocking** in the received data, either due to delay or loss*

Internet transport protocols

SCTP framing



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- Decoupling of reliable and ordered delivery
 - allowing unordered delivery eliminates **head-of-line blocking** delay

TCP receiver buffer



App must wait!




Internet transport protocols

SCTP framing



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- Applications may want logical data units (“chunks”)
- Viewing data as a byte stream is inefficient



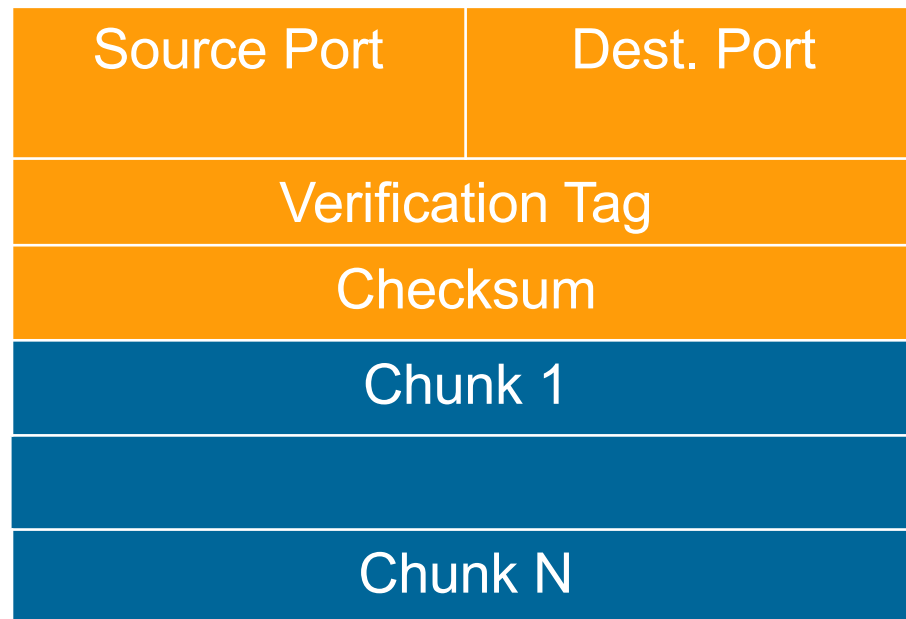
- If Packet 2 is lost, the other packets can be handled
- SCTP preserves Application Level Framing
 - each send/read is a “chunk” (App. Data Unit)

Internet transport protocols

SCTP framing



Universidade do Minho
Escola de Engenharia
Departamento de Informática



- SCTP PDU (max. 64 kbytes)
- SCTP may fragment user data according to minimum MTU

Examples of SCTP chunks:



0x00	DATA	User data
0x01	INIT	~ SYN
0x02	INIT-ACK	
0x03	SACK	Selective ACK
0x04	HEARTBEAT	Keep-alive message
0x05	HEARTBEAT-ACK	
0x07	SHUTDOWN	~FIN
0x08	SHUTDOWN-ACK	

Internet transport protocols

SCTP packets



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- **Common header**; three basic functions:
 - Source and destination ports, used together with the IP addresses and other state information to identify the **association** (connection)
 - Verification tag, random and negotiated at init
 - Checksum: CRC32 over the full SCTP packet
- followed by **one or more chunks**
 - chunks are concatenated building blocks containing either control or data information
 - a chunk header identifies length, type, and any special flags
 - **control chunks** transfer information needed for association functionality and **data chunks** carry application layer data.
 - current RFC specifies: 14 different **control chunks** for the association establishment, termination, ACK, destination failure recovery, ECN, and error reporting
- report to RFC4960 for SCTP PDU details


Internet transport protocols

SCTP 4-way handshake



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- SCTP improves protection against DoS attacks



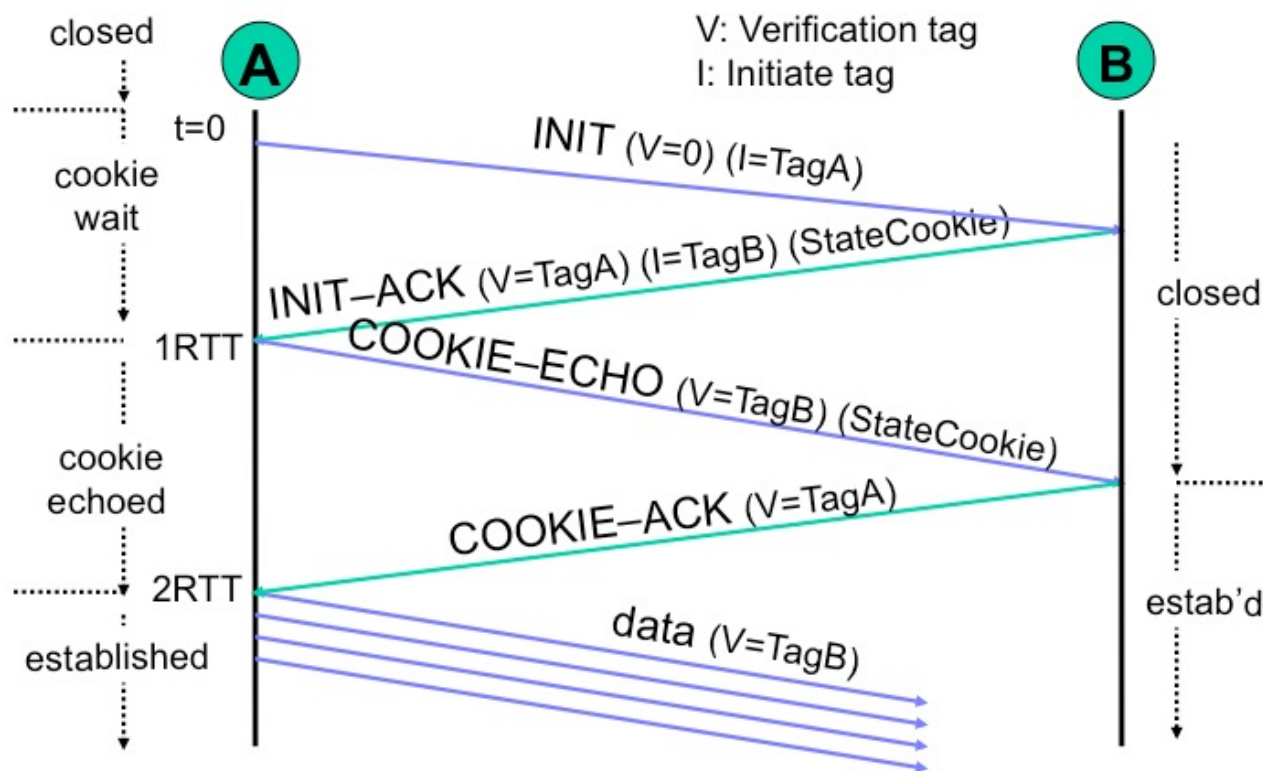
There is no ACK in response to the SYN-ACK, hence connection **remains half-open**. Other genuine clients cannot open connections.

Internet transport protocols

SCTP – 4-way handshake



Universidade do Minho
Escola de Engenharia
Departamento de Informática



- Host A sends INIT chunk to Host B, Host B returns INIT-ACK containing a cookie w/ state information that only Host B can verify; 32 bits random verification tags; **no memory allocated**
- Host A replies with COOKIE-ECHO chunk; may contain A's first data.
- Host B checks validity of cookie; connection is established

Internet transport protocols

SCTP multistreaming



Universidade do Minho
Escola de Engenharia
Departamento de Informática

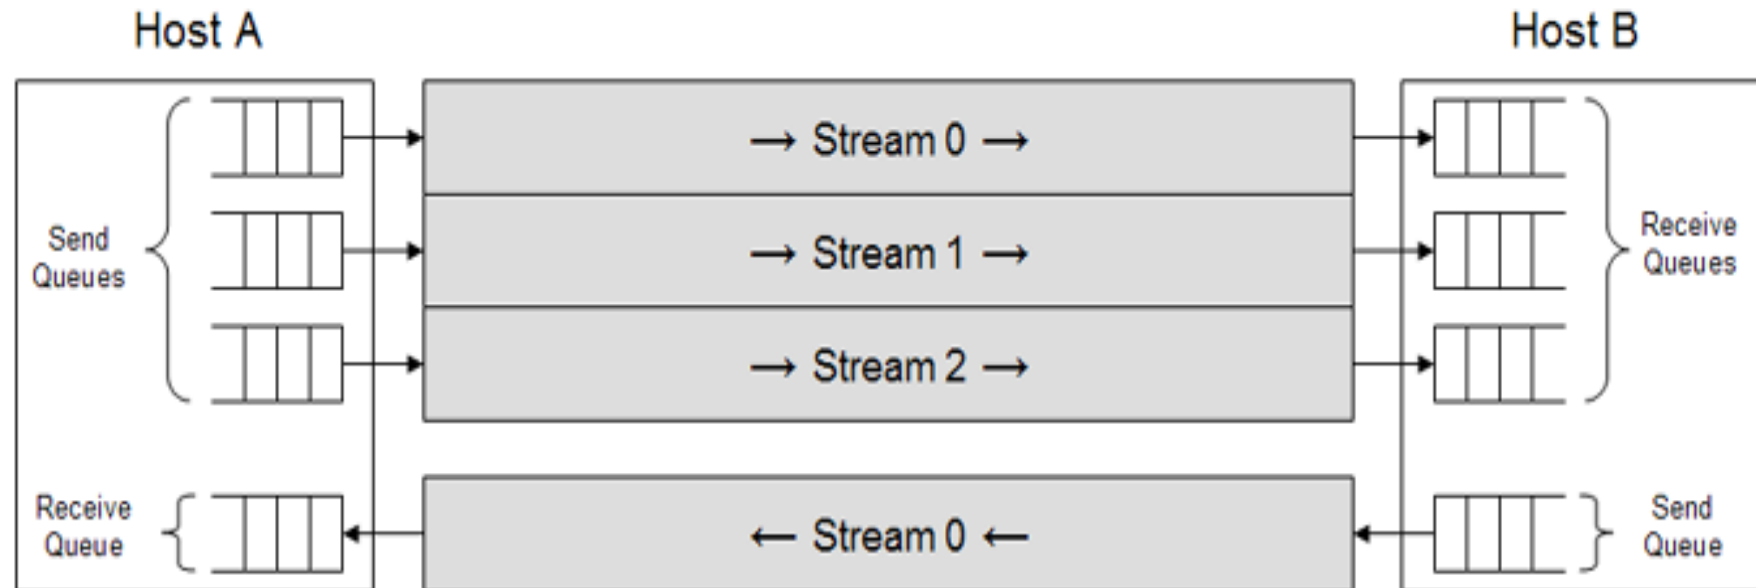
- TCP provides transport for a **single** data stream
- SCTP allows **multiple** independent streams per association
- Examples:
 - in **telephony signaling** (SS7), it is only necessary to maintain sequencing of messages that affect, e.g., the same call). Other messages are only loosely correlated and can be delivered without having to maintain overall sequence integrity.
 - **delivery of multimedia documents** over a single session. Multi-streaming allows transport of multimedia components to be partially ordered rather than strictly ordered.

Internet transport protocols

SCTP multistreaming



Universidade do Minho
Escola de Engenharia
Departamento de Informática




- Logical separation of data within an association
- Designed to prevent head-of-line blocking
- Can be used to deliver multiple objects of the same association
 - e.g.: objects within a webpage, multimedia streams (audio/video/text), files in an FTP *mget*

Internet transport protocols

SCTP multihoming



Universidade do Minho
Escola de Engenharia
Departamento de Informática




- SCTP is able to handle scenarios with multiple src/dst IP addresses (wired and wireless interfaces, multiple ISPs, etc.)

Internet transport protocols

SCTP multihoming



Universidade do Minho
Escola de Engenharia
Departamento de Informática



- Hosts pick 1 of 4 possible TCP connections:
 - $\{(A_1, B_1), (A_1, B_2), (A_2, B_1), (A_2, B_2)\}$
- Hosts use 1 SCTP association:
 - $(\{A_1, A_2\}, \{B_1, B_2\})$
 - Selectable “primary” destination
 - New data sent only to primary destination
 - Path status and reachability monitored (*heartbeats*)

Internet transport protocols

SCTP multihoming



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- TCP connection \Leftrightarrow SCTP association
 - 2 IP addresses, 2 port numbers \Leftrightarrow 2 sets of IP addresses, 2 port numbers
 - a common flow and congestion control mechanism is kept to reduce overhead
- Goal: robustness
 - automatically switch IP host addresses upon failure
 - eliminates effect of long routing reconvergence time
- **TCP**: no guarantee for “keepalive” messages when connection idle
- **SCTP** monitors each destination's reachability via ACKs of
 - data chunks
 - heartbeat chunks
- SCTP uses **multihoming for redundancy, not for load balancing!**

Internet transport protocols

SCTP comparison to other protocols



Universidade do Minho
Escola de Engenharia
Departamento de Informática

Services/Features	SCTP	TCP	UDP
• Full-duplex data transmission	yes	yes	yes
• Connection-oriented	yes	yes	no
• Reliable data transfer	yes	yes	no
• Partially reliable data transfer	optional	yes	no
• Ordered data delivery	yes	yes	no
• Unordered data delivery	yes	no	yes
• Flow and Congestion Control	yes	yes	no
• ECN support	yes	yes	no
• Selective acks	yes	optional	no
• Preservation of message boundaries	yes	no	yes
• Path MTU Discovery	yes	yes	no
• Application data fragmentation	yes	yes	no
• Multistreaming	yes	no	no
• Multihoming	yes	no	no
• Protection against SYN flooding attack	yes	no	n/a
• Half-closed connections	no	yes	n/a

Internet transport protocols

SCTP references



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- Current home: IETF TSVWG
(Transport Services Working Group)
 - <https://datatracker.ietf.org/wg/tsvwg/charter/>
- Main RFC and info
 - SCTP - RFC 4960, IETF, September 2007
 - Sockets API Extensions for the Stream Control Transmission Protocol (SCTP), RFC 6458, Dec. 2011
- Implementations
 - AIX, FreeBSD (7.0 onwards), Linux, Mac OS X, QNX, Solaris, True64, IOS (Cisco), Microsoft Windows



Universidade do Minho
Escola de Engenharia
Departamento de Informática

Eng. Serviços de Rede (ESR)

Transport Protocols Recent developments

QUIC

A UDP-Based Multiplexed and Secure Transport

(short overview)

[RFC9000,2021]

Internet transport protocols

QUIC motivation



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- Need to provide fast delivery and rendering at app level
- QUIC aims to solve TCP drawbacks (already identified), having a similar SCTP motivation by allowing:
 - HOL blocking removal
 - multiplexing different streams
 - increased security
- proposed by Jim Roslink at Google in 2012, being integrated as a Chromium component, and supported by major browsers
- currently a proposed standard, RFC 9000, May 2021

Internet transport protocols

QUIC



Universidade do Minho
Escola de Engenharia
Departamento de Informática




Figure 1: QUIC in the traditional HTTPS stack.

Internet transport protocols

QUIC



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- QUIC principles:
 - provides applications with **flow-controlled streams** for **structured**, communication, **low-latency** connection establishment, and network **path migration**.
 - includes **security measures** that ensure confidentiality, integrity, and availability based on TLS.
 - runs over TLS/UDP, implementing also **congestion control**.
 - is a connection-oriented protocol that creates a stateful interaction between a client and server.

Internet transport protocols

QUIC



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- QUIC operation:
 - **handshaking** combines negotiation of cryptographic and transport parameters (with low RTT concerns);
 - endpoints communicate exchanging **QUIC packets**;
 - packets contain **frames**, which carry **control info and application data** between endpoints;
 - app protocols send data over a QUIC connection via streams, which are ordered sequences of bytes.
 - streams can be bidirectional or unidirectional (only one endpoint can send data)

Internet transport protocols

QUIC



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- QUIC operation:
 - **flow control** relies on a credit-based scheme to limit stream creation and to bound the amount of data that can be sent.
 - provides the necessary feedback to implement **reliable delivery** and **congestion control**.
 - **path migration**, connections are not strictly bound to a single network path... a client may migrate to a new path if available (requires probing and path validation) (... e.g., endpoint with Wi-Fi and 4G access).

Internet transport protocols

QUIC info



Universidade do Minho
Escola de Engenharia
Departamento de Informática

- IETF QUIC working group
 - <https://quicwg.org>
- Current implementations: <https://github.com/quicwg/base-drafts/wiki/Implementations>
- Additional Reading
 - Adam Langley, et al. The QUIC Transport Protocol: Design and Internet-Scale Deployment. In ACM SIGCOMM '17. ACM, New York, NY, USA, 183–196. DOI:<https://doi.org/10.1145/3098822.3098842>
 - RFC 9000, May 2021.
- Must have QUIC on radar in 2022... 😊