

BigData Introduction

Ricardo Vilaça

HASLab/MACC


University of Minho

rmvilaca@di.uminho.pt



Presentation

- <https://rmpvilaca.github.io/>
- MACC & HASLab Researcher @ UM & INESC TEC
- Research Interests:
 - Energy-efficient scheduling
 - HPC, Big Data, and AI Convergence
 - Edge-Cloud-HPC Computing Continuum
 - Scalable data processing



Minho
Advanced
Computing
Center



Roadmap

- BigData and Data Lakes
- Distributed Storage
 - Object Storage
 - File Systems
- File Formats



Big Data






Image from shutterstock



Big Data

- Data Independence
- Full scan vs point queries
 - Projection and selection
- Data denormalization
 - Good for read-intensive scenarios
 - Not having to perform join improves performance



Scale challenges

- Very large amount of data:
 - Disk size
 - Caching performance
- Very large number of queries (reading):
 - Available CPU
- Very large number of queries (writing):
 - Available CPU and disk bandwidth
- Heterogeneous data
 - Nested data
 - Variety of data structures
- I/O bottleneck
 - High disk access time with respect to main memory access time
 - Ever growing processor speeds
- (And high availability...)



ETL vs ELT

- ETL (Extract, Transform, Load)
 - Imported into the database
 - Transform before load to a data warehouse
 - Schema on write
 - Complex transformations of smaller data sets
- ELT (Extract, Load, Transform)
 - Read from a file system (data lake)
 - Schema on read
 - More flexibility
 - Both structured and unstructured data

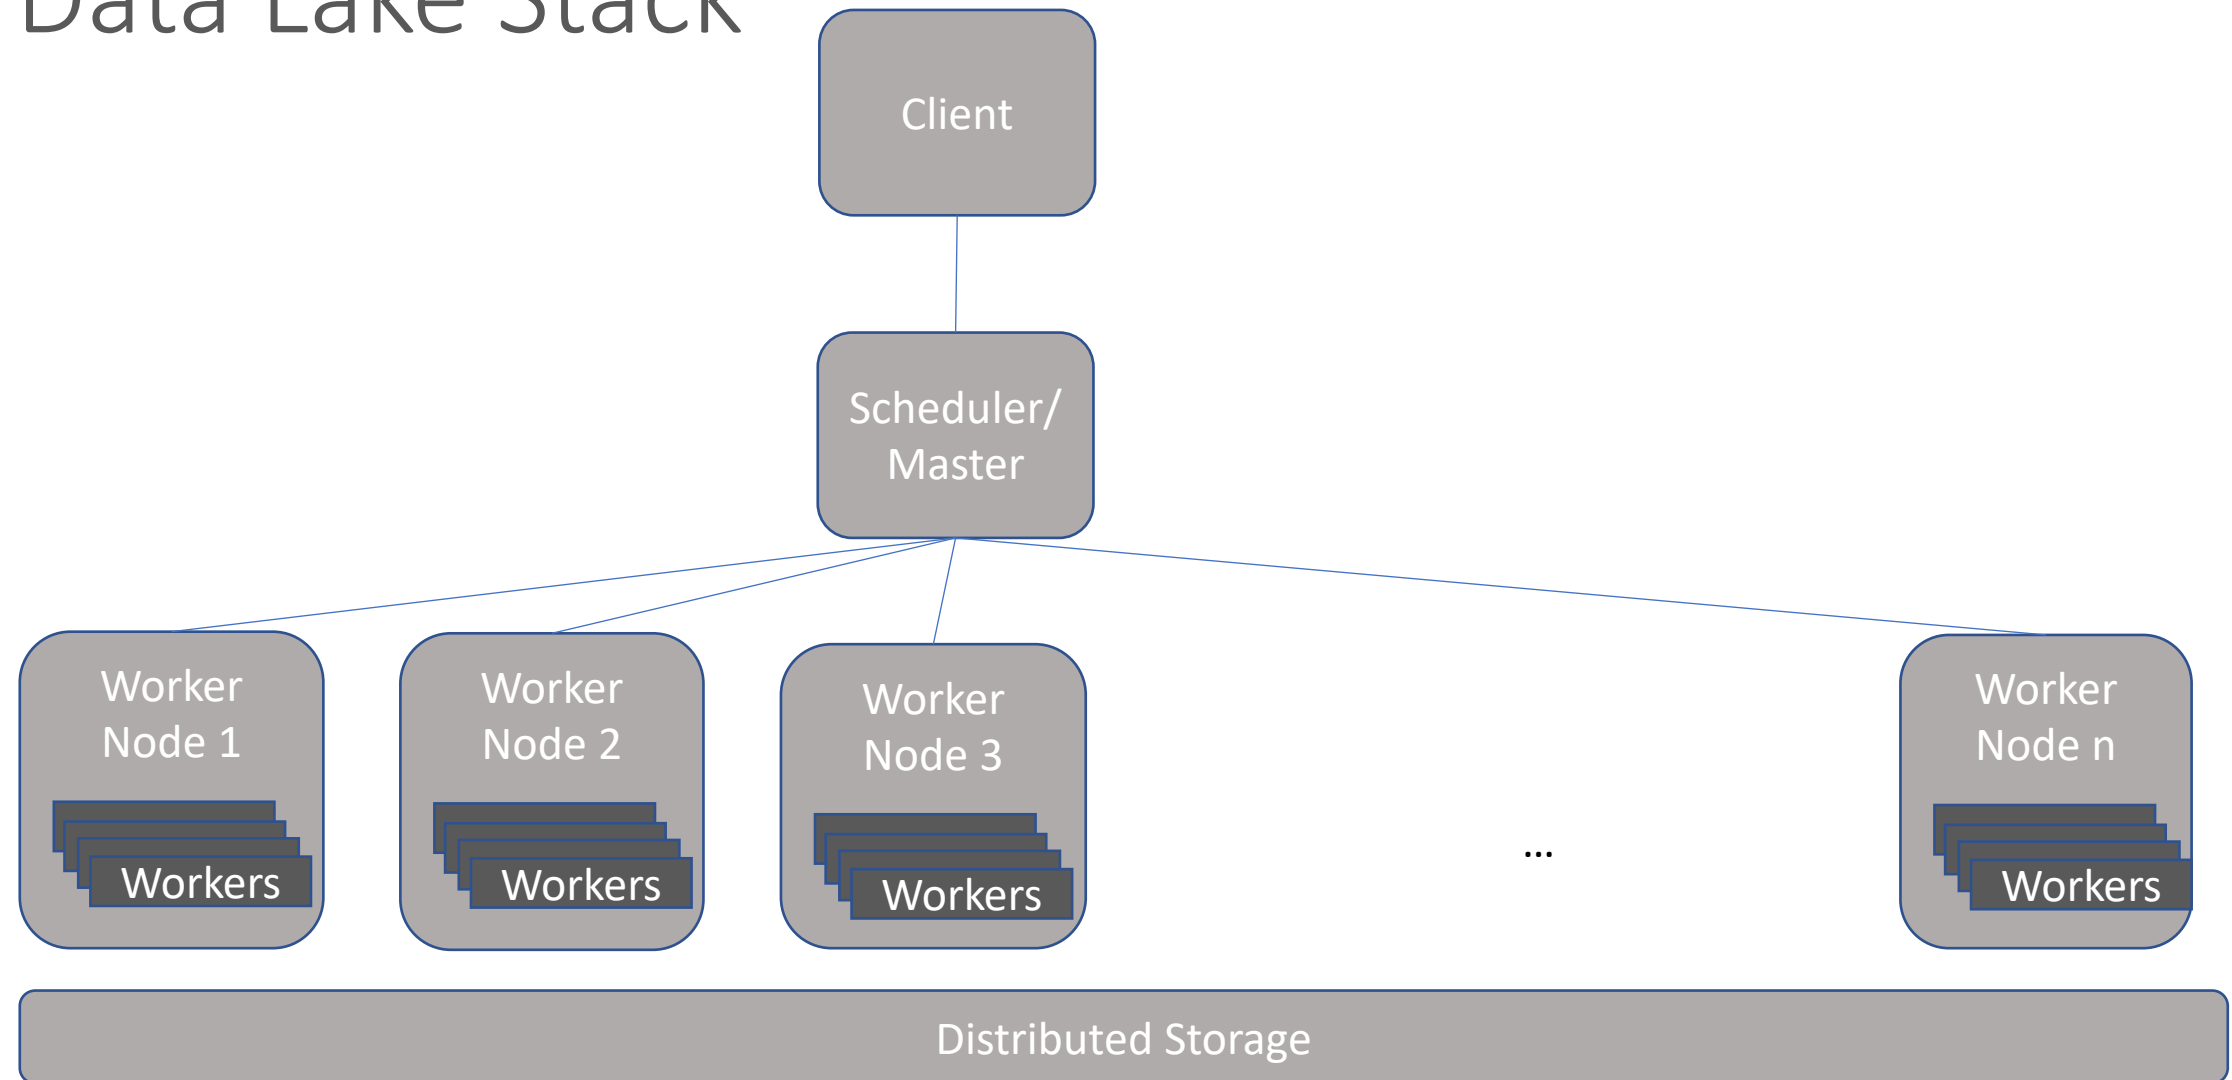


Push vs Pull

- Push Query to Data
 - Send the query (or a portion of it) to the node that contains the data
 - Perform as much filtering and processing as possible where data resides before transmitting over network
- Pull Data to Query
 - Bring the data to the node that is executing a query that needs it for processing
 - This is necessary when there is no compute resources available where persistent data files are located




Data Lake Stack



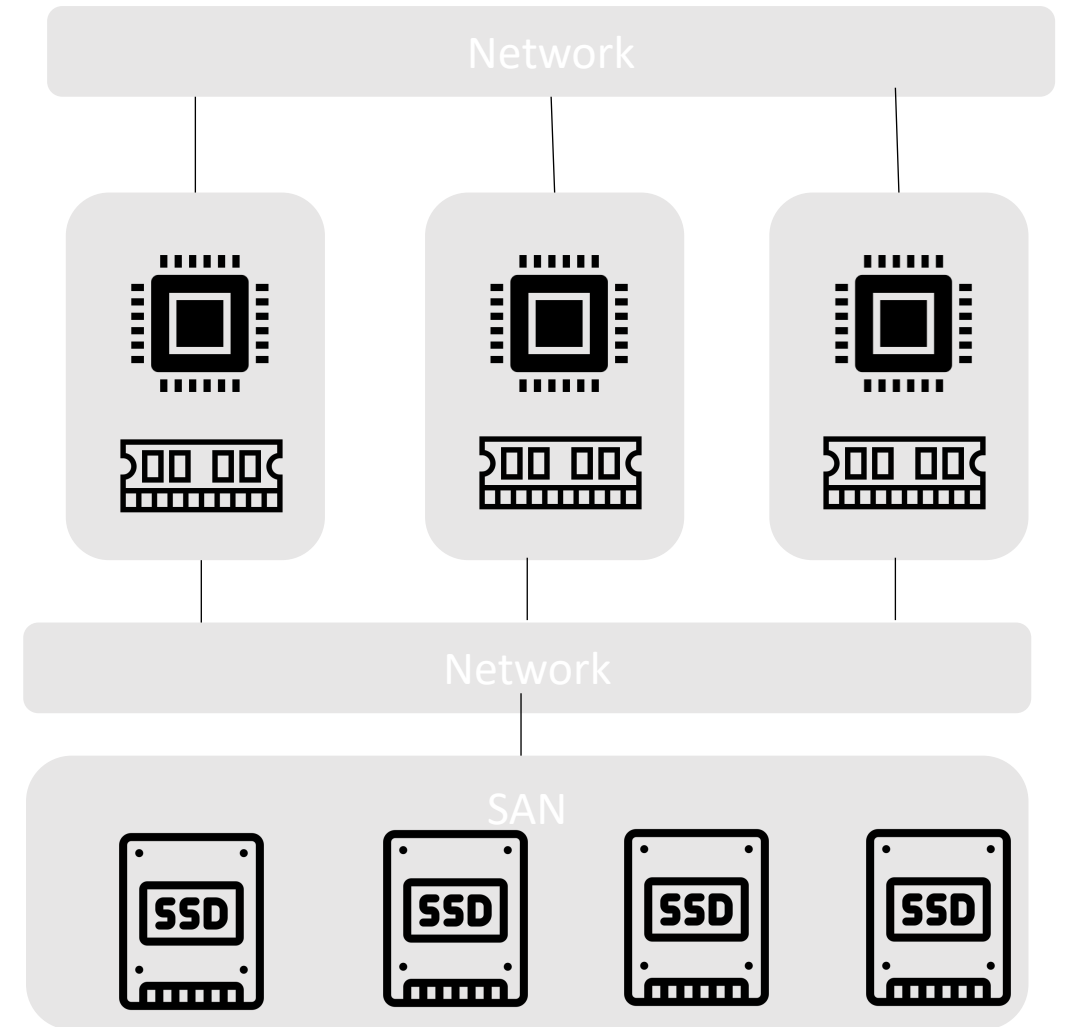
Shared Nothing

- Each instance has its own CPU, memory, locally-attached disk.
- Nodes only communicate with each other via network.
- Data partitioned into disjoint subsets across nodes.
- Cost-effective
 - off-the-shelf components
- Coupling with distributed file systems



Shared Disk

- Each node accesses a single logical disk via an interconnect, but also have their own private memory and ephemeral storage.
 - Must send messages between nodes to learn about their current state
- Cloud Object stores



Shared Nothing vs Shared Disk

- Shared Disk
 - Scale compute layer independently from the storage layer
 - Easy to shutdown idle compute layer resources
 - May need to pull uncached persistent data from storage layer to compute layer before applying filters
- Shared Nothing
 - Dominant parallel architecture for big data systems
 - Harder to scale capacity (data movement)
 - Potentially better performance & efficiency
 - Apply filters where the data resides before transferring



Distributed Challenges

- Nodes fail
 - 1 in 1000 nodes fail a day
 - *Duplicate Data*
- Network is a bottleneck
 - Typically 1-10 Gb/s throughput
 - **Bring computation to nodes, rather than data to nodes**
- Traditional distributed programming is
 - often ad-hoc and complicated
 - **Stipulate a programming system that can easily be distributed (MapReduce/Spark)**



Distributed Storage



Distributed Storage System

Storing and managing data across the nodes of a cluster

- Object-based
 - Object = $\langle \text{oid}, \text{data}, \text{metadata} \rangle$
 - Metadata can be different for different object
 - Easy to move
 - Flat object space \rightarrow billions/trillions of objects
 - Easily accessed through REST-based API (get/put)
- File-based
 - Data in files of fixed- or variable-length records
 - Metadata-per-file stored separately from file
 - For large data, a file needs to be partitioned and distributed



Object Store

- Partition the persistent data into large, immutable files stored in an object store
 - All attributes for a tuple are stored in the same file in a columnar layout (PAX)
 - Header (or footer) contains meta-data about columnar offsets, compression schemes, indexes, and zone maps
- No hierarchy
- Each cloud vendor provides their own proprietary API to access data (PUT, GET, DELETE).
 - Some vendors support predicate pushdown (S3)

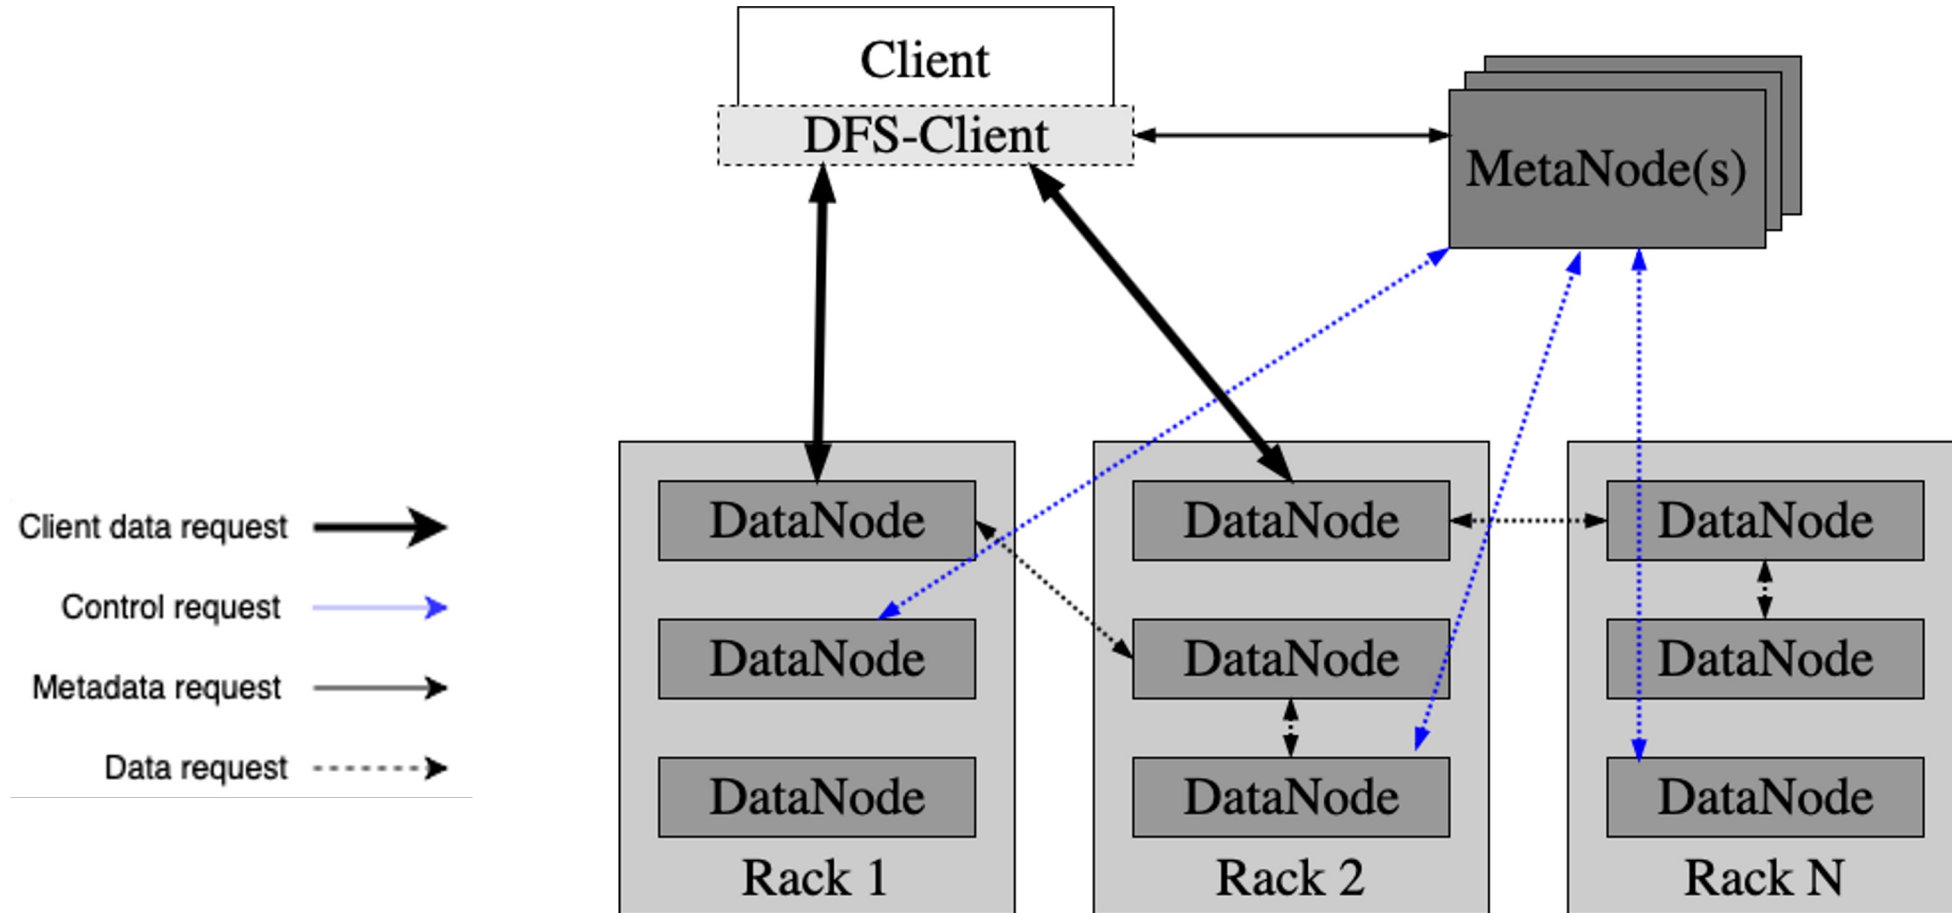


Hadoop Distributed File System (HDFS)

- Distributed architecture
- Master/slave topology
- Independent data and metadata handling
- POSIX-like Interface
- Designed for large data sets
- Throughput over latency



HDFS Architecture



NameNode

- Maintains the file system metadata
 - *e.g.*, namespace, access control, file-to-block mapping, block locations
- Manages and controls system-wide activities (*e.g.*, load balancing, locking)
- HDFS-client <-> NameNode
 - Exchange metadata requests
- NameNode <-> DataNode
 - Periodically exchange block reports and heartbeats

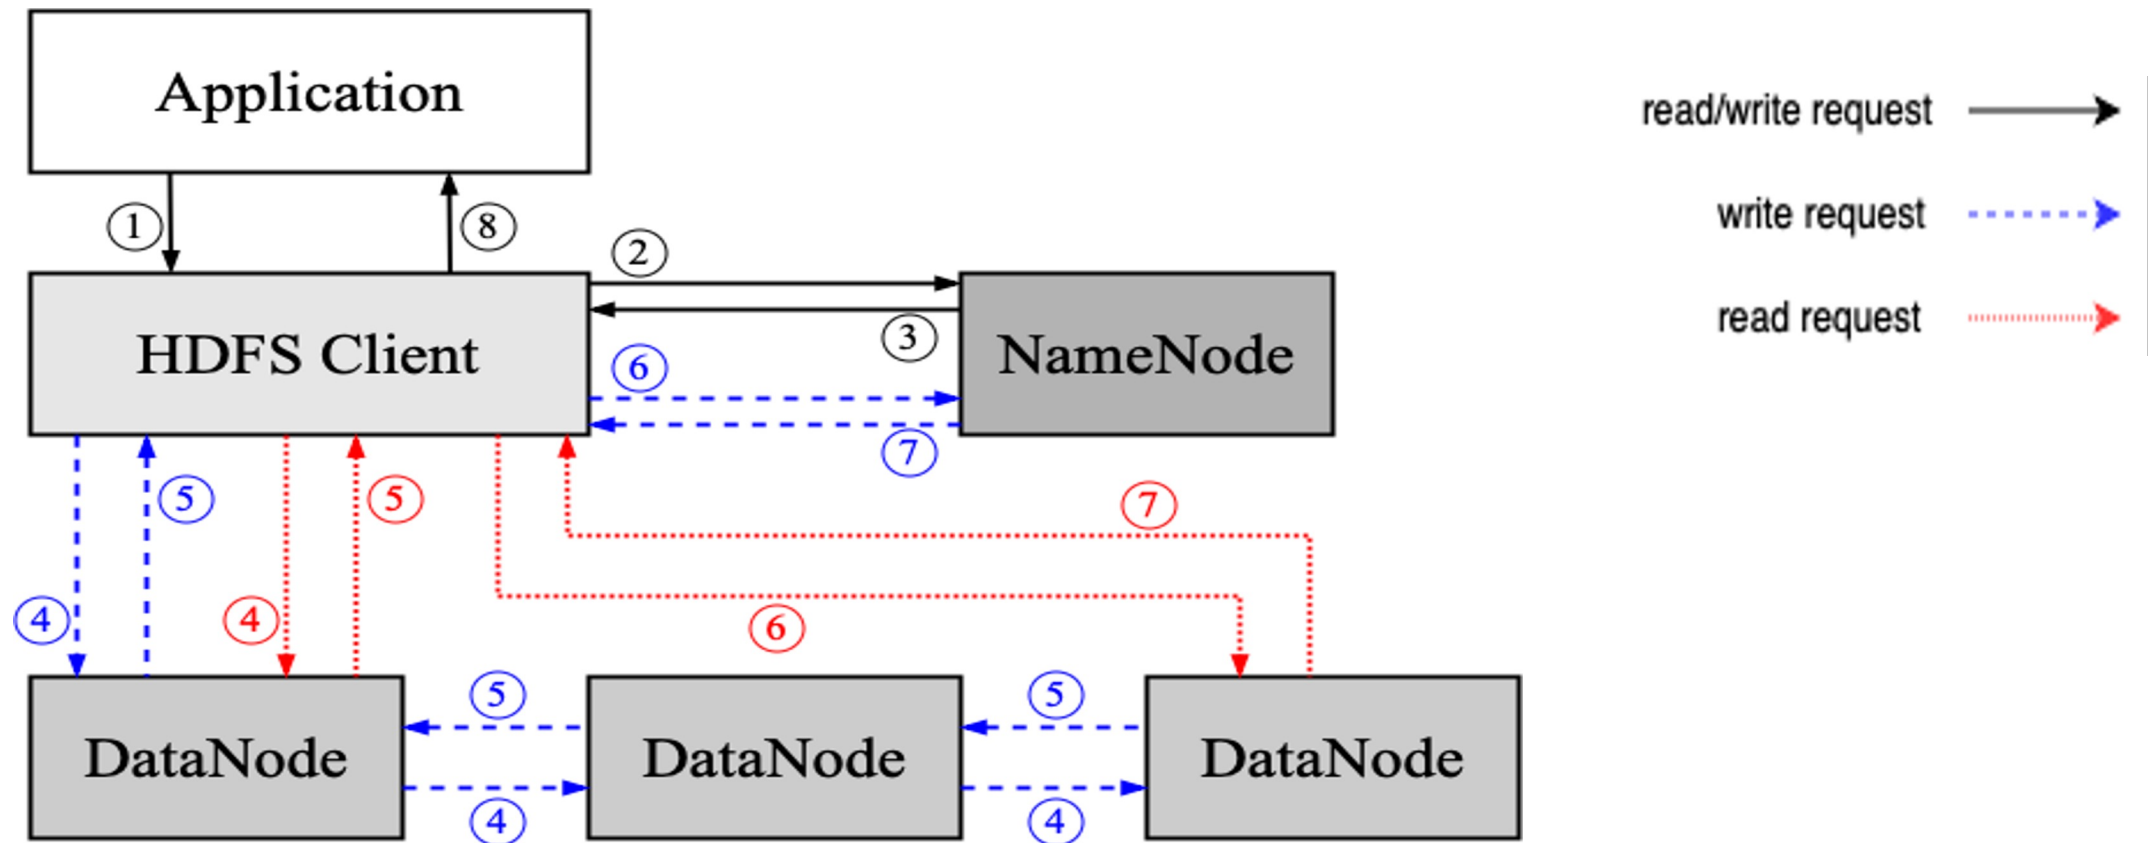


DataNode

- Storage units of the file system
- Store and retrieve files/blocks to HDFS-clients
- Files are split into blocks
 - *dfs.blocksize* is typically 64MiB or 128MiB
 - Blocks are stored across different DataNodes for increased availability and parallelism (defined by *dfs.replication*)
 - Default replication factor is 3



Requests lifecycle



File vs Objects

- File systems
 - tied to a particular hierarchical organisation and storage device
 - files can be modified or accessed arbitrarily
 - size limit determined by underlying storage facility
 - accessed by programs running in a host that mounts the file system
- Objects
 - immutable
 - data storage not bound logically to any storage facility
 - whole-object operations
 - accessible over the Internet



File Formats



Relational Data Format

- Tabular data
- Multiple data types
- Optional (null) values
- No nested or repeated values
- Large number of columns

<i>Id</i>	<i>Name</i>	<i>Location</i>
1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	
5	eee	Lisboa
...



Text (CSV)

- Row-oriented
- Simple to produce and consume
- Schema can be inferred
- Redundancy and verbose representation (numbers)
- Ambiguity in separators and missing fields
- Only primitive types
- Difficult to page, especially when compressed

<i>Id</i>	<i>Name</i>	<i>Location</i>
1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	
5	eee	Lisboa
...		...

data.csv "1","aa","Braga"
"2","bbb","Porto"
"3","cc","Porto"
"4","dddddd",
"5","eee","Lisboa"
...,...,...



Data Formats Challenges


- Representation of types
 - Compactness and ambiguity
- Data that needs to be moved for:
 - Selection (range scan)
 - Projection
- Compression



Binary rows

- Compact and unambiguous
- Efficient I/U/D
- Can be paged and compressed
 - Not efficient as different data types are interleaved
- All data is read for projections

<i>Id</i>	<i>Name</i>	<i>Location</i>
1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	
5	eee	Lisboa
...		...



data

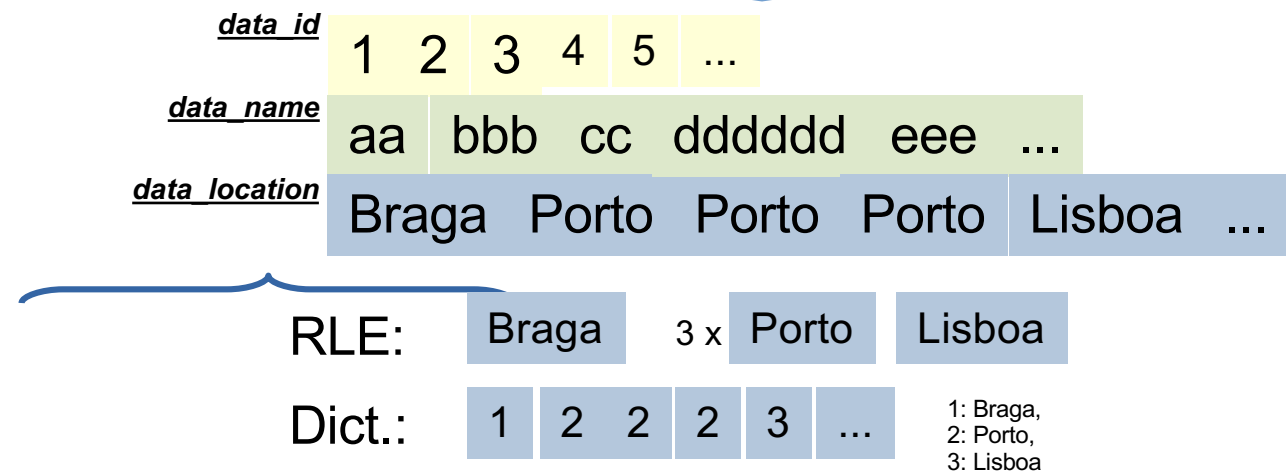
1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	Porto
5	eee	Lisboa
...



Columnar

- Efficient projections
- Compressed very efficiently
 - Dictionary and/or
 - Run Length Encoding (RLE)
- Inefficient I/U/D
- Inefficient range scan


<i>Id</i>	<i>Name</i>	<i>Location</i>
1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	
5	eee	Lisboa
...		...



Hybrid

- Columnar segments, that can be accessed and compressed separately
- Good trade-off:
 - I/U/D updates only one segment
 - Range scans can read only some segments
 - Projections can easily skip columns

<i>Id</i>	<i>Name</i>	<i>Location</i>
1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	
5	eee	Lisboa
...		...



Hierarchical data

- Data that is not normalized (in a relational sense)
 - Nested structures
 - Repeated fields
- Useful as it avoids multiple files and foreign keys

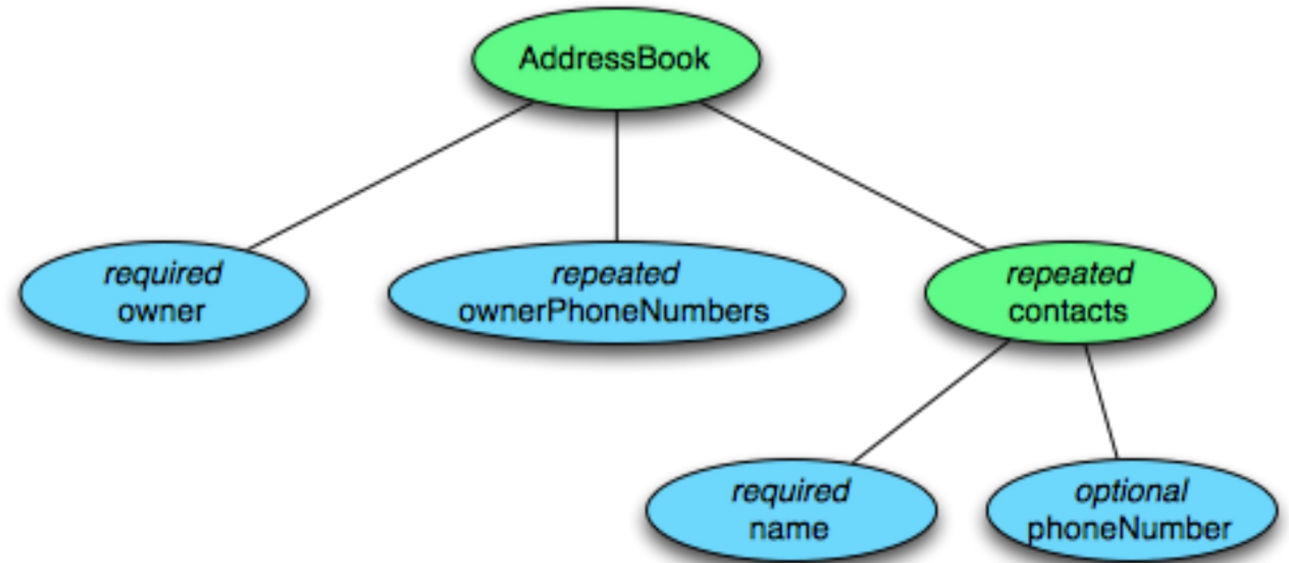


Image from: https://blog.twitter.com/engineering/en_us/a/2013/dremel-made-simple-with-parquet.html



JSON

- Well-known and widely supported
- No explicit schema
- Row-based
- Not splittable
- Complex data like structs and arrays

```
{
  "AddressBook": [
    {
      "owner": "Jason F.",
      "ownerPhoneNumbers": [
        "123456789",
        "987654321"
      ],
      "contacts": [
        { "name": "John" },
        { "name": "Joe", "number":
"214365879" }
      ]
    },
    {
      "owner": "Joe G.",
      "ownerPhoneNumbers": [
        "214365879"
      ]
    }
  ]
}
```



Types of metadata

- Technical
 - Types, representation, ...
- Operational
 - Location (indexing), cardinality, ...
- Business
 - What it means, quality, ...




Schema

- Information about data items and types
- Implicit, central or embedded:

.java

```
"1", "aa", "Braga"  
"2", "bbb", "Porto"  
"3", "cc", "Porto"  
"4", "dddddd",  
"5", "eee", "Lisboa"  
..., ..., ...
```



data

1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	Porto
5	eee	Lisboa
...

data_schema

```
"id": "integer",  
"name": "string",  
"location": "string"
```

data


1	aa	Braga
2	bbb	Porto
3	cc	Porto
4	dddddd	Porto
5	eee	Lisboa
...



Partitions

- Partition files by a low cardinality column
- Encode partition key in the file name
- Used often with locations and dates
- Useful to avoid reading data

<u>data</u>	1	aa	Braga
	2	bbb	Porto
	3	cc	Porto
	4	dddddd	Porto
	5	eee	Lisboa



<u>data-Braga</u>	1	aa	Braga

<u>data-Porto</u>	2	bbb	Porto
	3	cc	Porto
	4	dddddd	Porto

<u>data-Lisboa</u>	5	eee	Lisboa



Value summaries / indexes

- Range [min,max] of values in each column
- Compact representation (e.g., Bloom filter) of values in each column
- Useful to avoid reading data

Ranges (min/max)

id: [1,3]

id: [4,...]

data 0001

1	2	3
aa	bbb	cc
Braga	Porto	Porto
metadata		

data 0002

4	5	...
dddddd	eee	...
Porto	Lisboa	...
metadata		

Bloom filter

Aveiro	001010001011001
Braga	
Coimbra	
Lisboa	
Porto	
Setúbal	
	100101000111001



Cardinality summaries

- Number of distinct values in each column
- Compact representation (e.g., histogram) of repetitions of values in intervals, for each column
- Useful to predict how much data will be processed and stored

Counts

#id: 3

#id: ...

data_0001

1 2 3

aa bbb cc

Braga Porto Porto

metadata

data_0002

4 5 ...

dddddd eee ...

Porto Lisboa ...

metadata

Histograms



Aveiro
Braga
Coimbra
Lisboa
Porto
Setúbal



Compression tradeoffs

- Compression ratio vs splittable
- Cold data vs hot data
- Codecs
 - SNAPPY
 - GZIP
 - LZO



Big Data File Formats

- Open-source binary file formats that make it easier to access data across systems and libraries for extracting data from files
 - Libraries provide an iterator interface to retrieve (batched) columns from files
- Highly efficient data compression techniques
- Support for schema evolution
- Faster analytics workloads
 - Less I/O usage
- Splittable file formats
 - Spread between more than one worker node



Avro file

- JSON for data types and protocols
- Row-based
- Compacts binary format but not efficient data compression
- Language-neutral data serialization system
- Schemas: Primitive, Records, Enums, Arrays, Maps ...
 - Supports evolution of schemas
- Efficient for use with write-intensive, big data operations.



ORC File

- Compressed columnar storage from Apache Hive
- Block-mode compression
- Data type support
- Ordered data store (within one stripe)
- Indices with column-level aggregated values (min, max, sum, ...)



Apache Parquet

- Compressed columnar storage from Cloudera/Twitter
- Highly integrated with Apache Spark
- Supports (page) compression and splitting
- Supports nested columns (Dremel encoding)
- Supports complex nested data structures in a flat columnar format
- Supports minimal number of types
- Uses data skipping to locate specific column values



Apache Parquet File Format

- Row group: A logical horizontal partitioning of the data into rows
- Column chunk: A chunk of the data for a particular column.
- Page: Column chunks are divided up into pages written back to back.

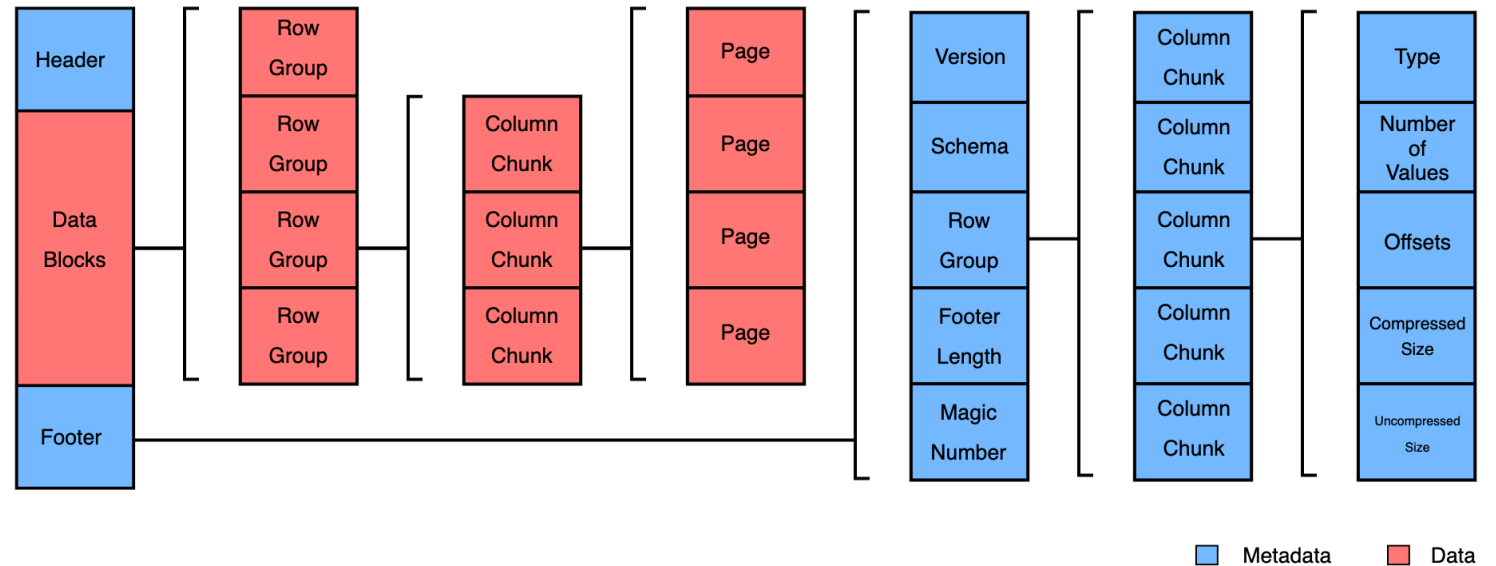



Image from <https://dkharazi.github.io/blog/parquet>



Apache Parquet File

Column-Oriented Data



Parquet Data Block

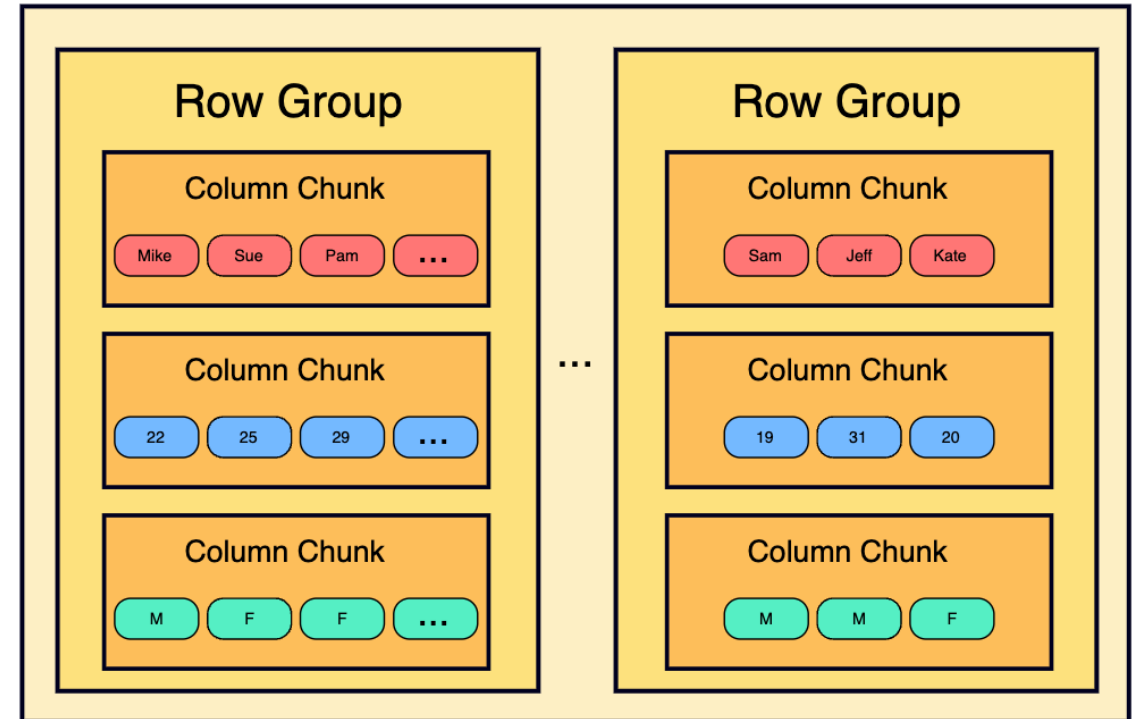


Image from <https://dkharazi.github.io/blog/parquet>



Apache Arrow

- Language-independent columnar memory format made for flat and hierarchical data
- Efficient analytic operations on modern hardware, CPUs and GPUs
- Primarily in-memory compressed columnar storage for vectorized processing
- Complementary to Parquet
 - Allows easier and more efficient movement of data from RAM to disk




Image from <https://arrow.apache.org/overview/>



Others formats

- [Apache CarbonData](#)
 - Compressed columnar storage with indexes from Huawei
- [Apache Iceberg](#)
 - Flexible data format that supports schema evolution from Netflix
- [Delta Lake](#)
 - Enables building lakehouses, with ACID transactions, time travel, ...
- Not everything tabular
 - Array
 - [HDF5](#)
 - Multi-dimensional arrays for scientific workloads
 - [Zarr](#)
 - Zarr is a format for the storage of chunked, compressed, N-dimensional arrays
 - Graph
 - [RDF](#)
 - [JSON-LD](#)



More information

- <https://www.researchgate.net/publication/361334530> The Big Data Textbook - teaching large-scale databases in universities
- <https://github.com/apache/parquet-format/blob/master/BloomFilter.md>
- <https://github.com/apache/parquet-format/blob/master/Compression.md>
- Sergey Melnik, Andrey Gubarev, Jing Jing Long, Geoffrey Romer, Shiva Shivakumar, Matt Tolton, and Theo Vassilakis. 2010. Dremel: interactive analysis of web-scale datasets. Proc. VLDB Endow.

