

Relational Operators

Ricardo Vilaça

HASLab

University of Minho

rmvilaca@di.uminho.pt



Roadmap

- Relational Operators in MapReduce
- SparkSQL Catalyst



MapReduce Implementations of relational operators

- Select and Project can be easily implemented in the map function
- Aggregation is not difficult
- Join requires more work



Select

- Selections do not need a full-blown MapReduce implementation
 - They can be implemented in the map phase alone
 - Or could also be implemented in the reduce portion
- Map
 - For each tuple t in R , check if t satisfies C
 - If so, emit a key/value pair $(t; t)$
- Reduce
 - Identity reducer



Select B ≤ 3

Map Worker 1

<i>File 1</i>		<i>File 2</i>	
<i>A</i>	<i>B</i>	<i>A</i>	<i>B</i>
1	2	2	8
2	3	4	4
5	6	6	1

Map Worker 2

<i>File 1</i>		<i>File 2</i>	
<i>A</i>	<i>B</i>	<i>A</i>	<i>B</i>
6	2	9	8
6	3	3	3
7	6	0	1



Select B ≤ 3

Map Worker 1

<i>key</i>	<i>value</i>
(1,2)	(1,2)
(2,3)	(2,3)
(6,1)	(6,1)

Map Worker 2

<i>key</i>	<i>value</i>
(6,2)	(6,2)
(6,3)	(6,3)
(3,3)	(3,3)
(0,1)	(0,1)



Select B ≤ 3

Map Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	(1,2)
(2,3)	(2,3)

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,1)	(6,1)

Map Worker 2

<i>RW1</i>	
<i>key</i>	<i>value</i>
(3,3)	(3,3)
(0,1)	(0,1)

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,2)	(6,2)
(6,3)	(6,3)



Select B ≤ 3

Reduce Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	(1,2)
(2,3)	(2,3)

<i>RW1</i>	
<i>key</i>	<i>value</i>
(3,3)	(3,3)
(0,1)	(0,1)

Reduce Worker 2

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,1)	(6,1)

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,2)	(6,2)
(6,3)	(6,3)



Select B ≤ 3

Reduce Worker 1

<i>File 1</i>	
<i>A</i>	<i>B</i>
0	1
1	2
2	3
3	3

Reduce Worker 2

<i>File 1</i>	
<i>A</i>	<i>B</i>
6	1
6	2
6	3



Projection

- Similar process to selection
- Projection may cause same tuple to appear several times
- Map
 - For each tuple t in R , construct a tuple t_0 by eliminating those components whose attributes are not in S
 - Emit a key/value pair $(t_0; t_0)$
- Reduce
 - For each key t_0 produced by any of the Map tasks, fetch $t_0; [t_0; \dots ; t_0]$
 - Emit a key/value pair $(t_0; t_0)$
- NOTE: the reduce operation is duplicate elimination



Projection (A,B)

Map Worker 1

<i>File 1</i>			<i>File 2</i>		
<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>
1	2	3	4	2	1
2	2	2	6	8	4
1	2	1	3	2	2

Map Worker 2

<i>File 1</i>			<i>File 2</i>		
<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>
1	2	5	3	2	1
2	3	2	6	8	9
1	3	1	3	4	2



Projection (A,B)

Map Worker 1

key	value
(1,2)	[(1,2),(1,2)]
(2,2)	[(2,2)]
(4,2)	[(4,2)]
(6,8)	[(6,8)]
(3,2)	[(3,2)]

Map Worker 2

key	value
(1,2)	[(1,2)]
(2,3)	[(2,3)]
(1,3)	[(1,3)]
(3,2)	[(3,2)]
(6,8)	[(6,8)]
(3,4)	[(3,4)]



Projection (A,B)

Map Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[(1,2),(1,2)]
(2,2)	[(2,2)]
(4,2)	[(4,2)]

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,8)	[(6,8)]
(3,2)	[(3,2)]

Map Worker 2

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[(1,2)]
(2,3)	[(2,3)]
(1,3)	[(1,3)]

<i>RW2</i>	
<i>key</i>	<i>value</i>
(3,2)	[(3,2)]
(6,8)	[(6,8)]
(3,4)	[(3,4)]



Projection (A,B)

Reduce Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[(1,2),(1,2)]
(2,2)	[(2,2)]
(4,2)	[(4,2)]

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[(1,2)]
(2,3)	[(2,3)]
(1,3)	[(1,3)]

Reduce Worker 2

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,8)	[(6,8)]
(3,2)	[(3,2)]

<i>RW2</i>	
<i>key</i>	<i>value</i>
(3,2)	[(3,2)]
(6,8)	[(6,8)]
(3,4)	[(3,4)]



Projection (A,B)

Reduce Worker 1

key	value
(1,2)	[(1,2),(1,2),(1,2)]
(1,3)	[(1,3)]
(2,2)	[(2,2)]
(2,3)	[(2,3)]
(4,2)	[(4,2)]

Reduce Worker 2

key	value
(3,2)	[(3,2),(3,2)]
(3,4)	[(3,4)]
(6,8)	[(6,8),(6,8)]



Projection (A,B)

Reduce Worker 1

<i>File 1</i>	
<i>A</i>	<i>B</i>
1	2
1	3
2	2
2	3
4	2

Reduce Worker 2

<i>File 1</i>	
<i>A</i>	<i>B</i>
3	2
3	4
6	8



Union

- Suppose relations R and S have the same schema
 - Map tasks will be assigned chunks from either R or S
 - Mappers don't do much, just pass by to reducers
 - Reducers do duplicate elimination
- A MapReduce implementation of union
 - Map
 - For each tuple t in R or S, emit a key/value pair $(t; t)$
 - Reduce
 - For each key t there will be either one or two values
 - Emit $(t; t)$ in either case



Intersection

- Very similar to computing unions
 - Suppose relations R and S have the same schema
 - The map function is the same (an identity mapper) as for union
 - The reduce function must produce a tuple only if both relations have that tuple
- A MapReduce implementation of intersection
 - Map
 - For each tuple t in R or S, emit a key/value pair (t; t)
 - Reduce
 - If key t has value list [t; t] then emit the key/value pair (t; t)
 - Otherwise, emit the key/value pair (t; NULL)



GroupBy A AGG(B)

- Let $R(A;B;C)$
 - The map operation prepares the grouping
 - The grouping is done by the framework
 - The reducer computes the aggregation
 - Simplifying assumptions: one grouping attribute and one aggregation function
- Map
 - For each tuple $(a; b; c)$ emit the key/value pair $(a; b)$
- Reduce
 - Each key a represents a group
 - Apply AGG to the list $[b1; b2; \dots; bn]$
 - Emit the key/value pair $(a; x)$ where $x = AGG([b1; b2; \dots; bn])$



GroupBy (A,B) Sum(C)

Map Worker 1

<i>File 1</i>				<i>File 2</i>			
<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
1	2	3	1	4	2	1	3
2	2	3	2	6	8	4	4
1	2	1	3	3	2	2	4

Map Worker 2

<i>File 1</i>				<i>File 2</i>			
<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
1	2	5	2	3	2	1	3
2	3	2	4	2	3	9	2
1	3	1	3	3	4	2	1



GroupBy (A,B) Sum(C)

Map Worker 1

<i>key</i>	<i>value</i>
(1,2)	[3,1]
(2,2)	[3]
(4,2)	[1]
(6,8)	[4]
(3,2)	[2]

Map Worker 2

<i>key</i>	<i>value</i>
(1,2)	[5]
(2,3)	[2,9]
(1,3)	[1]
(3,2)	[1]
(3,4)	[2]



GroupBy (A,B) Sum(C)

Map Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[3,1]
(2,2)	[3]
(4,2)	[1]

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,8)	[4]
(3,2)	[2]

Map Worker 2

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[5]
(2,3)	[2,9]

<i>RW2</i>	
<i>key</i>	<i>value</i>
(3,2)	[1]
(3,4)	[2]
(1,3)	[1]



GroupBy (A,B) Sum(C)

Reduce Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[3,1]
(2,2)	[3]
(4,2)	[1]

<i>RW1</i>	
<i>key</i>	<i>value</i>
(1,2)	[5]
(2,3)	[2,9]

Reduce Worker 2

<i>RW2</i>	
<i>key</i>	<i>value</i>
(6,8)	[4]
(3,2)	[2]

<i>RW2</i>	
<i>key</i>	<i>value</i>
(3,2)	[1]
(3,4)	[2]
(1,3)	[1]



GroupBy (A,B) Sum(C)

Reduce Worker 1

key	value
(1,2)	[3,1,5]
(2,2)	[3]
(2,3)	[2,9]
(4,2)	[1]

Reduce Worker 2

key	value
(1,3)	[1]
(3,2)	[1,2]
(3,4)	[2]
(6,8)	[4]



GroupBy (A,B) Sum(C)

Reduce Worker 1

<i>A</i>	<i>B</i>	<i>Sum</i>
1	2	9
2	2	3
2	3	11
4	2	1

Reduce Worker 2

<i>A</i>	<i>B</i>	<i>Sum</i>
1	3	1
3	2	3
3	4	2
6	8	4



Natural Join

- Let's look at two relations $R(A;B)$ and $S(B;C)$
 - We must find tuples that agree on their B components
 - We shall use the B-value of tuples from either relation as the key
 - The value will be the other component and the name of the relation
 - That way the reducer knows from which relation each tuple is coming from
- A MapReduce implementation of Natural Join
 - Map
 - For each tuple $(a; b)$ of R emit the key/value pair $(b; (R; a))$
 - For each tuple $(b; c)$ of S emit the key/value pair $(b; (S; c))$
 - Reduce
 - Each key b will be associated to a list of pairs that are either $(R; a)$ or $(S; c)$
 - Emit key/value pairs of all possible combinations for the values where one value is from table R and the other value is from table S



Natural Join

Map Worker 1

<i>Table 1</i>	
<i>A</i>	<i>B</i>
1	2
2	3
5	6

<i>Table 2</i>	
<i>B</i>	<i>C</i>
2	3
4	4
6	1

Map Worker 2

<i>Table 1</i>	
<i>A</i>	<i>B</i>
6	1
6	3
7	6

<i>Table 2</i>	
<i>B</i>	<i>C</i>
9	8
3	4
2	1



Natural Join

Map Worker 1

<i>key</i>	<i>value</i>
2	[(T1,1), (T2,3)]
3	[(T1,2)]
6	[(T1,5), (T2,1)]
4	[(T1,4)]

Map Worker 2

<i>key</i>	<i>value</i>
1	[(T1,6)]
3	[(T1,6),(T2,4)]
6	[(T1,7)]
9	[(T2,8)]
2	[(T2,1)]



Natural Join

Map Worker 1

<i>RW1</i>		<i>RW2</i>	
<i>key</i>	<i>value</i>	<i>key</i>	<i>value</i>
2	[(T1,1),(T2,3)]	6	[(T1,5), (T2,1)]
3	[(T1,2)]	4	[(T1,4)]

Map Worker 2

<i>RW1</i>		<i>RW2</i>	
<i>key</i>	<i>value</i>	<i>key</i>	<i>value</i>
1	[(T1,6)]	6	[(T1,7)]
3	[(T1,6),(T2,4)]	9	[(T2,8)]
2	[(T2,1)]		



Natural Join

Reduce Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
2	[(T1,1),(T2,3)]
3	[(T1,2)]

<i>RW1</i>	
<i>key</i>	<i>value</i>
1	[(T1,6)]
3	[(T1,6),(T2,4)]
2	[(T2,1)]

Reduce Worker 2

<i>RW2</i>	
<i>key</i>	<i>value</i>
6	[(T1,5),(T2,1)]
4	[(T1,4)]

<i>RW2</i>	
<i>key</i>	<i>value</i>
6	[(T1,7)]
9	[(T2,8)]



Natural Join

Reduce Worker 1

<i>RW1</i>	
<i>key</i>	<i>value</i>
1	[(T1,6)]
2	[(T1,1), (T2 ,3), (T2 ,1)]
3	[(T1,2), (T1,6),(T2 ,4)]

Reduce Worker 2

<i>RW2</i>	
<i>key</i>	<i>value</i>
6	[(T1,5), (T2 ,1), (T1,7)]
4	[(T1,4)]
9	[(T2 ,8)]



Natural Join

Reduce Worker 1

<i>B</i>	<i>A</i>	<i>C</i>
2	1	3
2	1	1
3	2	4
3	6	4

Reduce Worker 2

<i>B</i>	<i>A</i>	<i>C</i>
6	5	1
6	7	1



Other joins

- Applied to other complex operators such as duplicate elimination, union, intersection, etc. with minor adaptation

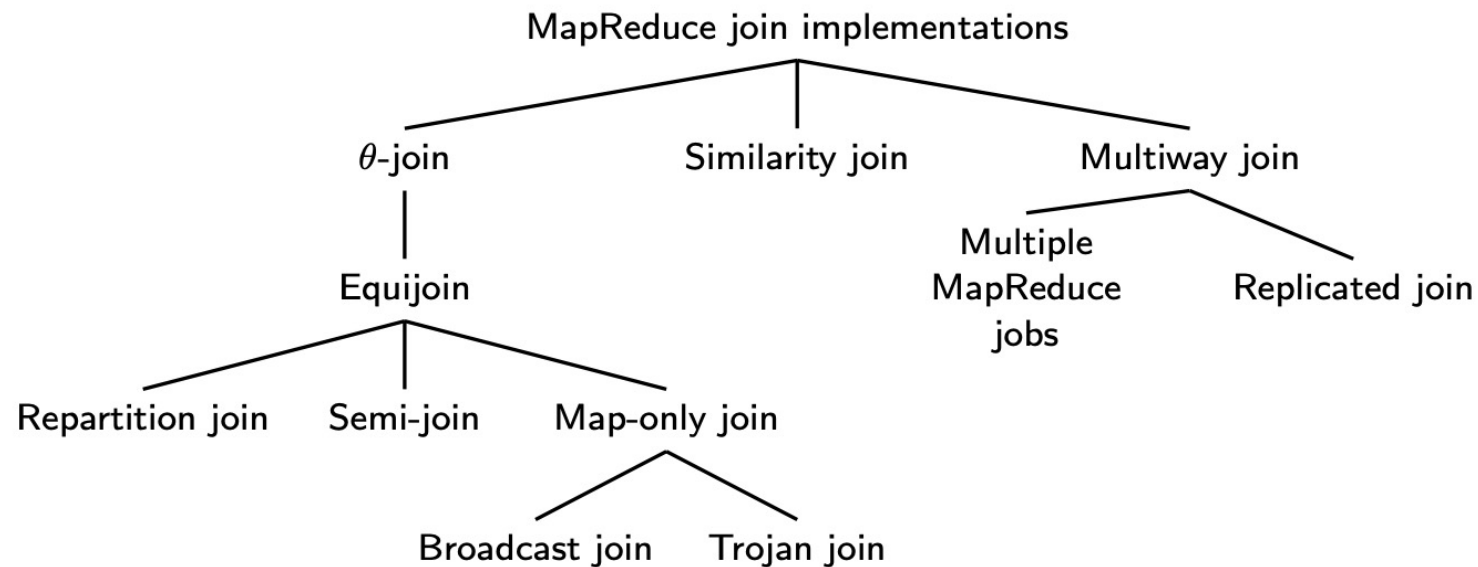


Image from M.T. Özsu & P. Valduriez (2020).
Principles of Distributed Database Systems



Spark Joins

- **SortMergeJoin**, ShuffleHashJoin, and BroadcastHashJoin
- SortMergejoin is composed of 2 steps
 - Sort the datasets
 - Merge the sorted data in the partition by iterating over the elements and according to the join key join the rows having the same value.
- BroadcastHashJoin
 - Optimum performance can be achieved
 - Strict limitations with the size of data frames
 - `spark.sql.autoBroadcastJoinThreshold=10MB`
 - Solves uneven sharding and limited parallelism
- ShuffleHashJoin
 - MapReduce based, similar to natural join



Catalyst

- Goals
 - Optimize logical plan
 - Convert logical to physical plan
 - Optimize physical plan
 - Code generation
- Scala language features
 - Pattern matching
 - Quasiquotes
 - Abstract syntax tree
 - Tree manipulation library
 - Optimizations rules implemented as tree transformations



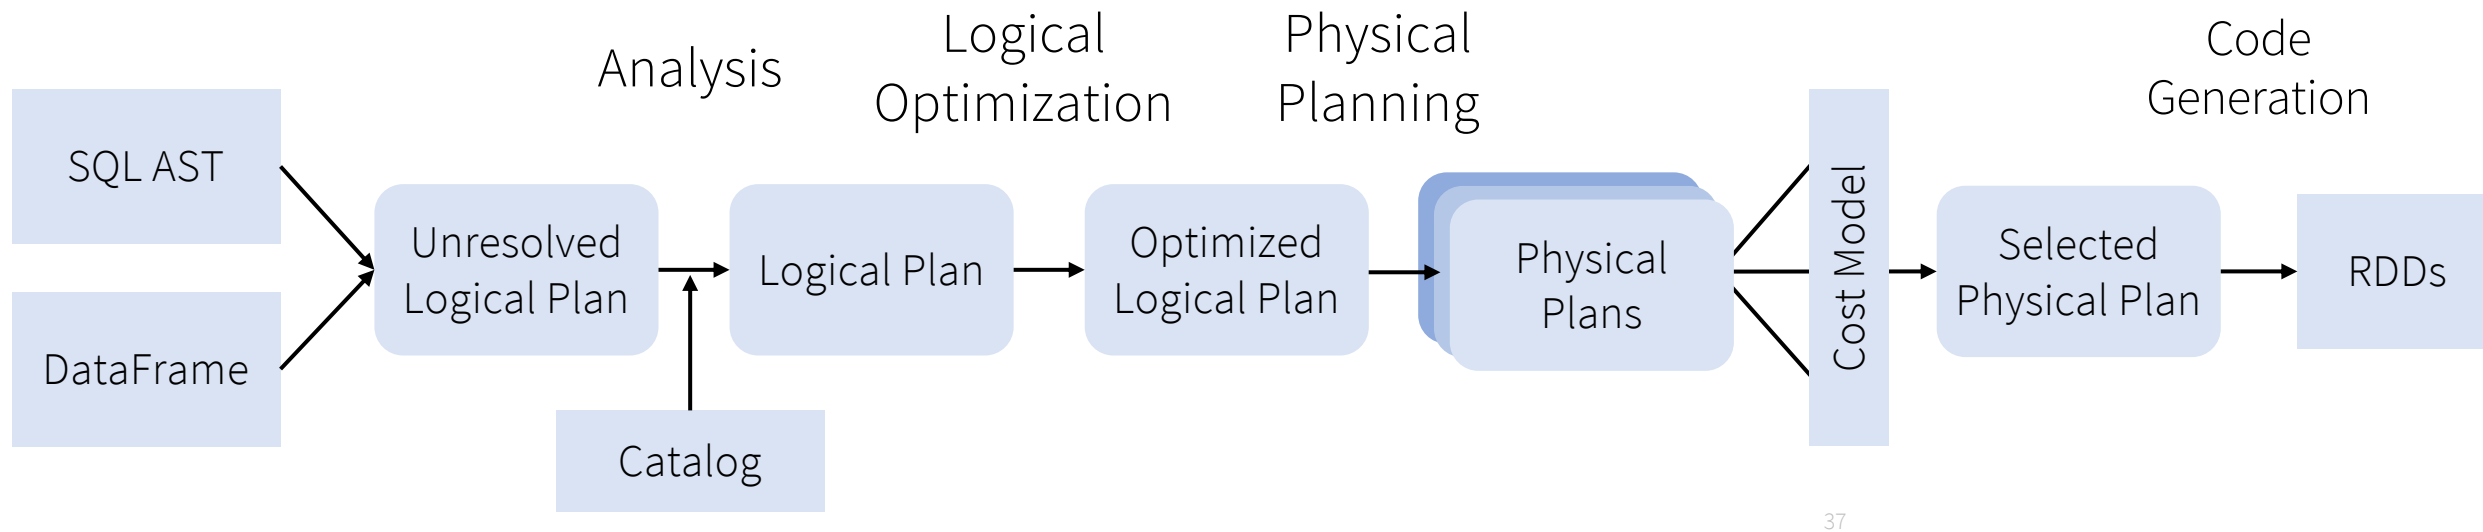
Extensibility

- Easily add new optimization techniques and features
- Enable external developers to extend the optimizer
 - e.g. adding data source specific rules, support for new data types, etc.
 - Data sources, E.g. CSV, Avro, Parquet, JDBC
 - Map user-defined types to structures composed of Catalyst's built-in types.

```
class PointUDT extends UserDefinedType[Point] {  
  def dataType = StructType(Seq( // Our native structure  
    StructField("x", DoubleType),  
    StructField("y", DoubleType)  
  ))  
  def serialize(p: Point) = Row(p.x, p.y)  
  def deserialize(r: Row) =  
    Point(r.getDouble(0), r.getDouble(1))  
}
```



Plan Optimization & Execution

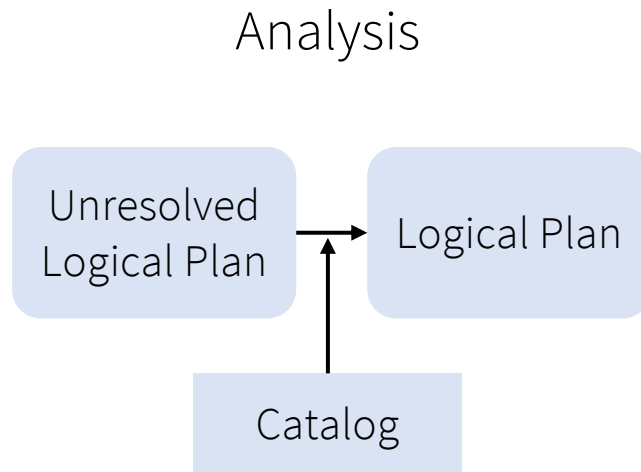


DataFrames and SQL share the same optimization/execution pipeline

Image from DataBricks SparkSQL presentation@SIGMOD 2015



Analysis



- An attribute is *unresolved* if its type is not known or it's not matched to an input table.
- To resolve attributes:
 - Look up relations by name from the catalog.
 - Map named attributes to the input provided given operator's children.
 - UID for references to the same value
 - Propagate and coerce types through expressions (e.g. $1 + col$)

Image from DataBricks SparkSQL presentation@SIGMOD 2015



Logical Optimization

- Applies standard rule-based optimization (constant folding, predicate-pushdown, projection pruning, null propagation, boolean expression simplification, etc)

Logical
Optimization





Image from DataBricks SparkSQL presentation@SIGMOD 2015

```
object DecimalAggregates extends Rule[LogicalPlan] {  
  /** Maximum number of decimal digits in a Long */  
  val MAX_LONG_DIGITS = 18  
  
  def apply(plan: LogicalPlan): LogicalPlan = {  
    plan transformAllExpressions {  
      case Sum(e @ DecimalType.Expression(prec, scale))  
        if prec + 10 <= MAX_LONG_DIGITS =>  
        MakeDecimal(Sum(LongValue(e)), prec + 10, scale)  
    }  
  }  
}
```




Physical Planning

Logical Plan



Physical Plan



Physical Plan
with Predicate Pushdown
and Column Pruning




Image from DataBricks SparkSQL presentation@SIGMOD 2015



Code Generation

```
def compile(node: Node): AST = node match {  
  case Literal(value) => q"$value"  
  case Attribute(name) => q"row.get($name)"  
  case Add(left, right) =>  
    q"${compile(left)} + ${compile(right)}"  
}
```




Figure 4: A comparison of the performance evaluating the expression $x+x+x$, where x is an integer, 1 billion times.

- Relies on Scala's quasiquotes to simplify code gen.
- Generating Java bytecode to run on each machine
- Whole-Stage CodeGen
 - Joins multiple physical operations together to form a single Java function
 - Leverages CPU registers for intermediate data

Images from DataBricks SparkSQL presentation@SIGMOD 2015



More information

- [Mining of Massive Datasets](#), Anand Rajaraman and Jeff Ullman, Cambridge University Press, Section 2.3
- [SparkSQL: Relational Data Processing in Spark](#), M. Armbrust, et al., SIGMOD, 2015

