

# Distributed Systems Paradigms

José Orlando Pereira

Departamento de Informática  
Universidade do Minho



# Motivation

- Handle a large number of clients and requests with a single server
- The “c10k problem” in 1999:
  - <http://www.kegel.com/c10k.html>
- Examples:
  - financial, games, ...
  - notifications in mobile apps
  - machine-to-machine (M2M)

# Case study

- Simple chat server:
  - Forward all messages to all clients
- Consider:
  - Large number of clients
  - Slow connections



# First threaded solution

- For each connection:
  - Handler thread
- When reading, write to all other connections
- Use buffering:
  - At user level (streams): To minimize system calls
  - In the kernel (socket): To cope with slow readers


# Sockets in java.net

```
ServerSocket ss=new ServerSocket(12345);  
  
while(true) {  
    Socket s=ss.accept();  
  
    // start handler thread  
  
    s.close();  
}
```

# Buffers in java.net

```
InputStream is=new BufferedInputStream(s.getInputStream());
OutputStream os=new BufferedOutputStream(s.getOutputStream());
while(true) {
    // i/o
    is.read(...)
    for(var os: connections) {
        os.write(...);
    }
}
```

# First threaded solution




# Buffers in java.net

```
InputStream is=new BufferedInputStream(s.getInputStream());
OutputStream os=new BufferedOutputStream(s.getOutputStream());
while(true) {
    // i/o
    is.read(...)
    for(var os: connections) {
        os.write(...);
        os.flush();
    }
}
```

What if flush blocks?



# Writer thread




# Memory

- Memory:  $n$  connections x messages in transit ( $\sim n^2$ )
  - Caused by data copying in stacked abstractions
    - Serialization!
  - Overhead in allocation and garbage collection
- Solution: Deal directly with byte buffers


# Buffers in java.nio

- Buffer = Array + Indexes:



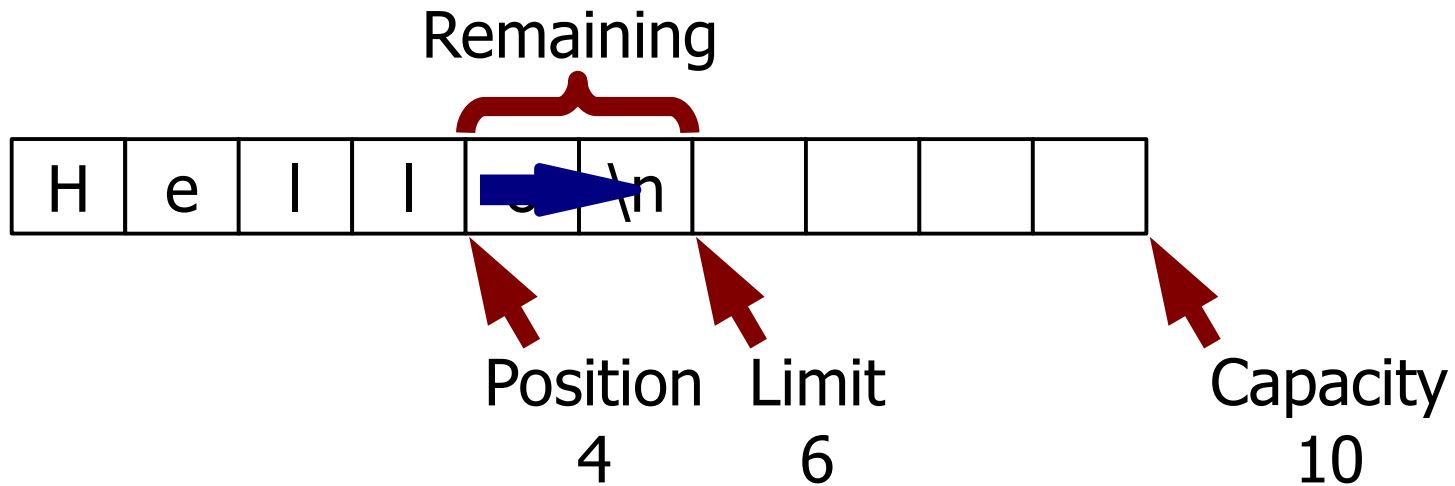
# Buffers in java.nio

- Put/read: advances position, sets content




# Buffers in java.nio

- Get/write: advances position, gets content

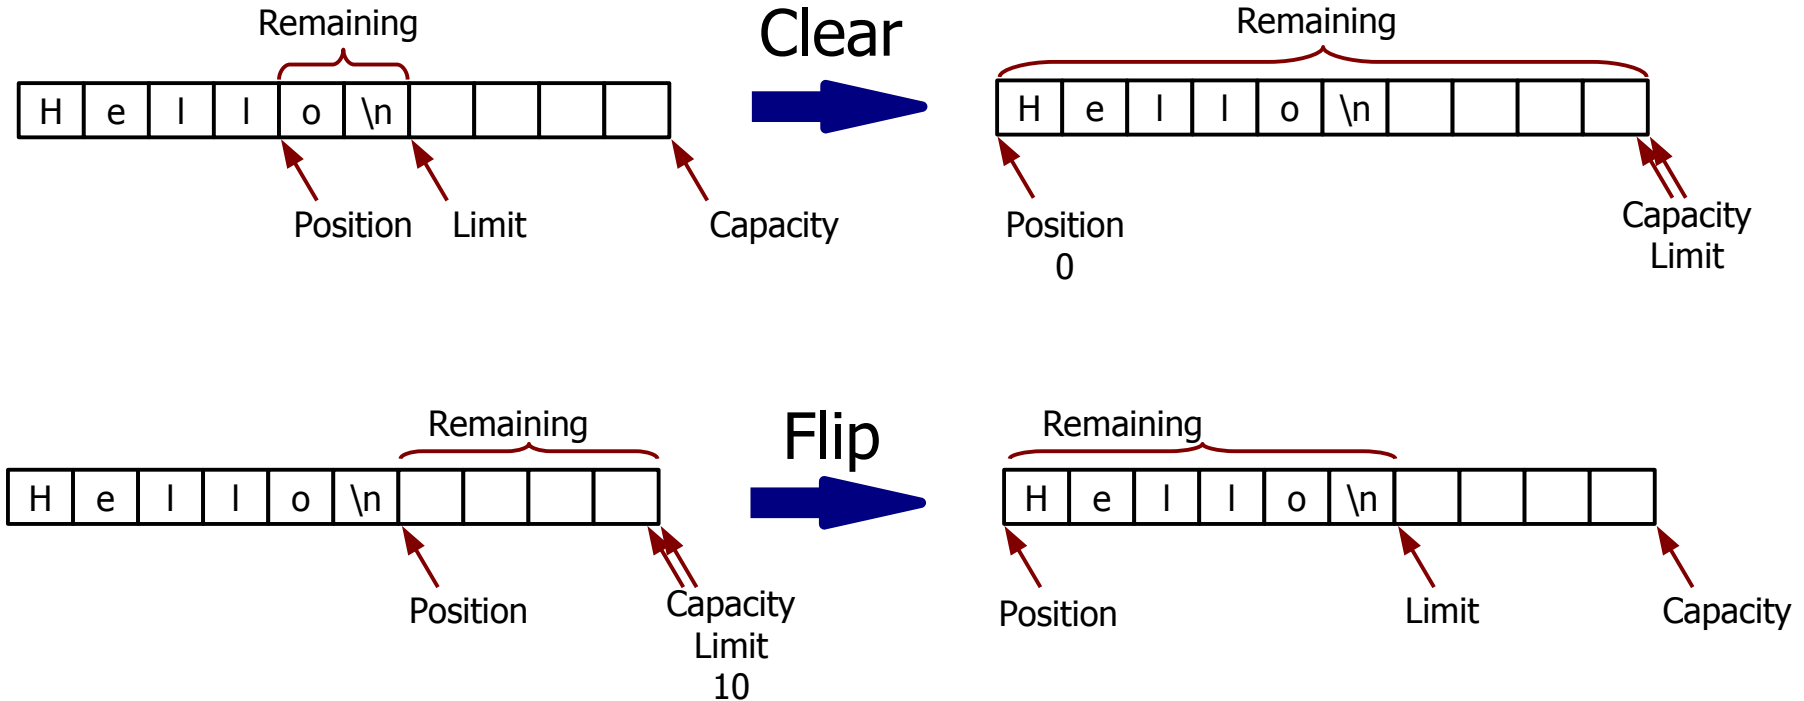


# Buffers in java.nio

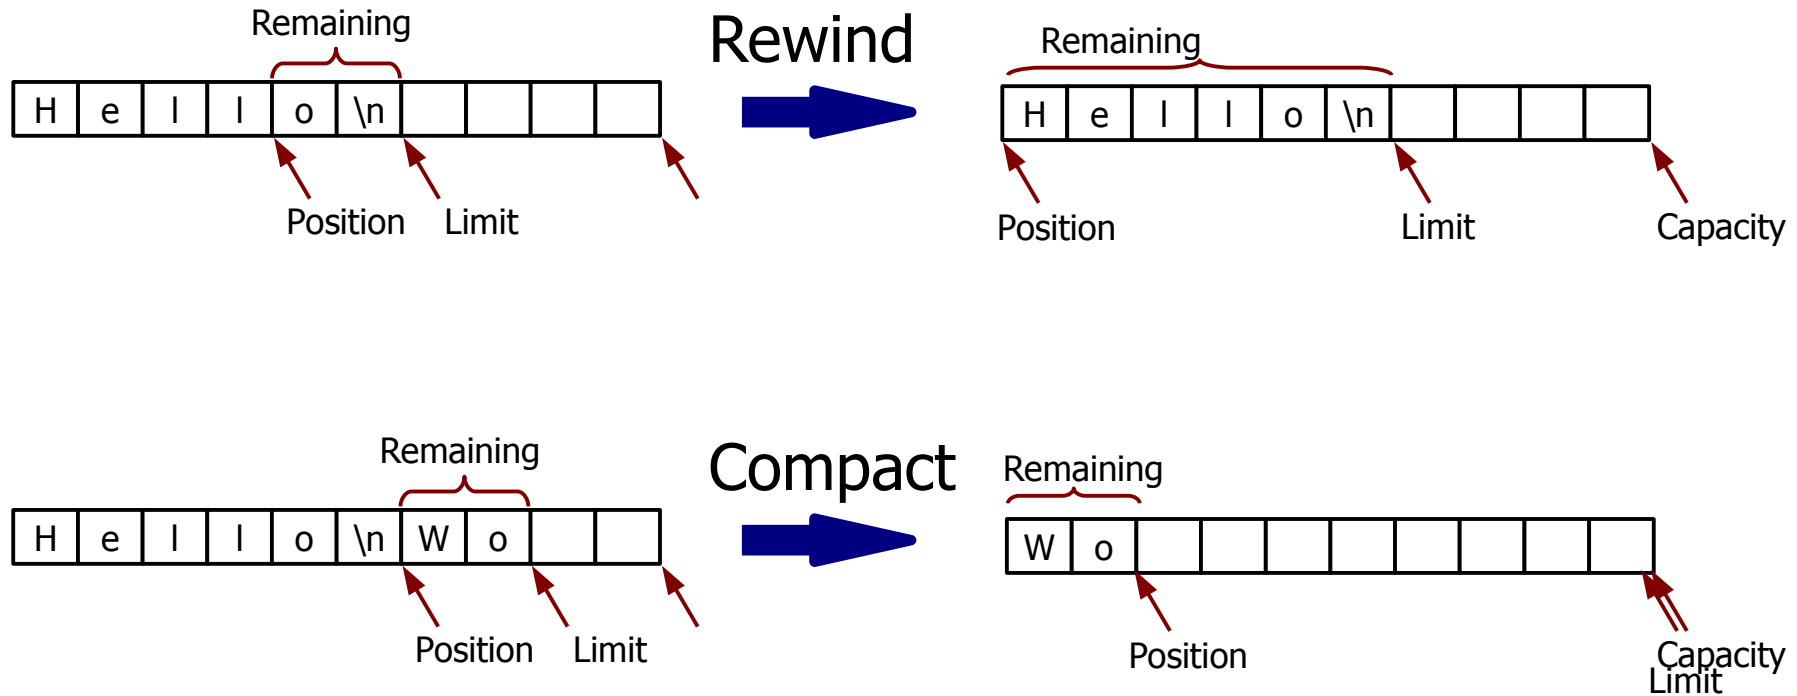
- Duplicate: multiple pointers into the same array



# Buffers in java.nio



# Buffers in java.nio





# Serialization

- Strings need to be encoded to UTF-8, ASCII, ... and decoded back when received

```
Charset utf8 = StandardCharsets.UTF8;  
String s = ...  
ByteBuffer b = ...  
  
b = utf8.encode(s);  
  
s = utf8.decode(b);
```

# Serialization

- Primitive types can be translated to and from bytes with `getInt/putInt`, `getFloat/putFloat`, etc...
- They are represented according to the current byte order of the byte buffer
- The byte order is “big endian” by default but can be changed

# Sockets in java.nio

```
ServerSocketChannel ss=ServerSocketChannel.open();  
ss.bind(new InetSocketAddress(12345));  
  
while(true) {  
    SocketChannel s=ss.accept();  
  
    // start i/o threads  
}
```

# Sockets in java.nio

```
ByteBuffer buf=ByteBuffer.allocate(100);
```

```
s.read(buf);  
buf.flip();
```

```
for(SocketChannel r: receivers) {  
    r.write(buf.duplicate());  
}  
buf.clear();
```



Use `write(ByteBuffer[] src)`  
to avoid copying!


# Shared buffers

- Memory used: messages in transit ( $\sim n$ )
- Ideally, never allocate or dispose of memory in normal operation:
  - No overhead, but...
  - Needs reference counting to know when to reuse

# Second threaded solution

- Keep a shared queue
- For each connection:
  - Reader thread + Pointer into queue + Writer thread
- When reading, insert in shared queues and notify writer threads
- When writing, advance pointer
- Lazily, remove prefix from shared queue

# Server architecture



# Threads summary

- Simple programming model
- Problems:
  - Memory overhead (stacks and buffers)
  - Context switches and lock contention
  - “Thundering herd”, hidden queues, and fairness