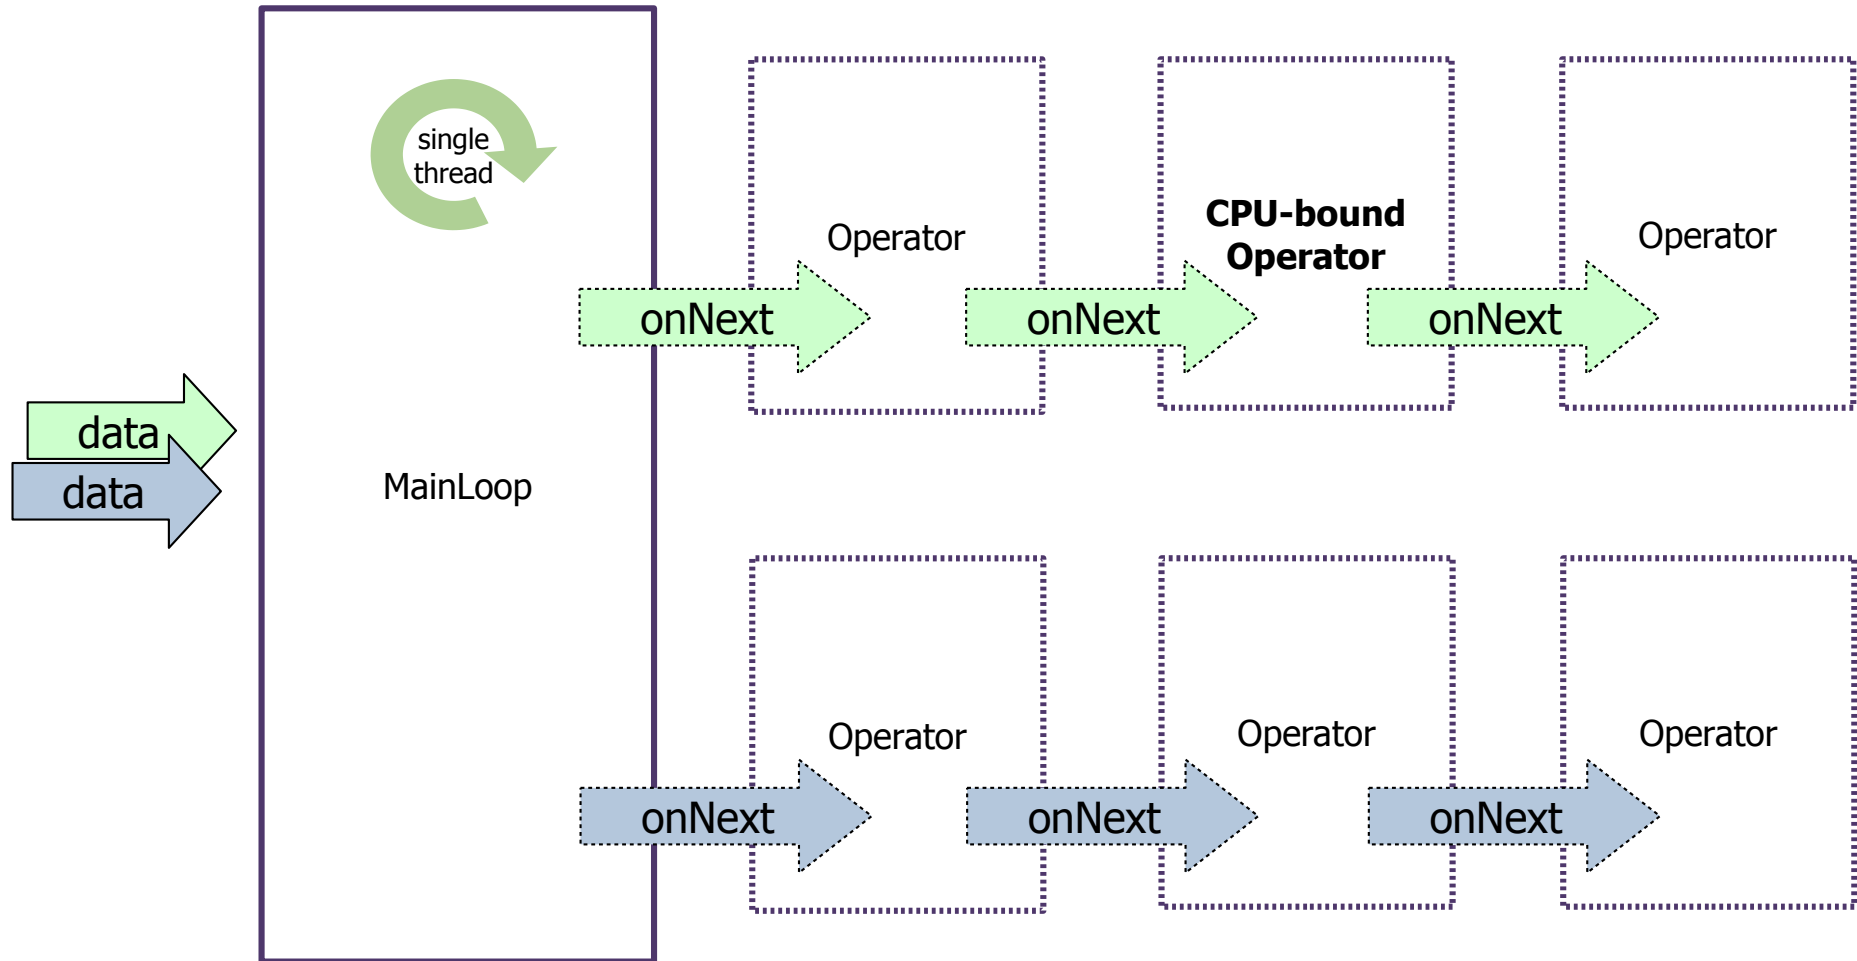


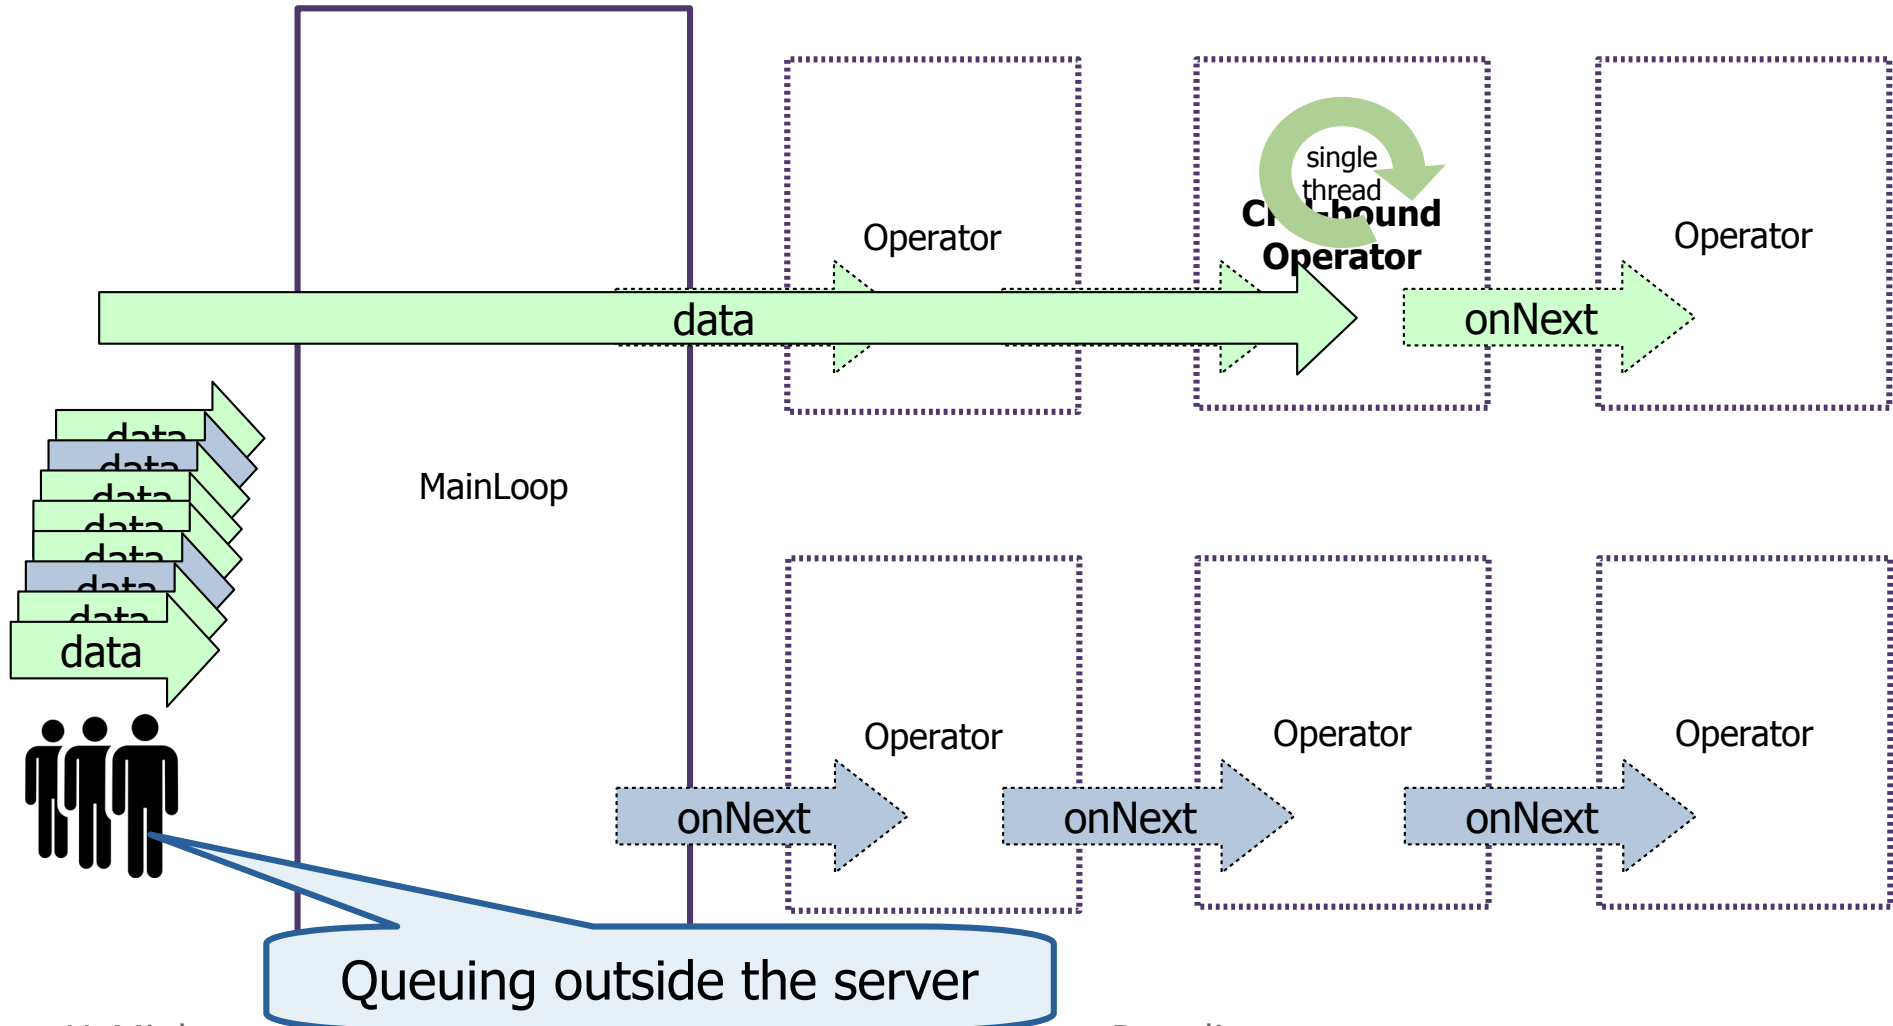
# Challenges

- Using a single thread avoids the complications of concurrent programming (races)
- What if a single CPU core is not enough?
- What if a single stream is using more than its share of resources?

# Buffer-based application



# Buffer-based application



# Naive parallel operator

```
public class NaiveParallelOperator implements ObservableOperator<T,R> {
```

```
    public Observer<...> apply(Observer<...> child) throws T {
```

```
        return new Observer<...>() {
```

```
            public void onNext(...) {
```

```
                new Thread(()->{
```

```
                    child.onNext(...);
```

```
                }).start();
```

```
            }
```

```
            public void onError(Throwable e) {
```

```
                ...
```

```
            }
```

```
            public void onComplete() {
```

```
                ...
```

```
            }
```

```
        }
```

```
    }
```

**Concurrent invocations  
of onNext()!!!**



# Naive parallel operator

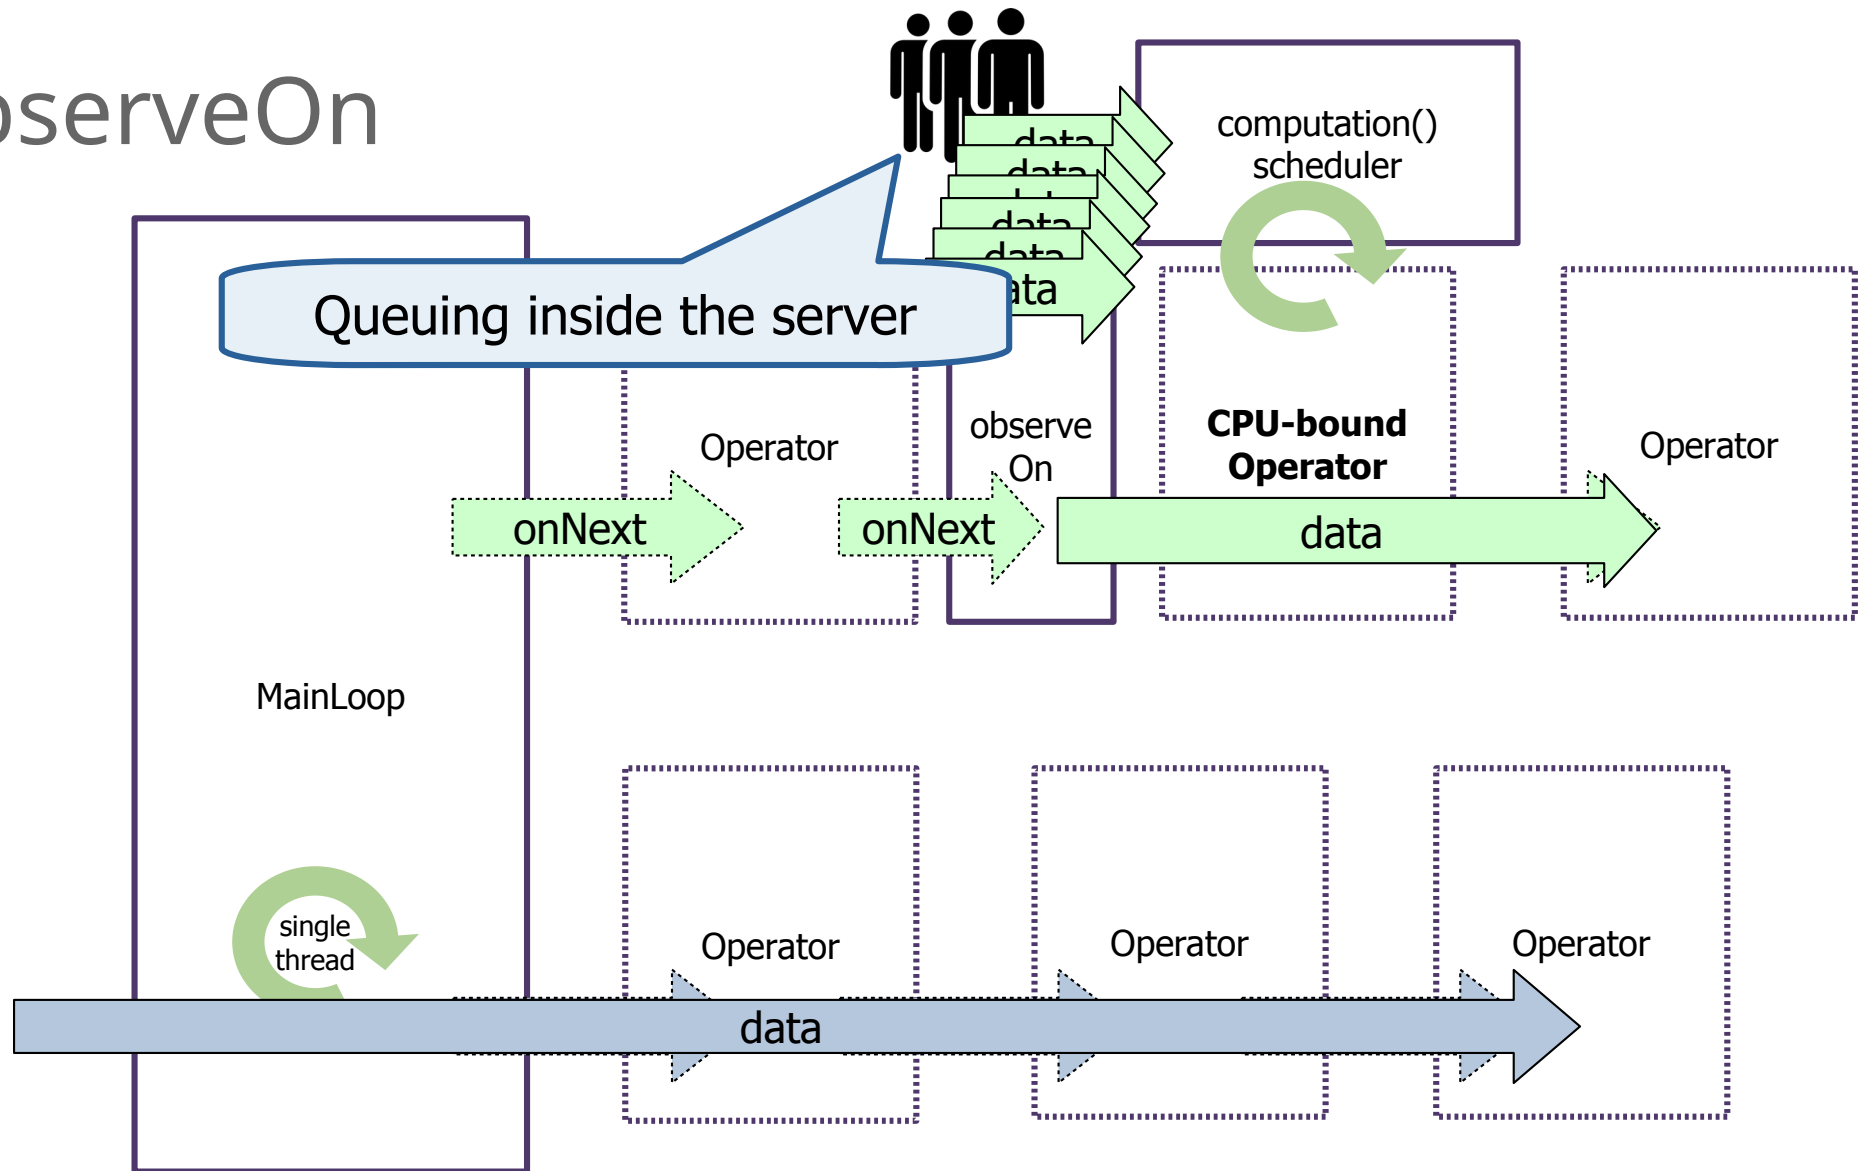
- Frees the calling thread to execute other requests
- Has many issues:
  - Concurrent invocations of `onNext` are not allowed
    - Races, out of order processing, ...
  - Synchronization of `onNext` with `onComplete/onError`
  - Does not limit the maximum number of threads
    - `#threads > #cores` is inefficient
    - potential resource exhaustion

# observeOn

- observeOn(...) operator redirects execution to a different thread, but ensures synchronization
- Threads are provided by a scheduler that manages threads according to a policy:
  - computation(), io(), ...

```
server.subscribe(conn -> {  
    conn  
        .op1(...)  
        .observeOn(computation())  
        .op2(...)  
        .subscribe(...);  
});
```

# observeOn




# observeOn

- Frees the main loop thread to process other requests
- Can be used to implement the parallel pipeline pattern:
  - Use observeOn multiple times in the same stream to delimit stages
  - Multiple stages execute in parallel
- Main problem: the waiting queue is now inside the server, consuming memory
- Limiting the queue blocks the main loop again

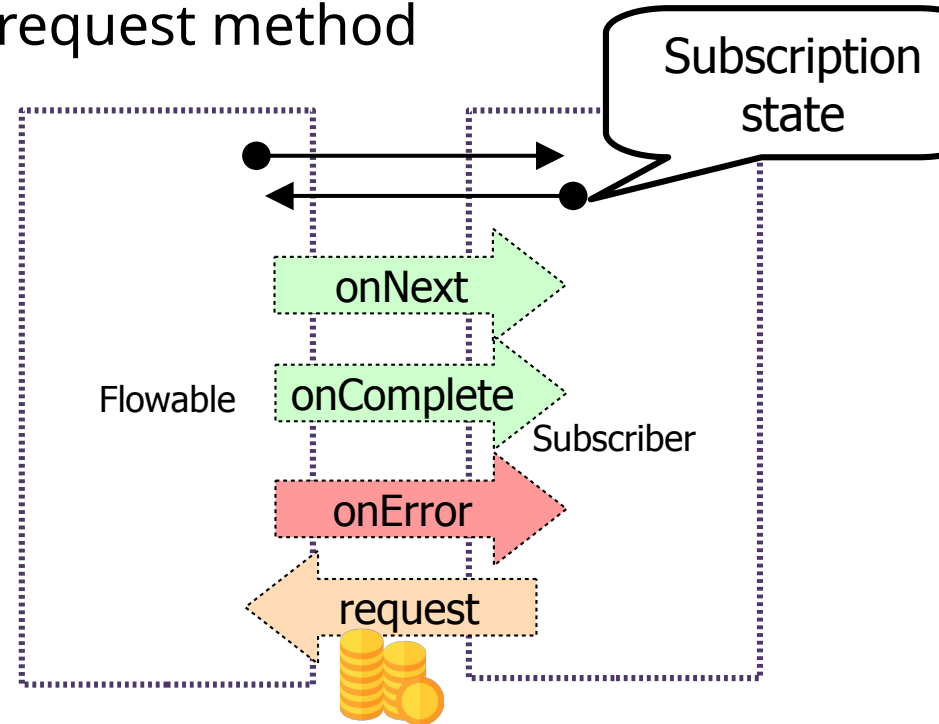


# Back pressure

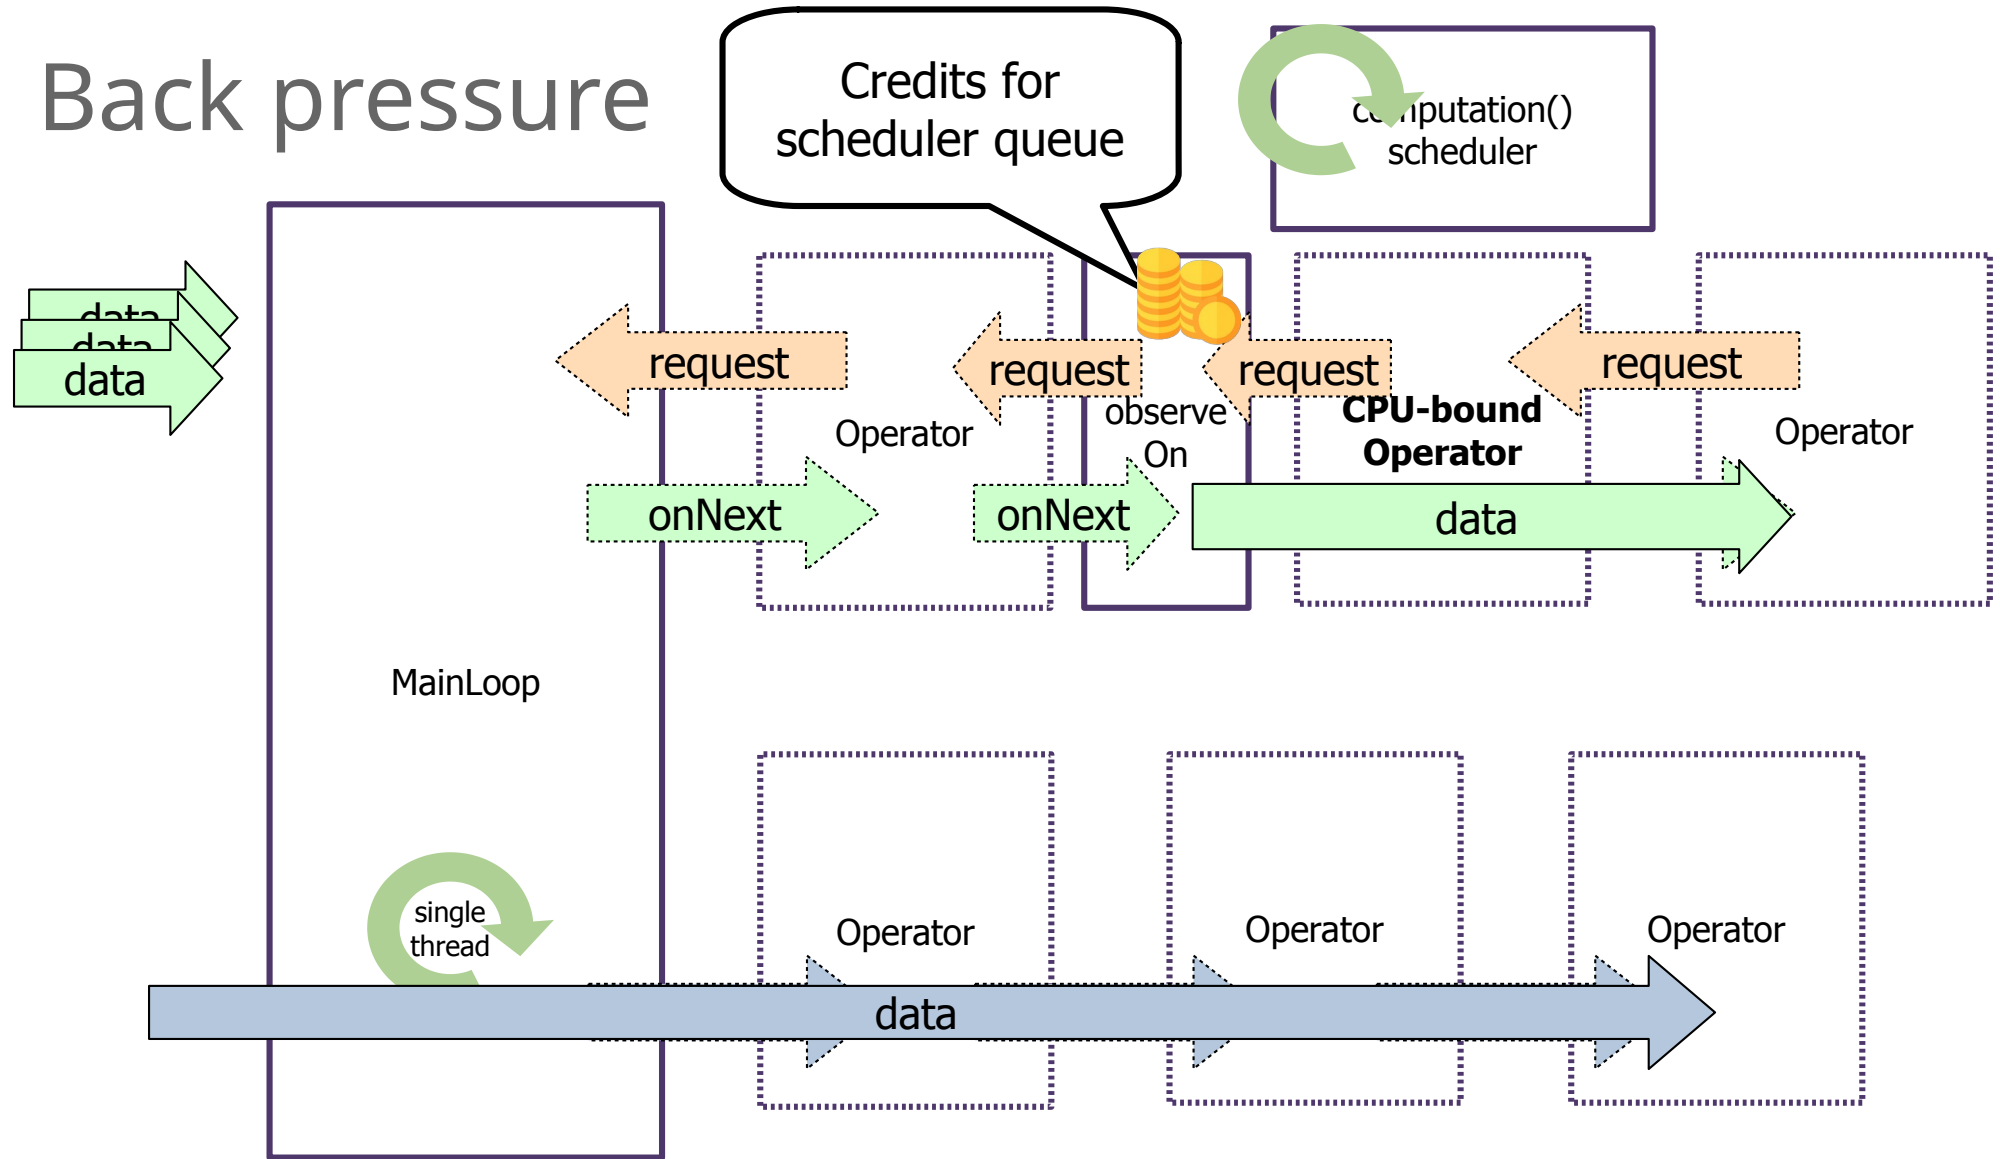
- The observer cannot influence the observable
  - No reference back and no interface



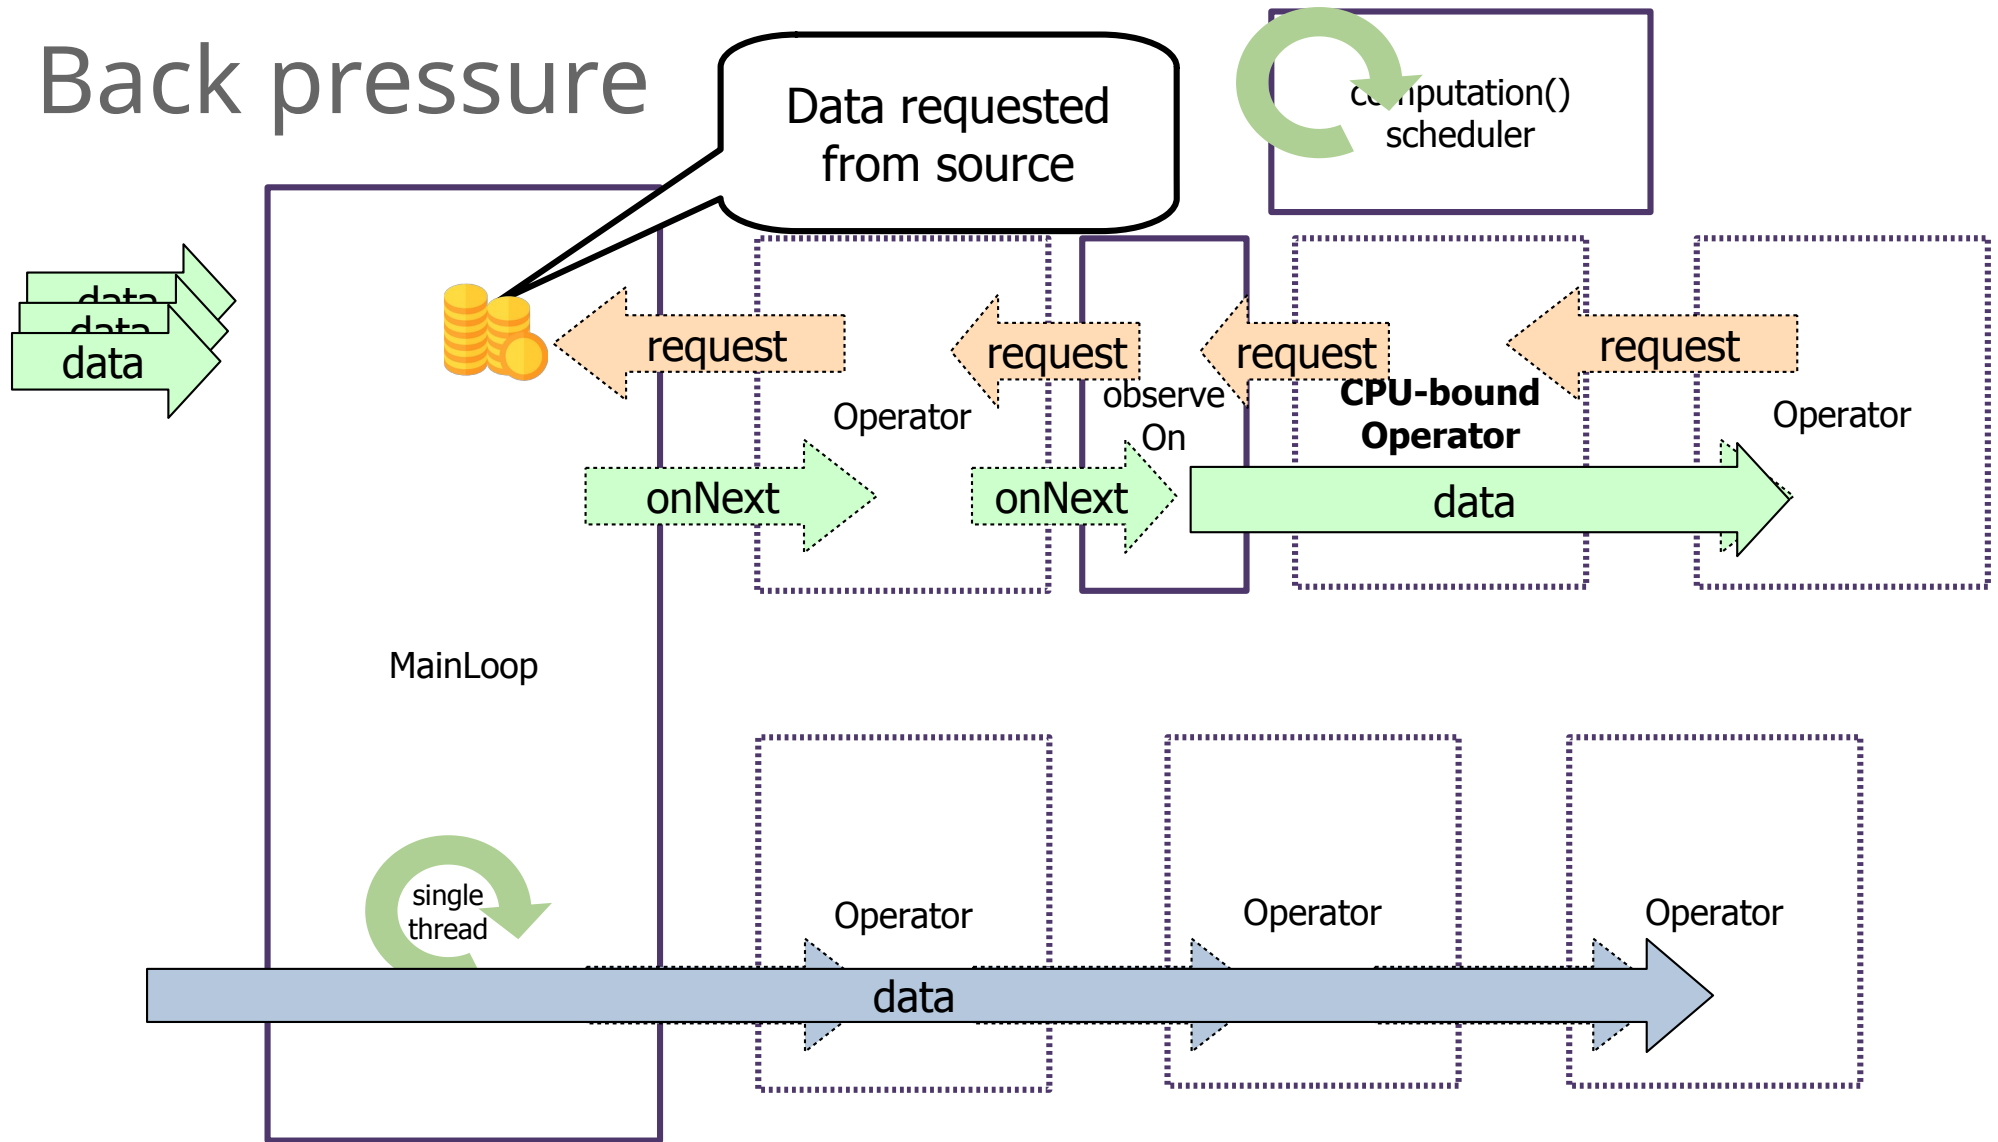
- Flowable streams can only invoke **onNext** when there is credit left
- Credit is provided by invoking **request** method



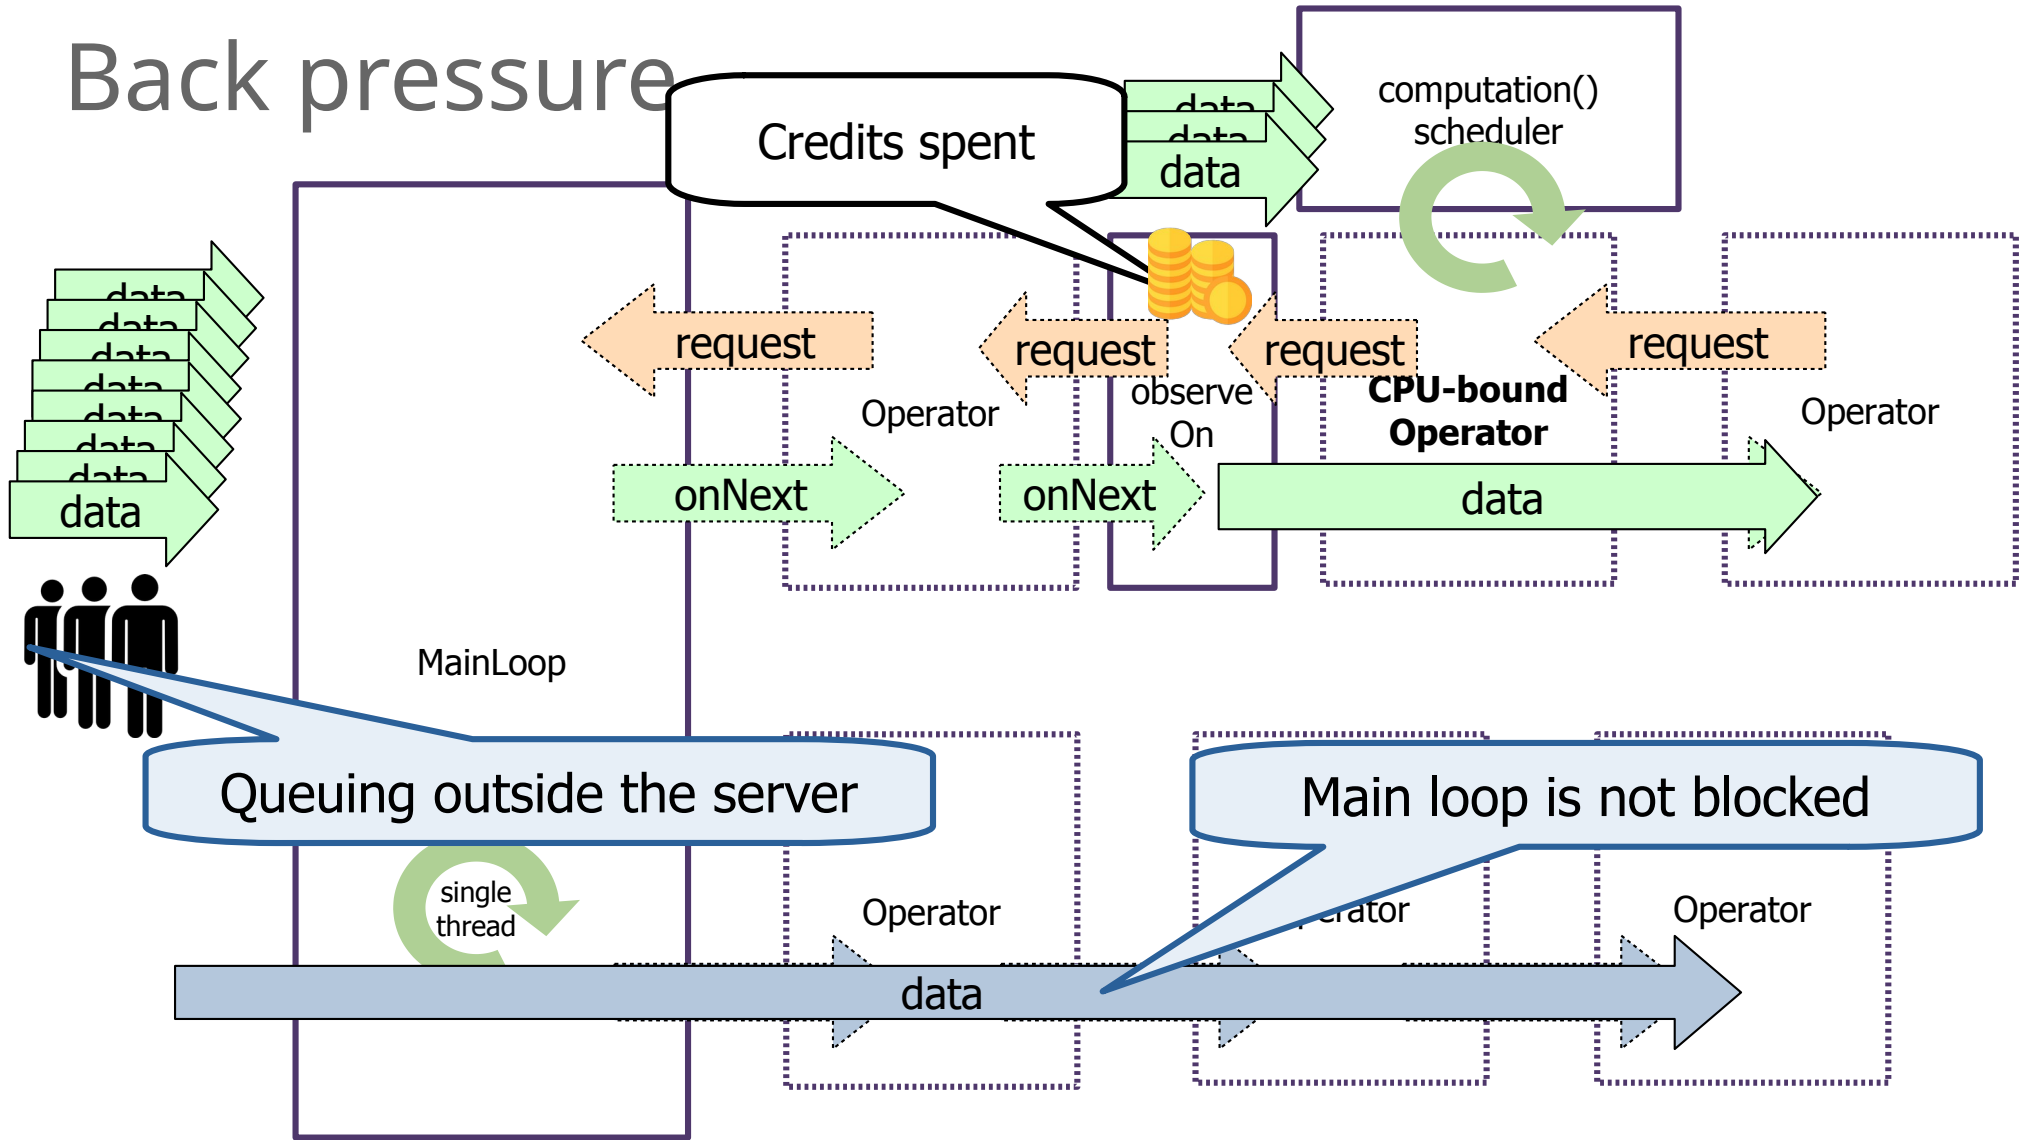
# Back pressure



# Back pressure



# Back pressure



# Subscribing flowables

- Sample subscription:

```
f.subscribe(new DefaultSubscriber<...>() {  
    public void onStart() {  
        request(1);  
    }  
    public void onNext(...) {  
        ...  
        request(1);  
    }  
    ...  
});
```

Initial credits

Refill when needed

- Subscription with a simple callback provides infinite credits when subscribing
  - Useful when the bottleneck is in the middle of the pipeline

# Subscribing flowables

```
public class Mainloop {
    public void write(Flowable<ByteBuffer> flow, SocketChannel s) {
        flow.subscribe(new DefaultSubscriber<ByteBuffer>() {
            public void onStart() {
                request(1);
            }
            public void onNext(ByteBuffer bb) {
                s.write(bb);
                if (bb.hasRemaining()) {
                    key.interestOpsOr(OP_WRITE);
                } else {
                    request(1);
                }
            }
        });
    }
}
```

(\*) Implement “Publisher” for compatibility with other frameworks.


# Subscribing flowables

```
public class Mainloop {  
    public run() {  
        ...  
        if (key.isWritable()) {  
            if (there is a pending write) {  
                s.write(bb);  
            } else {  
                flow.request(1);  
                key.interestOpsAnd(~OP_WRITE);  
            }  
        }  
    }  
}
```



Refill credits

# Complete client-server application





```
public class Server {  
    public static void main(String[] args) throws Exception {  
        var ss = ServerSocketChannel.open(new InetSocketAddress(12345));  
        var loop = new MainLoop();  
        var ss_obs = loop.accept(ss);  
        ss_obs.subscribe(s -> {  
            var s_flow = loop.read(conn)  
                .observeOn(computation())  
                .lift(new LineSplitOperator())  
                .map(bb -> StandardCharsets.UTF_8.decode(bb))  
                .map(String::toUpperCase)  
                .map(s -> StandardCharsets.UTF_8.encode(s));  
            loop.write(s_flow, conn);  
        });  
    }  
}
```

Business logic



# Combining flowable and observable

- Convert flowable to observable:
  - `f.toObservable().op(...).subscribe();`
  - subscribes with infinite credit
- Observable to flowable needs to deal with back pressure:
  - `o.toFlowable(strategy).op(...).subscribe(...);`
  - multiple strategies available, for example:




# Implementing flowables

```
public class Mainloop {  
    public Flowable<ByteBuffer> read(SocketChannel s) {  
        return subscriber -> {  
            s.configureBlocking(false);  
            var key = s.register(sel, 0);  
            var sub = new ReadSubscription(key, subscriber)  
            key.attach(sub);  
            subscriber.onSubscribe(subscription);  
        }  
    }  
    public run() {  
        ...  
        if (key.isReadable()) {  
            var sub = (MySubscription) k.attachment();  
            ...  
            subscription.subscriber.onNext(bb);  
            if (sub.credits == 0) key.interestOpsAnd(~OP_READ);  
            ...  
        }  
    }  
}
```

No OP\_READ until we  
have credits

No OP\_READ if we  
exhausted credits

# Implementing flowables

```
public class Mainloop {  
    private class ReadSubscription implements Subscription {  
         private long credits = 0;  
  
        public void request(long n) {  
            credits += n;  
            if (not currently reading) {  
                key.interestOpsOr(OP_READ);  
                key.selector().wakeup();  
            }  
        }  
    }  
}
```



needs thread  
synchronization

Restore OP\_READ

Make the selector  
consider new interests

# Custom operator



```
public class LineSplitOperator implements FlowableOperator<ByteBuffer,ByteBuffer> {  
    public Subscriber<...> apply(Subscriber<...> child) throws Throwable {  
        return new Subscriber<ByteBuffer>() {  
            public void onSubscribe(Subscription parent) {  
                this.parent = parent;  
                child.onSubscribe(new Subscription() {  
                    public void request(long l) {  
                        ..  
                    }  
                });  
            }  
            public void onNext(ByteBuffer bb) {  
                ...  
                if (have downstream credits)  
                    child.onNext(...);  
                if (can receive more data)  
                    parent.request(...);  
            }  
        }  
    }  
}
```

The diagram illustrates the credit management logic in the `onSubscribe` and `onNext` methods of the `LineSplitOperator`. Annotations include:

- Store reference to give credits**: Points to `this.parent = parent;` in the `onSubscribe` method.
- Receive credits from child**: Points to the `..` (ellipsis) in the `request` method of the inner `Subscription` class, which is enclosed in a red warning triangle.
- Consume credits from child**: Points to `child.onNext(...);` in the `onNext` method, associated with the condition *if (have downstream credits)*.
- Give credits to parent**: Points to `parent.request(...);` in the `onNext` method, associated with the condition *if (can receive more data)*.

# Parallelism

- In addition to parallel pipeline with `observeOn()`, back pressure allows parallel processing within one stage:



# References

- ReactiveX / RxJava documentation: <https://reactivex.io/>
- Tomasz Nurkiewicz and Ben Christensen. *Reactive Programming with RxJava: Creating Asynchronous, Event-Based Applications*. O'Reilly, 2017.
  - Chap. 6
  - *Warning*: Uses an older version of the RxJava API.
- For Python: <https://github.com/MichaelSchneeberger/rxbackpressure>