

Consistency

Paulo Sérgio Almeida

February 2023



Universidade do Minho

Abstractions


- ▶ In a process we can have abstract sequential datatypes
- ▶ In a distributed system we can also have abstract datatypes
 - registers (read, write)
 - counters (inc)
 - sets (add, remove)
- ▶ But what can we count on using such abstractions?
- ▶ Many possible design choices and tradeoffs
- ▶ Motivated by distribution, architecture, goals
- ▶ Looking at some architectures helps understanding the issue

Architectures

What is a consistency model


Some consistency models

Single server




- ▶ A single server, used by all clients
- ▶ Easy to emulate sequential datatype: server serves in sequence
- ▶ Single server a single point of failure (SPOF)
- ▶ Does not scale to many clients

Multiple replicated servers



- ▶ Multiple servers run symmetrical distributed algorithm
- ▶ Each replicates part or the whole data
- ▶ Each client picks one server, and possibly change server
- ▶ Can possibly scale to many clients and avoid a SPOF
- ▶ Costly to emulate a sequential datatype

Pure peer-to-peer system



- ▶ There are no “servers”
- ▶ Each “client” node maintains full replica of state
- ▶ Each node talks directly to (possibly a subset of) others
- ▶ Can possibly scale to many clients and avoid a SPOF
- ▶ Costly to emulate a sequential datatype

Architectures

What is a consistency model

Some consistency models

An ideal world vs actual possible outcomes

- ▶ Everything would be easy if operations
 - executed instantly
 - had effects propagated globally, instantly
- ▶ The above is physically impossible
- ▶ Not relevant only for distributed systems
- ▶ Example: in a Java program, two threads reading and writing global variables, initially $a = b = 0$, without synchronization

(thread 1)	(thread 2)
a = 1;	b = 1;
x = b;	y = a;

- ▶ Is outcome $x = 0$ and $y = 0$ possible?
 - cannot be explained by “things happening in some order”
 - we would prefer that to be impossible – for reasoning
 - why would JVM designers make it possible?
- ▶ A consistency model describes possible outcomes

Why relax consistency in distributed systems

- ▶ In a distributed system (remember example architectures)
 - messages take time to propagate
 - data may be replicated in several nodes
 - replicas receive updates concurrently
 - there can be network partitions
- ▶ If we try to emulate sequential datatype semantics
 - must slow everything down, to avoid replicas diverging
 - network partitions may cause unavailability (no response)
- ▶ Practically, we must make some tradeoffs, relaxing consistency
 - must allow more possible, strange, outcomes
 - to allow better performance and availability

Processes, operations, local order

- ▶ Each process issues operations
 - operations of each process are totally ordered
 - process is a logical entity; can be thread, OS process, node
 - we assume, reasonably, a finite number of processes
- ▶ Session order (so):
 - session: operations from the same process
 - $a \xrightarrow{\text{so}} b$ if a is issued before b , by the same process
 - session order is itself a partial order
- ▶ Operations are broadly divided into
 - writes: update operations, that change state, may also read it
 - reads: query operations, only read

Global Orders

► Visibility (vis)

- $a \xrightarrow{\text{vis}} b$ if effect of a visible to process performing b
- example: read b seeing value written by write a
- characterizes how updates are propagated and seen

► Happens-before (hb)

- $a \xrightarrow{\text{hb}} b$ if a is in the potential causal past of b
- analogous to classic one, for observed behavior; “semantic hb”
- combines session order and visibility

$$\text{hb} \doteq (\text{so} \cup \text{vis})^+$$

► Arbitration (ar)

- total order on all operations
- explains how conflicts are resolved; better if not needed
- should at least be compatible with visibility

$$\text{vis} \subseteq \text{ar}$$

- preferably, we want *causal arbitration*, compatible with hb

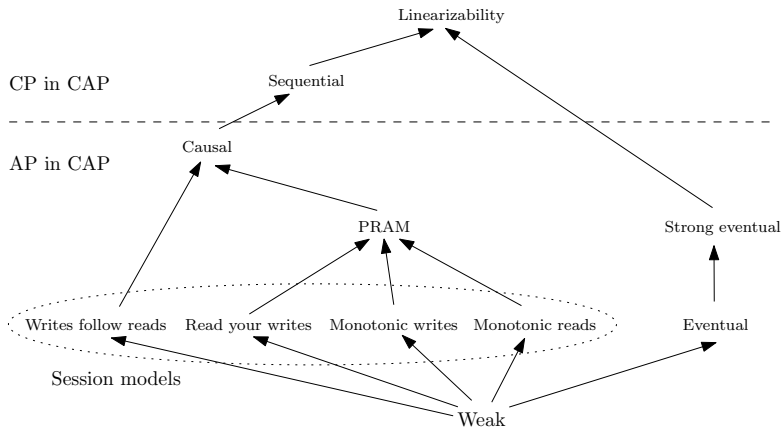
$$\text{hb} \subseteq \text{ar}$$

Architectures

What is a consistency model

Some consistency models

A taxonomy of some consistency models



Session models

- ▶ Basic desirable guarantees for weakly replicated data
- ▶ Each process, individually, should “see and propagate progress”
- ▶ Otherwise, very weird things will happen
- ▶ The 4 sessions guarantees are
 - RYW - read your writes
 - MR - monotonic reads
 - MW - monotonic writes
 - WFR - writes follow reads

Session models – read your writes

- ▶ RYW: reads reflect previous writes from the same process

$$\text{RYW} \doteq \forall o, \text{ write } w \cdot w \xrightarrow{\text{so}} o \Rightarrow w \xrightarrow{\text{vis}} o$$

- ▶ If it doesn't hold, it confuses clients/applications
 - a process updates a database
 - it then reads from database
 - the update is missing
- ▶ Reads restricted to replicas that include previous own writes
- ▶ Easy to hold with *server affinity*
 - if client process can stick to same server

Session models – monotonic reads

- ▶ MR: successive reads see at least the same writes

$$\text{MR} \doteq \forall o, o', \text{write } w \cdot w \xrightarrow{\text{vis}} o \wedge o \xrightarrow{\text{so}} o' \Rightarrow w \xrightarrow{\text{vis}} o'$$

- ▶ Example, with MR
 - a distributed database maps keys to records
 - records are inserted or updated by writes, never removed
 - a read returns a record for a given key
 - a subsequent read always finds that key present
 - it does not matter who is doing writes

Session models – monotonic writes

- ▶ MW: writes by same process are propagated in session order

$$\text{MW} \doteq \forall o, \text{writes } w, w'.$$

$$\begin{array}{l} w \xrightarrow{\text{so}} w' \wedge w' \xrightarrow{\text{vis}} o \Rightarrow \\ w \xrightarrow{\text{vis}} o \wedge w \xrightarrow{\text{ar}} w' \end{array}$$

where $w \xrightarrow{\text{ar}} w'$ is redundant in the normal case when $\text{so} \subseteq \text{ar}$

- ▶ Guarantee useful also to other processes
- ▶ Example, with MW
 - consider a distributed file system
 - a process writes a new version of a library file
 - new version upward compatible and extends features
 - the process updates executable using new library features
 - if a process reads updated executable, it reads updated library

Session models – writes follow reads

- WFR: writes after a read are propagated after observed writes

$\text{WFR} \doteq \forall o, \text{ writes } w, w', \text{ read } r.$

$$\begin{array}{c} w \xrightarrow{\text{vis}} r \wedge r \xrightarrow{\text{so}} w' \wedge w' \xrightarrow{\text{vis}} o \Rightarrow \\ w \xrightarrow{\text{vis}} o \wedge w \xrightarrow{\text{ar}} w' \end{array}$$

- Guarantee useful also to other processes
- A form of causality propagation
- Example, with WFR
 - a process reads a database record
 - notices that a field has a wrong value
 - writes the updated record
 - the updated version is propagated after original
 - replaces the previous version at all replicas

PRAM (or FIFO) - Pipelined Random Access Memory

- ▶ PRAM: writes seen in session order, own writes immediately

$$\text{PRAM} \doteq \forall o, o', \text{writes } w, w'.$$

$$\begin{aligned} & (w \xrightarrow{\text{so}} o \Rightarrow w \xrightarrow{\text{vis}} o) \wedge \\ & (w \xrightarrow{\text{vis}} o \wedge o \xrightarrow{\text{so}} o' \Rightarrow w \xrightarrow{\text{vis}} o') \wedge \\ & (w \xrightarrow{\text{so}} w' \wedge w' \xrightarrow{\text{vis}} o \Rightarrow w \xrightarrow{\text{vis}} o) \end{aligned}$$

- ▶ Equivalent to combination of RYW, MR, MW

$$\text{PRAM} \equiv \text{RYW} \wedge \text{MR} \wedge \text{MW}$$

- ▶ As if writes from each process propagate in FIFO order
- ▶ Writes from different processes may be seen in different orders
- ▶ No global total order exists

Causal consistency

- ▶ One of the most important consistency models
- ▶ Broadly speaking
 - all potential semantic causal past is available to an operation
 - it is the strongest model that remains available under partitions
- ▶ The essential aspect

$$\text{vis} = \text{hb}$$

- write w that potentially influences w' becomes visible first
- a read operation sees all writes that potentially influence it

Causal consistency – arbitration

- ▶ For some datatypes no order is needed at all
 - aggregation-like semantics, commutative and associative
 - just need to have the set of visible operations
 - example: counters, with an “increment” operation
- ▶ For most datatypes we need order
 - but not necessarily arbitration (a total order)
 - may be enough the set of maximal operations under hb
 - ▶ set of more recent concurrent operations
 - many conflict-free replicated datatypes
- ▶ In some cases we can resort to a causal arbitration ($hb \subseteq ar$)
 - example: classic registers (memory) with read/write
 - use hb-compatible “last-writer-wins” (not physical time)

Causality does not imply convergence

- ▶ All previous consistency models allow replicas diverging
- ▶ And even never converging again
- ▶ Even with causal consistency, replicas may diverge forever
- ▶ Example: causal memory
 - two processes concurrently write $x = 1$ and $x = 2$
 - one process may read forever x as 1 and another x as 2
- ▶ Something more is desirable using a weak consistency model
 - convergence

Eventual visibility

- First ingredient for convergence: eventual visibility
 - each write becomes eventually visible, i.e.,
 - a write can take some time to become visible ...
... but sooner or later it will become visible everywhere

Eventual visibility

For any write w in an infinite history

$$\left| \{o \mid w \xrightarrow{\text{vis}} o\} \right| < \infty$$

Convergence

► Strong Convergence

- two reads that see same write operations return same value

StrongConvergence $\doteq \forall$ reads r, r' .

$$\{\text{write } w \mid w \xrightarrow{\text{vis}} r\} = \{\text{write } w \mid w \xrightarrow{\text{vis}} r'\} \Rightarrow \\ \text{value}(r) = \text{value}(r')$$

► Convergence

- if we keep issuing reads that see same write operations ...
... the reads will eventually return same value

► The best, and what normally happens is strong convergence

- with automatic conflict resolution, as writes become visible
- why would we want otherwise?
- only if conflict resolution was manual or postponed

Eventual consistency and strong eventual consistency

- ▶ Strong eventual consistency (SEC) means having both
 - eventual visibility
 - strong convergence
- ▶ Eventual consistency (EC) means having both
 - eventual visibility
 - convergence
- ▶ Remarks
 - EC should have been defined as the SEC variant
 - using “strong” in the name is a source of confusion with strong consistency models

Combining order and convergence

- ▶ Order and convergence guarantees can be combined
- ▶ Causal+ consistency is the combination of
 - causal consistency
 - strong convergence

Ensuring the combination of causal consistency and strong convergence is the standard design goal for scalable highly available replicated datatypes

Strong consistency models

- ▶ Sequential consistency and Linearizability
- ▶ Same semantics as if all operations occurred in a sequence
- ▶ As if there is some arbitration such that
 - arbitration is compatible with session order
 - semantics as if all operations were issued by a single processso $\subseteq \text{vis} = \text{ar}$
- ▶ For sequential consistency “real time” does not matter
- ▶ Linearizability ensures a stronger guarantee
 - as if each operation occurs instantly
 - in a point in time between invocation and return time

The CAP theorem

The CAP theorem

From strong Consistency, Availability, and Partition tolerance, we can have at most two

- ▶ As partitions happen, and cannot be avoided, we can either design AP systems or CP systems
- ▶ CP systems
 - aim typically for linearizability
 - may become unavailable (stall) under network partitions
- ▶ AP systems
 - operations remain available, even when there are partitions
 - at most we aim for causal consistency and strong convergence
 - AP more suitable for truly global large scale systems