

# Operation-based CRDTs

Paulo Sérgio Almeida

March 2023



Universidade do Minho

## Conflict-free Replicated Data Types

Operation-based CRDTs

# Conflict-free Replicated Data Types

- ▶ Provide operations, like standard abstract datatypes
- ▶ Each datatype object replicated and accessed locally
  - Mutator operations update state
  - Query operations look at state and return result
- ▶ Information propagated to other replicas asynchronously
- ▶ Object highly available even under partitions

# The C from CRDTs

- ▶ C... Replicated Data Types
  - Convergent?
  - Conflict-free?
  - Commutative?
- ▶ Convergence while resolving conflicts
- ▶ Replicas keep converging; world does not have to stop
- ▶ Conflicts are dealt with semantically: spec of datatype
- ▶ Availability is achieved by forgoing total orders
- ▶ Concurrent operations will become visible in different orders
- ▶ Some confusion about what is commutative
  - some operations are not (semantically) commutative
  - effect of executing concurrent operations must be

# Operation-based vs state-based approaches

- ▶ State-based approaches
  - propagate replica states
  - detect mutual inconsistency
  - reconcile (merge) concurrent replicas
  - anti-entropy by opportunistic, “background” communication
  - can be made more incremental by delta-state approach
- ▶ Operation-based approaches
  - propagate information about operations
  - use a reliable messaging algorithm for propagation
  - need ordering guarantees (typically causal)
- ▶ Here we address operation-based CRDTs

## Conflict-free Replicated Data Types

### Operation-based CRDTs

# Operation-based CRDTs

- ▶ Core concept
  - send operations, not state, to other replicas
  - operations applied at each replica
- ▶ Uses reliable causal broadcast
  - ensures exactly once for non-idempotent operations
  - respects order of causally dependent operations
- ▶ What about non-commutative operations?
  - applied in different orders would lead to divergence
  - dealt with by sending more than just the operation

# Standard execution model of op-based CRDTs

- ▶ **Prepare** performed at replica where operation is invoked
  - looks at state and op
  - does not have side effects (on abstract state)
  - returns message to be sent
- ▶ Message disseminated with reliable causal broadcast
- ▶ Upon message delivery, **effect** is applied at each replica
  - assumes immediate self-delivery on sender replica
- ▶ Effect designed to be commutative for concurrent ops
  - ensures convergence under different application orders



## Simple CRDTs – with commutative and associative ops

- ▶ If the operations are associative and commutative . . .  
. . . they can be grouped and applied in any order
- ▶ Examples:
  - Counter
  - Positive-Negative Counter (PN-Counter)
  - Grow-only set (GSet)
- ▶ Such op-based CRDTs are trivial, given exactly-once delivery
- ▶ Not even FIFO is needed for convergence
- ▶ Causal delivery normally used to achieve causal consistency

# Counter

**CRDT state:**

$n : \mathbb{N} = 0$

**query** value() :  $\mathbb{N}$

**return**  $n$

**update** inc()

**prepare**

**return** inc

**effect** inc

$n \leftarrow n + 1$

# Positive-negative counter (PNCounter)

**CRDT state:**

$v : \mathbb{Z} = 0$

**query** value() :  $\mathbb{Z}$

**return**  $v$

**update** inc()

**prepare**

**return** inc

**effect** inc

$v \leftarrow v + 1$

**update** dec()

**prepare**

**return** dec

**effect** dec

$v \leftarrow v - 1$

## Grow-only set ( $\text{Gset}\langle E \rangle$ )

**CRDT state:**

$s : \mathcal{P}(E) = \emptyset$

**query** elements() :  $E$

**return**  $s$

**query** contains( $e : E$ ) :  $\mathbb{B}$

**return**  $e \in s$

**update** add( $e : E$ )

**prepare**

**return** (add,  $e$ )

**effect** (add,  $e$ )

$s \leftarrow s \cup \{e\}$

## CRDTs with non-commutative operations

- ▶ Example: set with add and remove operations:

$$\text{add}(v, \text{rmv}(v, s)) \neq \text{rmv}(v, \text{add}(v, s))$$

- ▶ Concurrent operations can be delivered in different orders
- ▶ Applying them to a sequential datatype would cause divergence
- ▶ CRDT cannot be simply the sequential datatype
- ▶ Effect must be defined to be commutative for concurrent ops

# Defining concurrent semantics

- ▶ We must define how to handle conflicts
- ▶ Example: given a concurrent add and remove, which will “win”
- ▶ In general many options possible
- ▶ An elegant concept: act on observed (visible) ops

If some op cancels others, do it only to observed ops

- ▶ Example: an observed-remove set
  - has add and remove ops
  - a remove only cancels adds that are visible to it
  - with causal delivery, cancels adds in the causal past
  - concurrent adds will not be canceled, and will “win”
- ▶ Used elsewhere:
  - an observed-reset counter: resets cancel observed incs

# Observed-remove set ( $\text{ORSet}\langle E \rangle$ ), vanilla, sketch

## CRDT state:

$s : \mathcal{P}(E \times \dots) = \emptyset$

**query**  $\text{elements}() : E$   
**return**  $\{e \mid (e, \_) \in s\}$

**query**  $\text{contains}(e : E) : \mathbb{B}$   
**return**  $\exists x. (e, x) \in s$

**update**  $\text{add}(e : E)$

**prepare**

**let**  $u = [\text{some unique id}]$

**return**  $(\text{add}, e, u)$

**effect**  $(\text{add}, e, u)$

$s \leftarrow s \cup \{(e, u)\}$

**update**  $\text{remove}(e : E)$

**prepare**

**let**  $r = \{(x, u) \in s \mid x = e\}$

**return**  $(\text{rmv}, r)$

**effect**  $(\text{rmv}, r)$

$s \leftarrow s \setminus r$

- Assumes reliable causal delivery

## Observed-remove set, issues, improvements

- ▶ Sketch assumes generation of unique ids; how?
  - can use counter per replica, incremented at each add
  - unique id obtained as pair (replica id, counter)
  - sometimes called a “dot” (from Dotted Version Vectors)
  - counter not used by queries or effect; auxiliary state
- ▶ Most operations need set traversal
  - solution: use a map from elements to sets of ids
- ▶ Adds accumulate entries, even if element already present
  - solution: replace current entries for given element
- ▶ In remove, prepare sends element repeated in each pair
  - solution: collect set of ids separately



## Observed-remove set ( $\text{ORSet}\langle E \rangle$ ), optimized

### types:

$\mathbb{I}$ , set of replica identifiers

### CRDT state:

$m : E \rightarrow \mathcal{P}(\mathbb{I} \times \mathbb{N}) = \emptyset$

$c : \mathbb{N} = 0$ , auxiliary state

**query**  $\text{elements}() : E$

**return**  $\text{dom } m$

**query**  $\text{contains}(e : E) : \mathbb{B}$

**return**  $m[e] \neq \emptyset$

**update**  $\text{add}(e : E)$

**prepare**

$c \leftarrow c + 1$

**return**  $(\text{add}, e, (i, c), m[e])$

**effect**  $(\text{add}, e, d, r)$

$m[e] \leftarrow m[e] \setminus r \cup \{d\}$

**update**  $\text{remove}(e : E)$

**prepare**

**return**  $(\text{remove}, e, m[e])$

**effect**  $(\text{remove}, e, r)$

$m[e] \leftarrow m[e] \setminus r$

- ▶ Algorithm for replica  $i \in \mathbb{I}$
- ▶ Assumes reliable causal delivery
- ▶ Map stores only non-“bottom” values (non-empty sets)
- ▶ Map returns  $\emptyset$  for unmapped keys

# CRDTs non-equivalent to serialized execution

- ▶ Most CRDTs resolve conflicts aiming to achieve behavior equivalent to some sequential execution
- ▶ In some cases the CRDT:
  - has behavior not possible by any sequential execution
  - the interface itself is different from sequential datatype
- ▶ Most well know case: multi-value register
  - made popular by the Amazon Dynamo paper
  - register keeps set of most recent concurrent writes
  - a read returns that set
  - a write overwrites the set into a singleton

# Multi-value register (MVReg $\langle E \rangle$ )

## types:

$\mathbb{I}$ , set of replica identifiers

## CRDT state:

$s : \mathcal{P}(E \times (\mathbb{I} \times \mathbb{N})) = \emptyset$

$c : \mathbb{N} = 0$ , auxiliary state

**query** read() :  $\mathcal{P}(E)$

**return**  $\{e \mid (e, \_) \in s\}$

**update** write( $e : E$ )

**prepare**

$c \leftarrow c + 1$

**let**  $r = \{d \mid (\_, d) \in s\}$

**return** (write, ( $e, (i, c)$ ),  $r$ )

**effect** (write,  $v, r$ )

$s \leftarrow \{(e, d) \in s \mid d \notin r\} \cup \{v\}$

- ▶ Algorithm for replica  $i \in \mathbb{I}$
- ▶ Assumes reliable causal delivery