

Large Scale Distributed Systems

José Orlando Pereira

Departamento de Informática
Universidade do Minho



Searching

- Store and translate arbitrary keys to values (i.e. $\text{Map}\langle K, V \rangle$)
 - Large number of (k, v) pairs
 - Large number n of nodes
- Abstracted to $\text{Map}\langle \text{byte}[], \text{InetSocketAddress} \rangle$
 - Key is hashed to a binary string (e.g., SHA-1 with 160 bits) with a uniform random distribution
 - Value is the the address of the node holding (k, v)

Flooding



- Broadcast query to all nodes using an epidemic algorithm
 - Nodes involved in each query: $O(n)$
 - State in each node for the overlay network: $O(\log n)$
- Not scalable as the number of queries grows
- Idea: Route each query towards the correct node...
How?

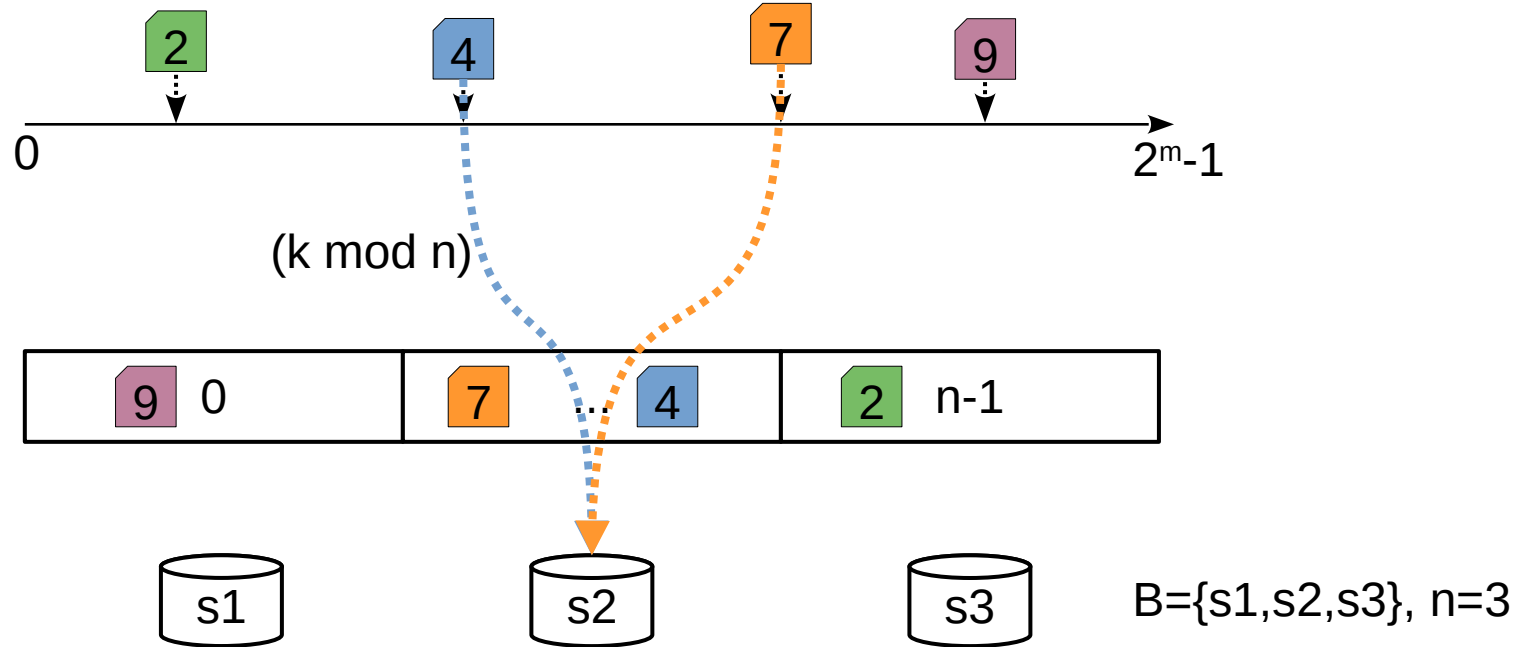
Naive distributed hashing

- Our key is already randomly distributed
- 1 node = 1 bucket
 - $(k \bmod n)$ gives the bucket number
 - Nodes involved in each query: $O(1)$
- Need to map buckets to server addresses
 - Routing state: $O(n)$
- Node churn: How to update data placement and routing information when nodes enter/leave?


Hashing properties

- Balance: Proportion of items in each bucket should be $O(1/n)$
- Monotonicity: When there is a new bucket, items are moved only to that new bucket (not between old buckets)
- In a distributed system, there is uncertainty about location of keys:
 - Wrong opinions about some key
 - Wrong opinions about some bucket

Naive distributed hashing




Naive distributed hashing



- Balance: Good, assuming random k
- Monotonicity: Bad
- Uncertainty while items move


Consistent hashing (key idea)

- Hash keys and bucket ids into the same space and assign by distance
- Balance is not perfect: Distribution and extremes




Consistent hashing (key idea)

- Monotonicity: Good! Keys are moved only to the new bucket
- Less uncertainty but all keys move to/from at most two other buckets




Consistent hashing

- Assume a circular space:
 - $\text{mod } 2^m$
- Clockwise distance
 - Successor bucket
 - Direct comparison, no need for arithmetic




Consistent hashing

- Multiple buckets in each node (virtual nodes)
- Same mean, lower variance



Consistent hashing

- On change: Keys move to/from k other nodes




Consistent hashing

- Balance: Good balance with $k \sim O(\log n)$ buckets in each node
- Monotonicity: Good! Keys are moved only to the new bucket
- Nodes involved in each query: $O(1)$
- Routing state: $O(n)$ or $O(n \log n)$ with virtual nodes


Partitioned routing state (Kelips)

- Split nodes in \sqrt{n} zones
- Each node keeps routing state for:
 - its zone
 - a few contact nodes in each other zone



Partitioned routing state (Kelips)

- Lookup:
 - 1 hop to some contact in zone
 - 1 hop to node
- Use epidemic dissemination to update routing tables within zones




Partitioned routing state (Kelips)

- Balance and Monotonicity unchanged
- Nodes involved in each query: still $O(1)$
- Routing state: down to $O(\sqrt{n})$
- Background traffic to synchronize routing state in each zone
- Uncertainty, due to synchronization and lost contact nodes


Scalability

- Can we do $O(1)$ state?
- Yes: Keep a pointer to the successor node
- Lookup is $O(n)$
 - average $n/2$ hops
- Fragile...



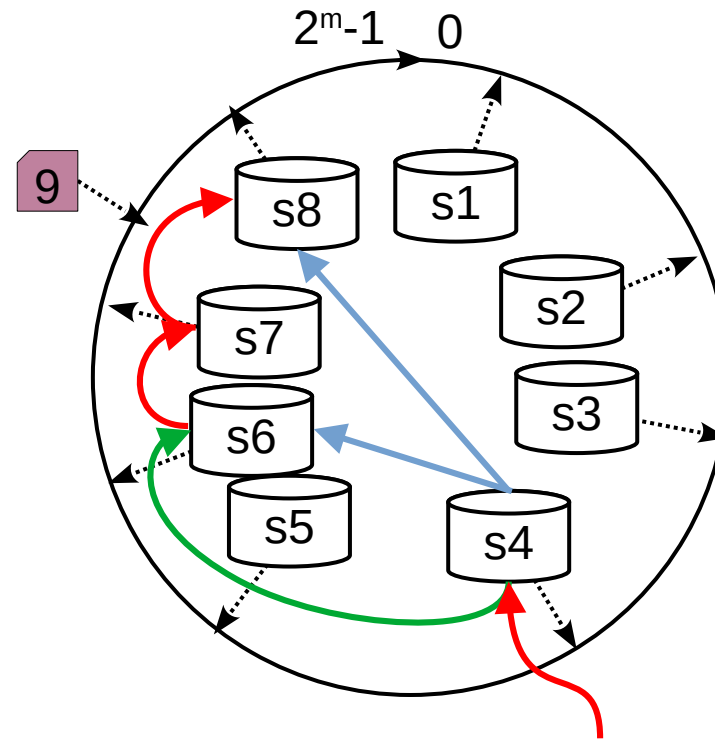
Scalability

- Idea: Keep a shortcut to “the other side of the ring”
- Lookup is $O(n)$
 - 2 pointers
 - average $n/4$ hops




Scalability

- What about items in the first half?
- Split again
- Lookup is $O(n)$
 - 3 pointers
 - average $n/8$ hops!
- Can use more than one shortcut...



Scalability

- How many times can we split it?
 - $m!$
- Routing state:
 $O(m) \sim O(\log n)$
- Lookup: $O(\log n)$
- How to setup and maintain these pointers?




Chord

- Node p keeps finger pointers for:
 - $q = \text{succ}(p + 2^i)$, for $0 \leq i < m$
 - $i = 0$ gives the direct successor of p
- Lookup $\text{succ}(k)$ at p :
 - if k in local interval, then return p
 - else forward to $q = \text{succ}(p + 2^i)$ such that:
 - if exists, largest i with $q \leq k$
 - else with $i = 0$

Chord

- A node q joins the ring by looking up its successor r
- Uses this to set its first pointer ($i = 0$)




Chord

- The ring is repaired by running two procedures:
 - Stabilize: Runs periodically and exchanges information with the (currently best) successor:
 - Informs successor of a possible new predecessor
 - Obtains or updates finger pointers, possibly discovering a better successor
 - Rectify: Runs when a new predecessor is discovered and selects the candidate with the largest value


Chord

- Node q stabilizes:
 - Informs r of its id
 - Learns the successor's predecessor p and adopts it as successor if $p > q$
 - Not in this case...
 - Learns r's pointer table and initializes its own




Chord

- Upon learning of a new predecessor, node r rectifies:
 - checks if $q > p$ and sets predecessor pointer to q




Chord

- Node p stabilizes:
 - Informs r of its id
 - Learns the successor's predecessor q and adopts it as successor if $q > p$
 - Yes in this case!
 - Learns r 's pointer table and updates its own



Chord

- Upon learning of a new predecessor, node q rectifies:
 - no current predecessor, so it adopts p



Chord

- If q fails or leaves: Periodical stabilization at p will repair the ring
- Why does it work?
 - Stabilize selects the smallest of successor candidates
 - Rectify selects the largest of predecessor candidates
 - It converges to the correct order as long as there are no ties (i.e. nodes with the same id)

Chord

- If state is volatile, it is lost when a node fails or disconnects
 - Caching
- If not, it needs to be replicated to $f+1$ nodes
 - Each node keeps $f+1$ predecessor pointers
 - State is replicated forward by owner node
- Note that replication also improves Balance, in the same way as virtual nodes!

Summary

	Balance	Monotonicity	Lookup	State
Naive hashing	Perfect	No	$O(1)$	$O(n)$
Consistent hashing	Good(*)	Yes	$O(1)$	$O(n)$
<i>Kelips</i> DHT	Good(*)	Yes	$O(1)$	$O(\sqrt{n})$
Sequential cons. hashing	Good(*)	Yes	$O(n)$	$O(1)$
<i>Chord</i> DHT	Good(*)	Yes	$O(\log n)$	$O(\log n)$

(*) Very good at the expense of $\times(\log n)$ state

- Typical choices:
 - Consistent hashing in the medium / data center scale
 - Chord (and Kademlia) DHTs in the large scale

References

- D. Karger, E. Lehman, T. Leighton, M. Levine, D. Lewin, and R. Panigrahy, “Consistent hashing and random trees: Distributed caching protocols for relieving hot spots on the world wide web,” in ACM Symposium on Theory of computing, 1997 (Mainly Section 4)
<https://www.cs.princeton.edu/courses/archive/fall09/cos518/papers/chash.pdf>
- I. Stoica et al., “Chord: a scalable peer-to-peer lookup protocol for internet applications,” IEEE/ACM Trans. Netw., vol. 11, no. 1, pp. 17–32, 2003
<http://dx.doi.org/10.1109/tnet.2002.808407>

More...

- I. Gupta, K. Birman, P. Linga, A. Demers, and R. van Renesse, “Kelips: Building an Efficient and Stable P2P DHT through Increased Memory and Background Overhead,” Peer-to-Peer Systems II. pp. 160–169, 2003
http://dx.doi.org/10.1007/978-3-540-45172-3_15
- P. Zave, “How to Make Chord Correct,” arXiv [cs.DC], Feb. 23, 2015
<http://arxiv.org/abs/1502.06461>
- P. Maymounkov and D. Mazières, “Kademlia: A Peer-to-Peer Information System Based on the XOR Metric,” in Peer-to-Peer Systems, 2002
http://dx.doi.org/10.1007/3-540-45748-8_5