

Motivation


- Example: Read a photo album and its access control list
- Option 1, concurrently:
 - Read photos, then read ACL
 - Delete private photo, make album public
 - Might read private photo and see it as public
- Option 2, concurrently:
 - Read ACL, then read photos
 - Make album private, add private photo
 - Might see album as public and then read private photo

Transactions

- Read transactions:
 - Avoid missing dependencies in values read
 - Solves the problem if writes issued in the correct order
- Write transactions:
 - Ensures atomicity (mutual dependency) of values written
 - Allows any write order


Read transactions: Challenges

- Dependency information is not stored in the server or returned to clients



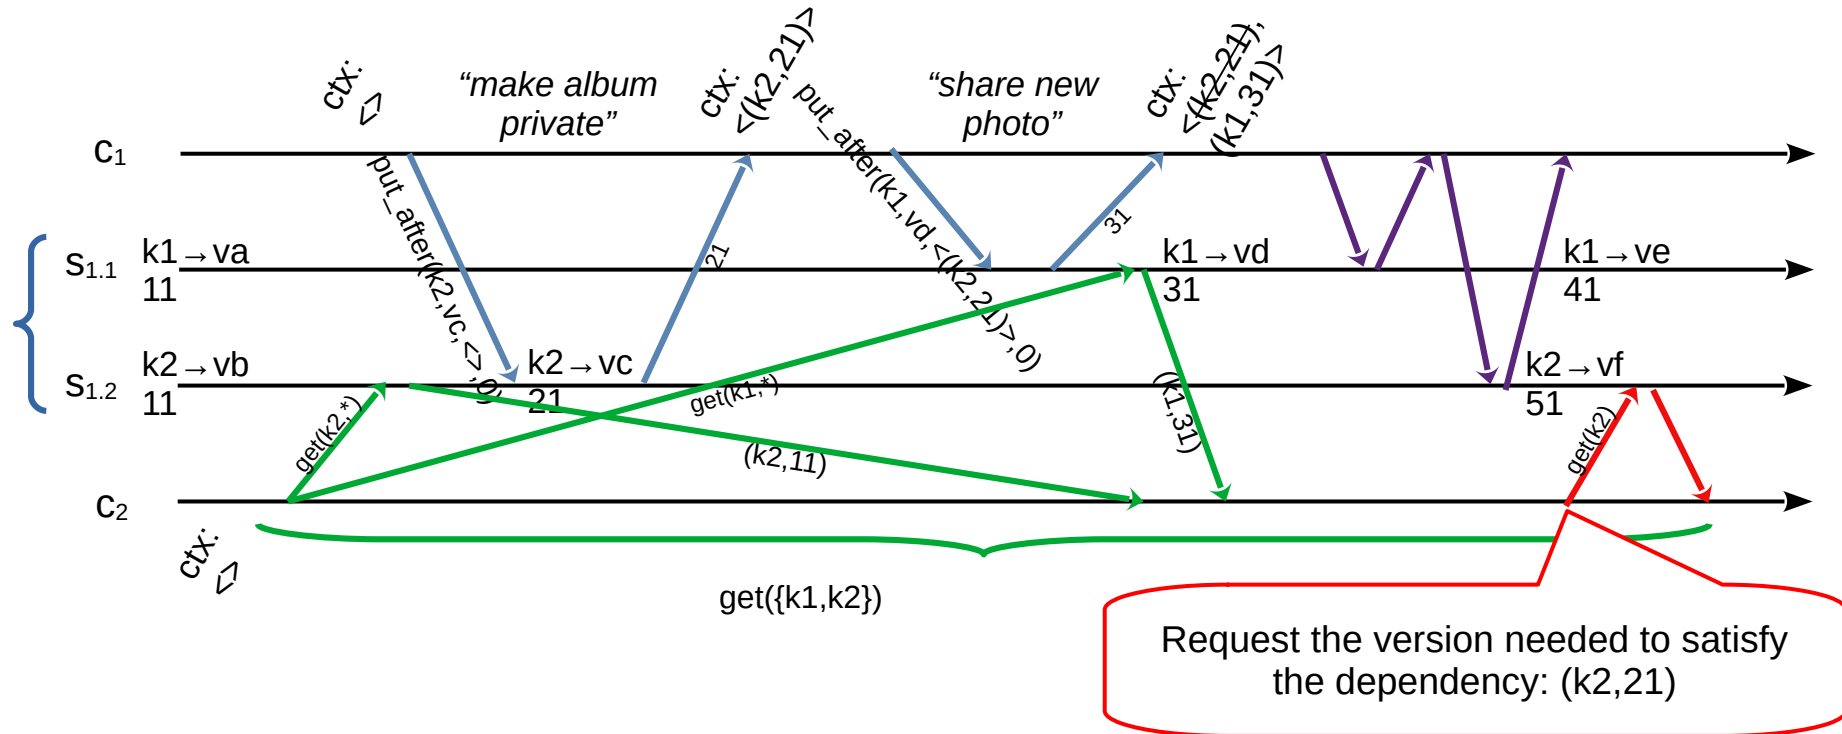
Stored dependencies

- Keep dependencies to detect what is missing
- How to get it without introducing new dependencies?



Multi-version

- Keep preceding versions of each item
- Return the exact version needed by the transaction




Read from snapshot

- How to avoid keeping detailed dependency information?
- The problem arises when reading some item that was written after the transaction has started:
 - (k1,31) in the exampleas it may introduce a dependency:
 - (k2,21) in the example
- Stable time: Latest time for which all dependencies are known

Read from snapshot

- Obtain a stable start timestamp st at the start of the transaction
- Read latest version $v \leq st$

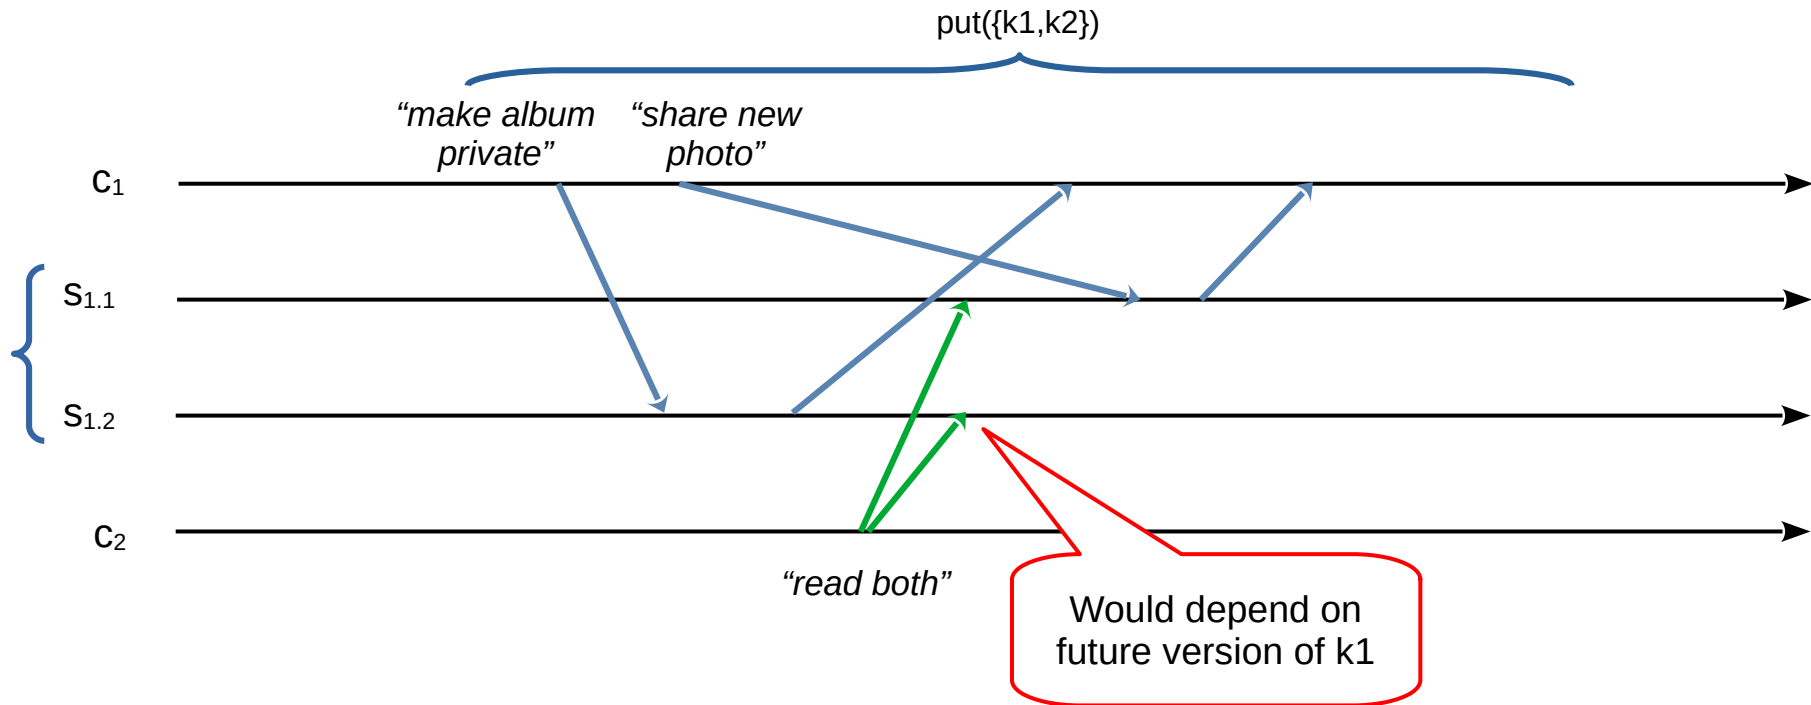


Snapshot assignment

- Start timesampt st is stable: all updates (hence, their dependencies) are available in all hosts
- Start timestamp st is greater than all previous reads and writes from the same client
 - Ensures RYOW and causality across different transactions
 - May need to wait for stability to catch up to recent updates
- The latest stable snapshot can be computed with epidemic protocols


Write transactions: Challenges

- Reads assume that all dependencies are known and committed
 - Not true with write transactions




Two phase commit

- An atomic update of two servers needs 2-phase commit
- Locking avoids reading incomplete transactions



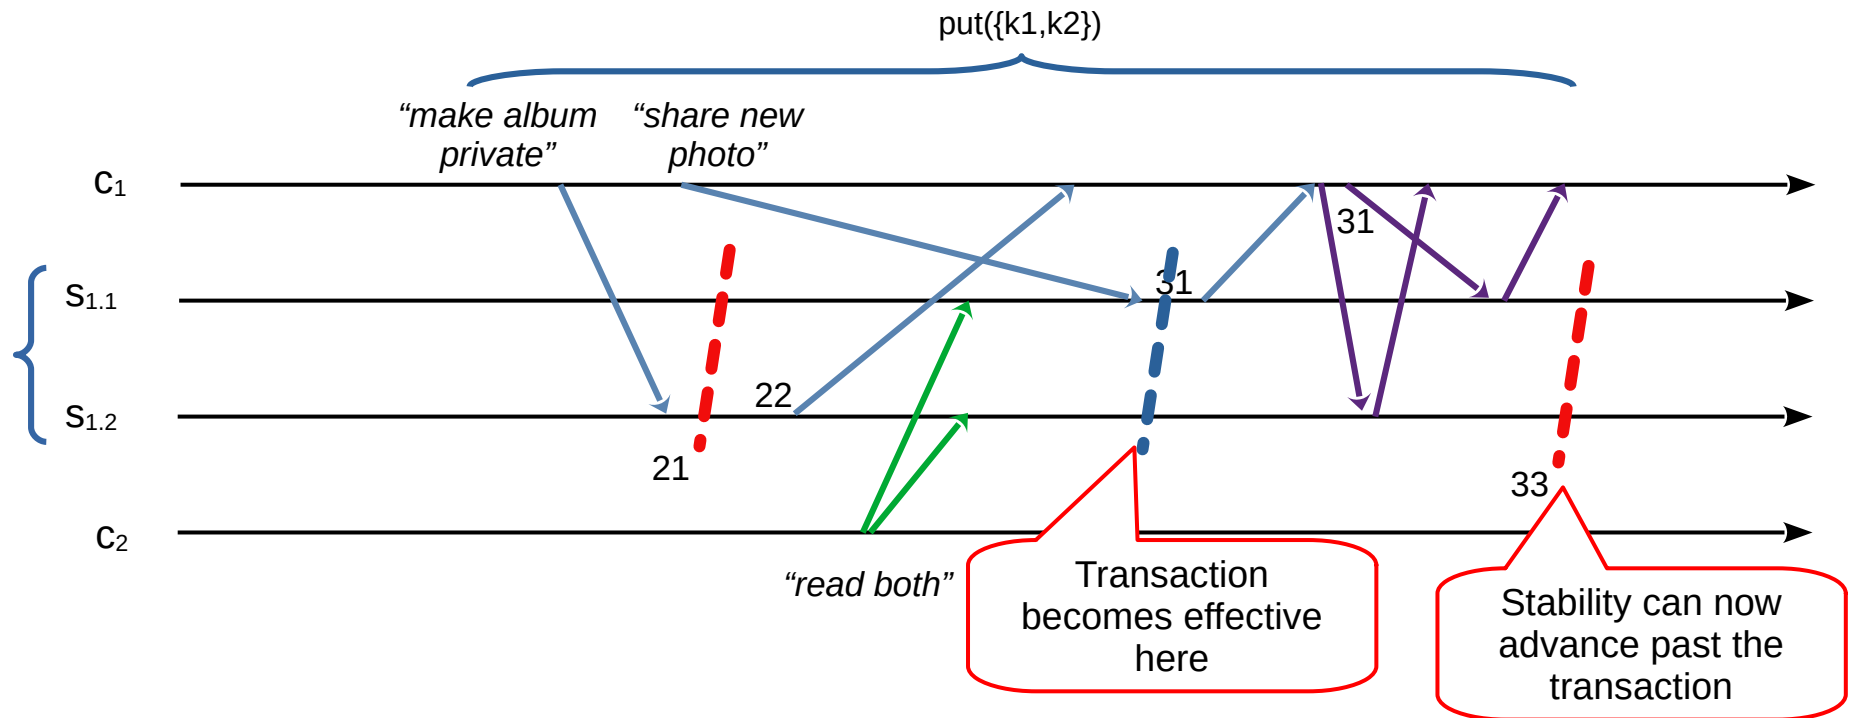
2PC with snapshots

- Each participant proposes a timestamp
- Prepared transactions block stability



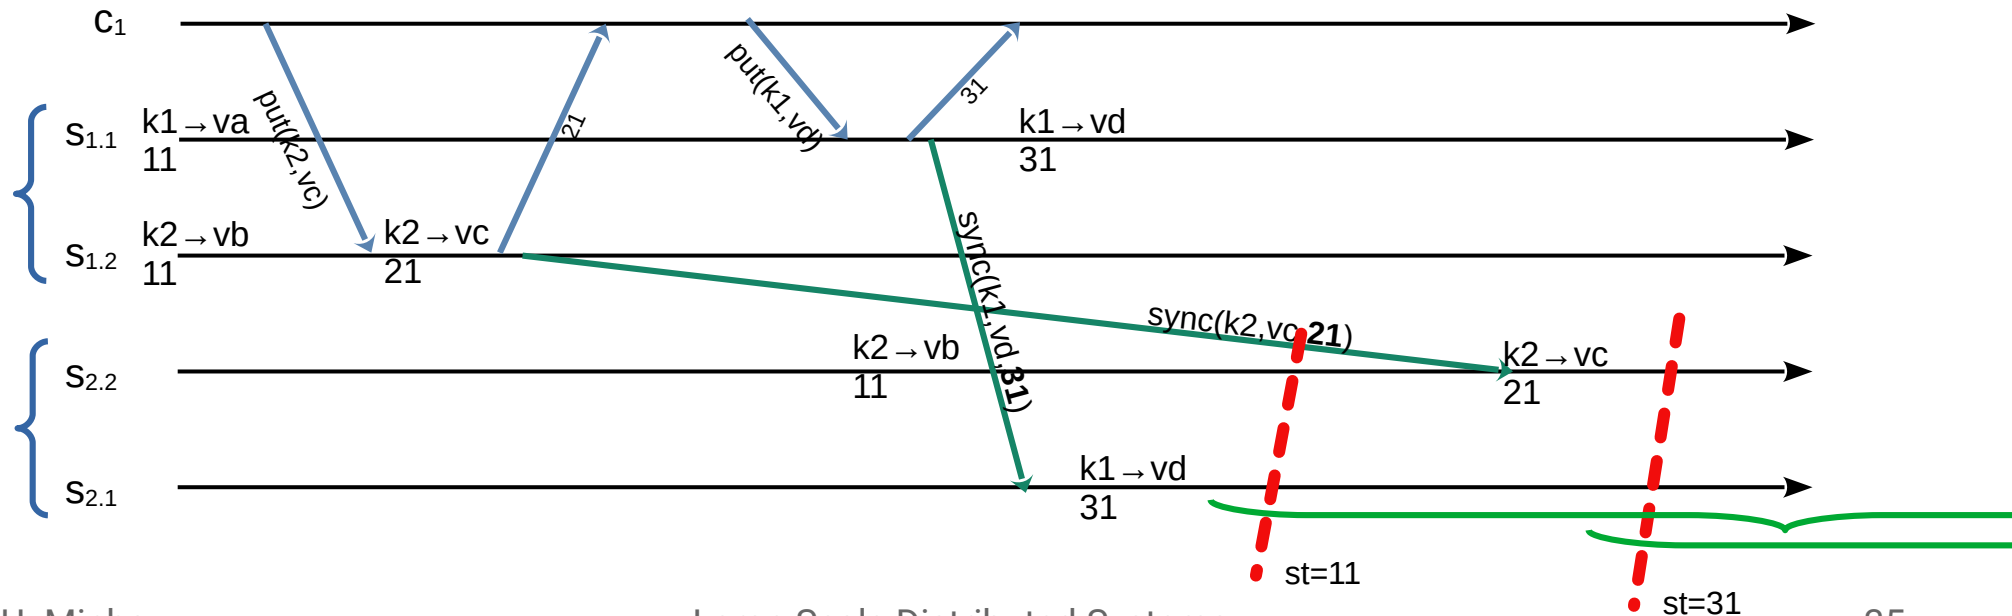
2PC with snapshots

- The latest timestamp is used as the global commit time
 - Known to be after current stable time

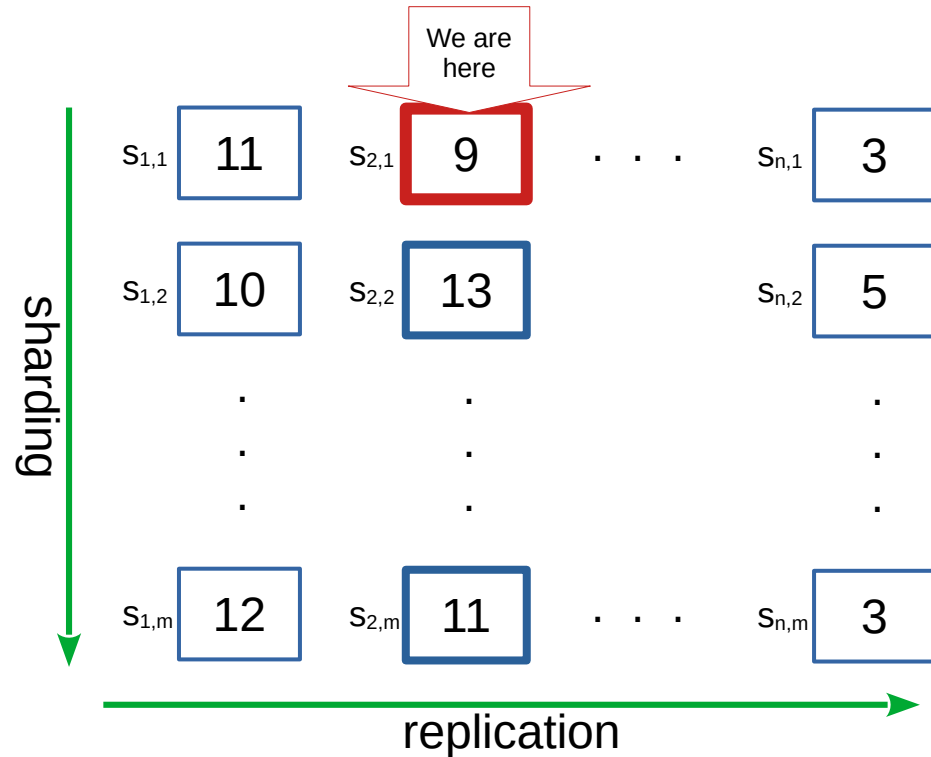


Write propagation with snapshots

- Bonus of snapshots: no need to wait before applying remote updates
 - They remain invisible until stable

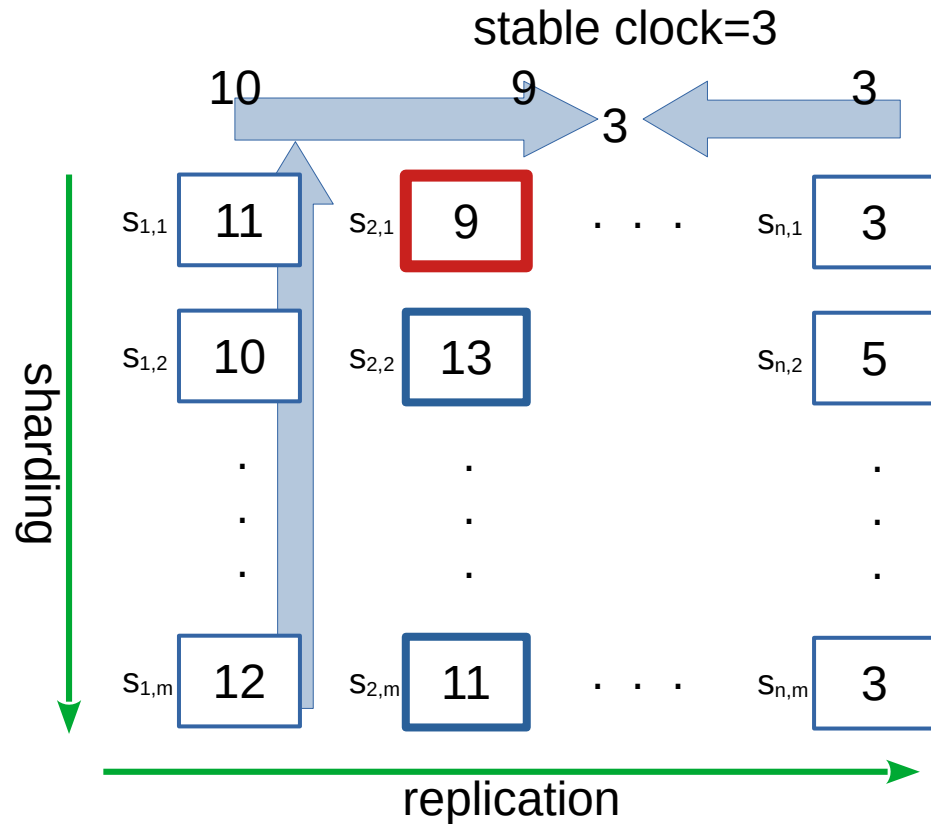


Stability



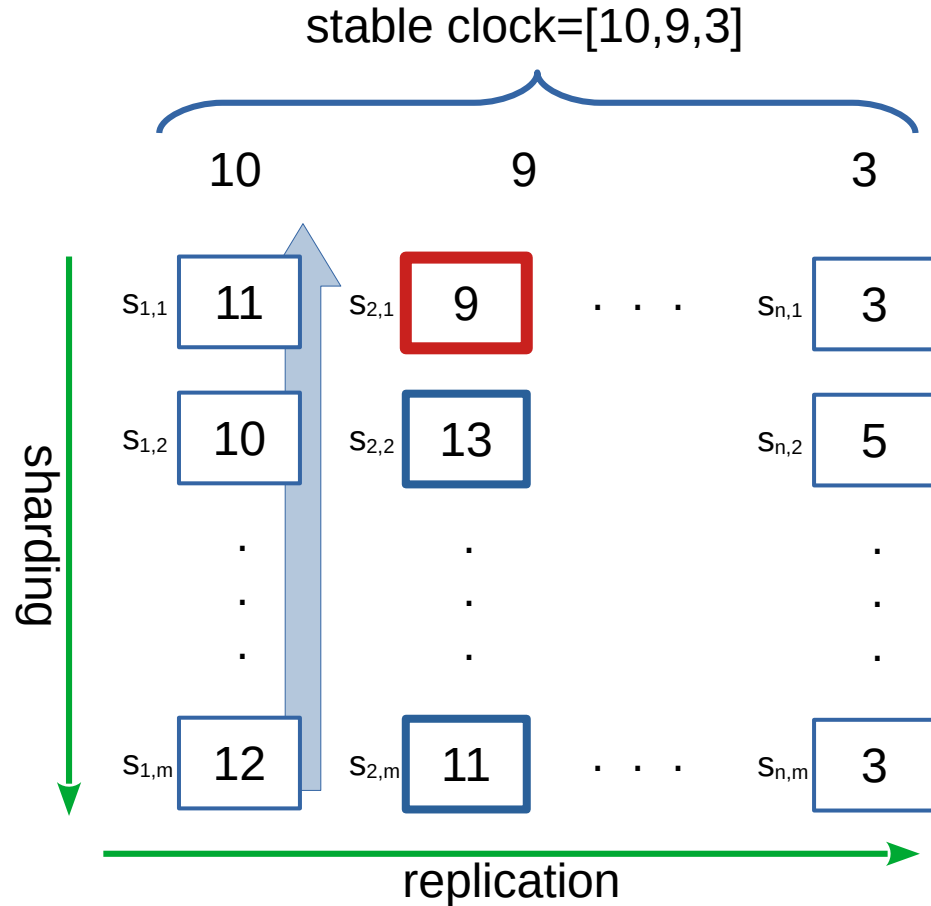
- Each server has its own clock
 - n replicas
 - m shards
- Knowledge about other servers (esp. other DCs) may be late....

Scalar timestamp




- Stable snapshot is the global minimum
- Visibility determined by oldest replica/shard
 - Blocks if a data center is disconnected
 - Not “AP”

Vector timestamp



- Stable snapshot is the minimum for each DC
- Commit timestamp is vectorial
- Values become visible when st is larger than the transactions ct
 - Remember that $st \leq \text{stable clock}$
- Does not block when DCs are partitioned

Two scalars



- Stable snapshot:
 - Local component (used for local transactions)
 - Remote component (used for others)
- Does not block local updates when DCs are partitioned
- Blocks all remote when a single DC is partitioned

Summary

- Interactive causal transactions
 - Consistent reads
 - Atomic writes
- Non-blocking reads and writes
- Trade-off: Freshness of values read waiting for stability detection

References

- D. D. Akkoorath et al., “**Cure: Strong Semantics Meets High Availability and Low Latency**,” in 2016 IEEE 36th International Conference on Distributed Computing Systems (ICDCS), Jun. 2016
<http://dx.doi.org/10.1109/ICDCS.2016.98>